

PhD Thesis

Koustuv Sinha

Acknowledgements

Abstract

Abstract in French

Contributions to Original Knowledge

Contributions of Authors

List of Figures

3.1	Data generation pipeline. Step 1: generate a kinship graph. Step 2: sample a target fact. Step 3: Use backward chaining to sample a set of facts. Step 4: Convert sampled facts to a natural language story.	7
3.2	Amazon Mechanical Turker Interface built using ParlAI which was used to collect data as well as peer reviews.	13
3.3	Illustration of how a set of facts can split and combined in various ways across sentences.	16
3.4	Noise generation procedures of CLUTRR.	17
3.5	Systematic generalization performance of different models when trained on clauses of length $k = 2, 3$ (Left) and $k = 2, 3, 4$ (Right).	26
3.6	Systematic Generalizability of different models on CLUTRR-Gen task (having 20% less placeholders and without training and testing placeholder split), when Left : trained with $k = 2$ and $k = 3$ and Right : trained with $k = 2, 3$ and 4	27

4.1	Graphical representation of the Permutation Acceptance class of metrics. Given a sample test set D_{test} with six examples, three of which originally predicted correctly (model predicts gold label), three incorrectly (model fails to predict gold label), with $n = 6$ permutations, $\Omega_{\max}, \Omega_{\text{rand}}, \Omega_{1.0}, \mathcal{P}^c$ and \mathcal{P}^f are provided. Green boxes indicate permutations accepted by the model. Blue boxes mark examples that crossed each threshold and were used to compute the corresponding metric.	40
4.2	Comparison of $\Omega_{\max}, \Omega_{\text{rand}}, \mathcal{P}^c$ and \mathcal{P}^f with the model accuracy \mathcal{A} on multiple datasets, where all models are trained on the MNLI corpus Williams et al. [2018c].	43
4.3	Ω_x threshold for all datasets with varying x and computing the percentage of examples that fall within the threshold. The top row consists of in-distribution datasets (MNLI, SNLI) and the bottom row contains out-of-distribution datasets (ANLI)	44
4.4	Average entropy of model confidences on permutations that yielded the correct results for Transformer-based models (top) and Non-Transformer-based models (bottom). Results are shown for D^c (orange) and D^f (blue). The boxes show the quartiles of the entropy distributions.	46
4.5	Length and Permutation Acceptance by Transformer-based models. . . .	48
4.6	Comparing the effect between randomizing both premise and hypothesis and only hypothesis on two Transformer-based models, RoBERTa and BART. Here, we observe the difference of Ω_{\max} is marginal in in-distribution datasets (SNLI, MNLI), while hypothesis-only randomization is worse for out-of-distribution datasets (ANLI).	49

4.7	Comparing the effect between randomizing both premise and hypothesis and only hypothesis on two Transformer-based models, RoBERTa and BART. In this figure, we compare the mean number of permutations which elicited correct response, and naturally the hypothesis-only randomization causes more percentage of randomizations to be correct.	50
4.8	BLEU-2 score versus acceptability of permuted sentences across all test datasets. RoBERTa and BART performance is similar but differs considerably from the performance of non-Transformer-based models, such as InferSent and ConvNet.	51
4.9	Example POS signature for the word ‘river’, calculated with a radius of 2. Probability of each neighbor POS tag is provided. Orange examples come from the permuted test set, and blue come from the original training data.	53
4.10	POS Tag Mini Tree overlap score and percentage of permutations which the models assigned the gold-label.	54
5.1	Perplexity of various models on Wiki 103 valid and test sets.	68
5.2	GLUE & PAWS task dev performance when finetuned on naturally (blue) and randomly ordered (orange) text, respectively, using pre-trained RoBERTa (base) models trained on different versions of BookWiki corpus.	70
5.3	GLUE results on various model ablations using BookWiki corpus.	79
5.4	BPPL scores per model per test scenario.	84
5.5	Rissanen Data Analysis Perez et al. [2021] on the GLUE benchmark and PAWS datasets. The lower minimum description length (MDL, measured in kilobits), the better the learning ability of the model.	86
5.6	Comparison among GLUE task performance from different steps in pre-training of RoBERTa on BookWiki Corpus.	88

5.7	The difference in word probabilities for stimuli in Marvin and Linzen [2018]: Simple Verb Agreement (SVA), In a sentential complement (SCM), Short VP Coordination (SVC), Long VP Coordination (LVC), Across a prepositional phrase (APP), Across a subject relative clause (ASR), Across an object relative clause (AOR), Across an object relative (no <i>that</i>) (AOR-T), In an object relative clause (IOR), In an object relative clause (no <i>that</i>) (IOR-T), Simple Reflexive (SRX), In a sentential complement (ISC), Across a relative clause (ARC), Simple NPI (SNP).	100
5.8	Linzen et al. [2016]	101
5.9	Gulordava et al. [2018b]	101

List of Tables

3.1	Statistics of the AMT paraphrases. Jaccard word overlap is calculated within the templates of each individual clause of length k	14
3.2	Human performance accuracies on CLUTRR dataset. Humans are provided the Clean-Generalization version of the dataset, and we test on two scenarios: when a human is given limited time to solve the task, and when a human is given unlimited time to solve the task. Regardless of time, our evaluators provide a score of difficulty of individual puzzles.	15
3.3	Snapshot of puzzles in the dataset for $k=2$	19
3.4	Snapshot of puzzles in the dataset for $k=3$	20
3.5	Snapshot of puzzles in the dataset for $k=4$	20
3.6	Snapshot of puzzles in the dataset for $k=5$	21
3.7	Snapshot of puzzles in the dataset for $k=6$	22
3.8	Testing the robustness of the various models when training and testing on stories containing various types of noise facts. The types of noise facts (supporting, irrelevant, and disconnected) are defined in Section . .	27
3.9	Testing the robustness of the various models when trained various types of noisy facts and evaluated on other noisy / clean facts.	29

3.10 Testing the robustness on toy placeholders of the various models when trained various types of noisy facts and evaluated on other noisy / clean facts.	30
4.1 Examples from the MNLI Matched development set. Both the original example and the permuted one elicit the same classification label (entailment and contradiction respectively) from RoBERTa (large). A simple demo is provided in an associated Google Colab notebook.	36
4.2 Statistics for Transformer-based models trained on MNLI corpus Williams et al. [2018c]. The highest values are bolded (red indicates the model most insensitive to permutation) per metric and per model class (Transformers and non-Transformers). A1*, A2* and A3* refer to the ANLI dev. sets [Nie et al., 2020].	45
4.3 Results on evaluation on OCNLI Dev set. All models are trained on OCNLI corpus Hu et al. [2020a]. Bold marks the highest value per metric (red shows the model is insensitive to permutation).	47
4.4 Human (expert) evaluation on 200 permuted examples from the MNLI matched development set. Half of the permuted pairs contained shorter sentences and the other, longer ones. All permuted examples were assigned the gold label by RoBERTa-Large.	54
4.5 NLI Accuracy (\mathcal{A}) and Permutation Acceptance metrics (Ω_{\max}) of RoBERTa when trained on MNLI dataset using vanilla (V) and Maximum Random Entropy (ME) method.	56
5.1 BLEU-2,3,4 scores (mean and std dev) on a sample of 1M sentences drawn from the corpus used to train \mathcal{M}_1 , \mathcal{M}_2 , \mathcal{M}_3 and \mathcal{M}_4 compared to \mathcal{M}_N	64

5.2	GLUE and PAWS-Wiki dev set results on different RoBERTa (base) models trained on variants of the BookWiki corpus (with mean and std). The top row is the original model, the middle half contains our primary models under investigation, and the bottom half contains the baselines.	69
5.3	Unlabeled Attachment Score (UAS) (mean and std) on the dependency parsing task (DEP) on two datasets, UD EWT and PTB, using the Pareto Probing framework Pimentel et al. [2020a].	74
5.4	SentEval Probing Conneau et al. [2018], Conneau and Kiela [2018] results (with mean and std) on different model variants.	75
5.5	Mean (and std) non-parametric probing accuracy on different datasets. * indicates rebalanced datasets, see ?? for more details.	77
5.6	GLUE and PAWS-Wiki dev set results on different ablations of the RoBERTa (base) models, trained on variants of the BookWiki corpus (with mean and std dev). The top row is the original model, the middle half contains the sentence randomization models, and the bottom half contains the ablations.	78
5.7	Reconstruction experiments on shuffled word order sentences by fixing the same seed for every sentence (\mathcal{M}_1) and having different seed for different shards of the corpus (\mathcal{M}_1^*). We observe minimal difference in the downstream GLUE and PAWS scores.	80
5.8	$\Delta_{\{D_i\}}(\mathcal{T})$, scaled by a factor of 100 for GLUE and PAWS tasks.	81

5.9 Fine-tuning evaluation by varying different sources of word order (with mean and std dev). We vary the word order contained in the pre-trained model ($\mathcal{M}_N, \mathcal{M}_1, \mathcal{M}_{UG}$); in fine-tuning training set (natural and shuffled); and in fine-tuning evaluation (natural and shuffled). Here, <i>shuffled</i> corresponds to unigram shuffling of words in the input. In case of fine-tune evaluation containing shuffled input, we evaluate on a sample of 100 unigram permutations for each data point in the dev set of the corresponding task.	82
5.10 Unlabeled Attachment Score (UAS) on Dependency parsing task on two datasets, UD EWT and PTB, using the Second order Tree CRF Neural Dependency Parser Zhang et al. [2020]	83
5.11 Accuracy on the part-of-speech labelling task (POS) on two datasets, UD EWT and PTB, using the Pareto Probing framework Pimentel et al. [2020b].	89
5.12 Accuracy on the dependency arc labelling task (DAL) on two datasets (with mean and std dev), UD EWT and PTB, using the Pareto Probing framework Pimentel et al. [2020a].	90
5.13 Pareto Hypervolume of dependency parsing task (DEP) on two datasets (with mean and std dev), UD EWT and PTB, using the Pareto Probing framework Pimentel et al. [2020b].	91
5.14 Linzen et al. [2016] stimuli results in raw accuracy. Values in parenthesis reflect the standard deviation over different seeds of pre-training. Values in square brackets indicate the mean probability difference among correct and incorrect words.	91
5.15 Gulordava et al. [2018b] stimuli results in raw accuracy. Values in parenthesis reflect the standard deviation over different seeds of pre-training. Values in square brackets indicate the mean probability difference among correct and incorrect words.	91

5.16 Marvin and Linzen [2018] stimuli results in raw accuracy. Values in parenthesis reflect the standard deviation over different seeds of pre-training. Values in square brackets indicate the mean probability difference among correct and incorrect words. Abbreviations: Simple Verb Agreement (SVA), In a sentential complement (SCM), Short VP Coordination (SVC), Long VP Coordination (LVC), Across a prepositional phrase (APP), Across a subject relative clause (ASR), Across an object relative clause (AOR), Across an object relative (no <i>that</i>) (AOR-T), In an object relative clause (IOR), In an object relative clause (no <i>that</i>) (IOR-T), Simple Reflexive (SRX), In a sentential complement (ISC), Across a relative clause (ARC), Simple NPI (SNP).	92
5.17 Linzen et al. [2016] stimuli results in raw accuracy on original, unbalanced data. Values in parenthesis reflect the standard deviation. S/P reflects the count of correct singular and plural focus words.	93
5.18 First 10 lines from the BookWiki corpus, and their respective n-gram permutations.	94
5.19 Fine-tuning hyperparam Learning rate of each model for each task in GLUE and PAWS	95
5.20 Finetuning hyperparam batch size of each model for each task in GLUE and PAWS	95

Contents

1	Introduction	1
2	Background	2
2.1	Early methods for text representation	2
2.2	Neural Inductive bias of text representation	2
2.2.1	Feed Forward Neural Networks	2
2.2.2	Recurrent Neural Networks	2
2.2.3	Transformer Models	2
2.3	Pre-training and the advent of Large Language Models	2
2.4	Systematicity and Generalization	3
2.4.1	Definitions	3
2.4.2	Tasks	3
3	Understanding semantic generalization through systematicity	4
3.1	Technical Background	6
3.1.1	Notations and Terminology	6
3.2	Overview and construction of CLUTRR	7
3.2.1	Graph generation	8
3.2.2	Backward chaining	10
3.2.3	Adding natural language	10

Paraphrasing using Amazon Mechanical Turk	10
Reusability and composition	13
3.2.4 AMT Template statistics	14
3.2.5 Human performance	14
3.2.6 Representing the question and entities	16
3.3 Experimental Setups	17
3.3.1 Systematic generalization	17
3.3.2 Robust Reasoning	18
3.3.3 Generated Datasets	19
3.4 Evaluated Models	19
3.4.1 Hyperparameters	23
3.5 Results	24
3.5.1 Systematic Generalization	25
3.5.2 The benefit of structure	26
3.5.3 Robust Reasoning	27
Learning from noisy data	28
Learning with synthetic placeholders	29
3.6 Related Work	30
3.6.1 Reading comprehension benchmarks	30
3.6.2 Systematic generalization	31
3.6.3 Question-answering with knowledge graphs	31
3.7 Discussion	31
3.8 Follow-up findings in the community	33
4 Quantifying syntactic generalization using word order	34
4.1 Technical Background	37
4.2 Experimental Setup	39
4.2.1 Constructing the permuted dataset.	39

4.2.2	Defining Permutation Acceptance	40
4.3	Evaluated Models	42
4.4	Results	42
4.4.1	Models accept many permuted examples	42
Investigating other Ω values	44	
4.4.2	Models are very confident	46
4.4.3	Similar artifacts in Chinese NLU	47
4.4.4	Other Results	48
4.5	Analysis	49
4.5.1	Analyzing Syntactic Structure Associated with Tokens	49
4.5.2	Human Evaluation	53
4.5.3	Training by Maximizing Entropy	55
4.6	Related Work	55
4.6.1	Models appear to have acquired syntax	56
4.6.2	Models appear to struggle with syntax	57
4.6.3	Insensitivity to Perturbation	58
4.6.4	NLI Models are very sensitive to words	59
4.7	Discussion	59
4.8	Follow-up findings in the community	60
5	Probing syntax understanding through distributional hypothesis	61
5.1	Technical Background	63
5.2	Experimental Setup	63
5.2.1	Sentence word order permutation	63
5.2.2	Corpus word order bootstrap resample	64
5.2.3	Further baselines	66
5.3	Evaluated Models & Tasks	66
5.3.1	Pre-training details	67

5.3.2	Fine-tuning tasks	67
5.4	Results	69
5.4.1	Downstream task results	69
Word order permuted pre-training	69	
Word order permuted fine-tuning	71	
5.4.2	Probing results	73
Parametric Probing	73	
SentEval Probes	75	
Non-Parametric Probing	76	
5.5	Analysis	78
5.5.1	Word-order pre-training ablations	78
5.5.2	Measuring Relative difference	80
5.5.3	Fine-tuning with randomized data	81
5.5.4	Dependency parsing using Second order Tree CRF Neural Dependency Parser	83
5.5.5	Perplexity analysis	84
5.5.6	The usefulness of word order	85
5.5.7	At what point do models learn word order during pre-training? .	87
5.5.8	More results from Syntactic Probes	88
5.5.9	Non parametric probes	89
5.6	Related Work	93
5.6.1	Sensitivity to word order in NLU	93
5.6.2	Randomization ablations	96
5.6.3	Synthetic pre-training	96
5.6.4	On the utility of probing tasks	97
5.7	Discussion	98
5.8	Follow-up findings in the community	99

6 Measuring systematic generalization by exploiting absolute positions	102
6.1 Technical Background	102
6.2 Systematic understanding of absolute position embeddings	102
6.3 Related Work	102
6.4 Experiments	102
6.5 Discussion	102
7 Conclusion	103
7.1 Summary	103
7.2 Limitations	103
7.3 Future Work	103
Bibliography	104
Glossary	143
Acronyms	143
8 Appendix	144
8.1 Org mode auto save	144
8.2 Remove “parts” from report	144
8.3 Add newpage before a heading	145
8.4 Glossary and Acronym build using Latexmk	145
8.5 Citation style buffer local	146
8.6 Org latex compiler options	146

Chapter 1

Introduction

Central Theme of the thesis : Understanding systematicity in pre-trained language models through semantic and syntactic generalization.

In this thesis I discuss my work on understanding systematicity in pre-trained language models.

Chapter 2

Background

2.1 Early methods for text representation

2.2 Neural Inductive bias of text representation

2.2.1 Feed Forward Neural Networks

2.2.2 Recurrent Neural Networks

2.2.3 Transformer Models

Large Language Models (LLMs) are the state-of-the-art in language models, which are based on Transformers.

2.3 Pre-training and the advent of Large Language Models

Success of pre-training and scale

2.4 Systematicity and Generalization

2.4.1 Definitions

1. Systematicity
2. Word Order Sensitivity

2.4.2 Tasks

Chapter 3

Understanding semantic generalization through systematicity

Natural Language Understanding (NLU) systems have been extremely successful at reading comprehension tasks, such as question answering (QA) and natural language inference (NLI). These tasks typically test for semantic generalization, where a model has to understand the meaning of the input sentence / passage in order to perform the given task. An array of existing datasets are available for these tasks. This includes datasets that test a system's ability to extract factual answers from text [Rajpurkar et al., 2016a, Nguyen et al., 2016, Trischler et al., 2016, Mostafazadeh et al., 2016, Su et al., 2016], as well as datasets that emphasize commonsense inference, such as entailment between sentences [Bowman et al., 2015b, Williams et al., 2018d].

However, there are growing concerns regarding the ability of NLU systems—and neural networks more generally—to generalize in a systematic and robust way [Bahdanau et al., 2019, Lake and Baroni, 2018, Johnson et al., 2017]. For instance, recent work has highlighted the brittleness of NLU systems to adversarial examples [Jia and Liang, 2017], as well as the fact that NLU models tend to exploit statistical artifacts in datasets, rather than exhibiting true reasoning and generalization capabilities [Guru-

rangan et al., 2018b, Kaushik and Lipton, 2018]. These findings have also dovetailed with the recent dominance of large pre-trained language models, such as BERT, on NLU benchmarks [Devlin et al., 2018, Peters et al., 2018b], which suggest that the primary difficulty in these datasets is incorporating the statistics of the natural language, rather than reasoning.

An important challenge is thus to develop NLU benchmarks that can precisely test a model’s capability for robust and systematic generalization. Ideally, we want language understanding systems that can not only answer questions and draw inferences from text, but that can also do so in a systematic, logical, and robust way. While such reasoning capabilities are certainly required for many existing NLU tasks, most datasets combine several challenges of language understanding into one, such as coreference/entity resolution, incorporating world knowledge, and semantic parsing—making it difficult to isolate and diagnose a model’s capabilities for systematic generalization and robustness.

In this work, we propose to use the properties of *systematicity* to test the limits of semantic generalization of modern neural networks. As defined by Fodor and Pylyshyn [1988], systematicity test the ability of a system to understand the recombination of known parts and rules. Thus, inspired by the classic AI challenge of inductive logic programming [Quinlan, 1990], in this chapter I discuss my work on developing semi-synthetic benchmark designed to explicitly test an NLU model’s ability for systematic and robust logical generalization [Sinha et al., 2019]. Our benchmark suite—termed **CLUTRR** (Compositional Language Understanding and Text-based Relational Reasoning)—contains a large set of semi-synthetic stories involving hypothetical families. Given a story, the goal is to infer the relationship between two family members, whose relationship is not explicitly mentioned. To solve this task, a learning agent must extract the relationships mentioned in the text, induce the logical rules governing the kinship relationships (e.g., the transitivity of the sibling relation), and use a

combination of these rules to infer the relationship between a given pair of entities. Crucially, the CLUTRR benchmark allows us to test a learning agent’s ability for *systematic generalization* by testing on stories that contain unseen combinations of logical rules. CLUTRR also allows us to precisely test for the various forms of *model robustness* by adding different kinds of superfluous *noise facts* to the stories.

3.1 Technical Background

3.1.1 Notations and Terminology

Following standard practice in formal semantics, we use the term *atom* to refer to a *predicate symbol* and a list of terms, such as $[\text{grandfatherOf}, X, Y]$, where the predicate `grandfatherOf` denotes the *relation* between the two *variables*, X and Y . We restrict the predicates to have an arity of 2, i.e., binary predicates. A logical *rule* in this setting is of the form $\mathcal{H} \vdash \mathcal{B}$, where \mathcal{B} is the *body* of the rule, i.e., a conjunction of two *atoms* ($[\alpha_1, \alpha_2]$) and \mathcal{H} is the *head*, i.e., a single atom (α) that can be viewed as the goal or query. For instance, given a knowledge base (KB) R that contains the single rule

$$[\text{grandfatherOf}, X, Y] \vdash [[\text{fatherOf}, X, Z], [\text{fatherOf}, Z, Y]], \quad (3.1)$$

the query $[\text{grandfatherOf}, X, Y]$ evaluates to true if and only if the body

$$\mathcal{B} = [[\text{fatherOf}, X, Z], [\text{fatherOf}, Z, Y]] \quad (3.2)$$

is also true in a given world. A rule is called a *grounded* rule if all atoms in the rule are themselves *grounded*, i.e., all variables are replaced with *constants* or entities in a world. A *fact* is a grounded binary predicate. A *clause* is a conjunction of two or more atoms ($\mathcal{C} = (\mathcal{H}_C \vdash \mathcal{B}_C = ([\alpha_1, \dots, \alpha_n]))$) which can be built using a set of rules.

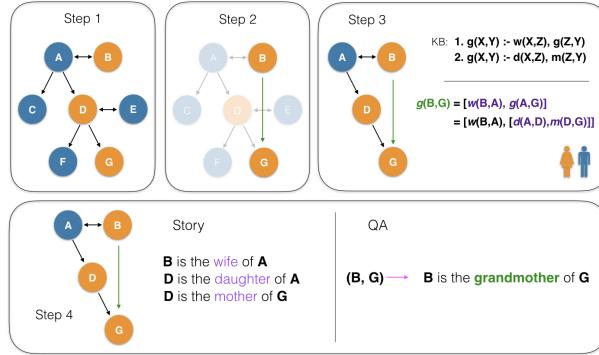


Figure 3.1 Data generation pipeline. Step 1: generate a kinship graph. Step 2: sample a target fact. Step 3: Use backward chaining to sample a set of facts. Step 4: Convert sampled facts to a natural language story.

3.2 Overview and construction of CLUTRR

The core idea behind the CLUTRR benchmark suite is the following: Given a natural language story describing a set of kinship relations, the goal is to infer the relationship between two entities, whose relationship is *not* explicitly stated in the story. To generate these stories, we first design a knowledge base (KB) with rules specifying how kinship relations resolve, and we use the following steps to create semi-synthetic stories based on this knowledge base:

- Step 1.** Generate a random kinship graph that satisfies the rules in our KB.
- Step 2.** Sample a target fact (i.e., relation) to predict from the kinship graph.
- Step 3.** Apply backward chaining to sample a set of facts that can prove the target relation (and optionally sample a set of “distracting” or “irrelevant” noise facts).
- Step 4.** Convert the sampled facts into a natural language story through pre-specified text templates and crowd-sourced paraphrasing.

Figure 3.1 provides a high-level overview of this idea, and the following subsections describe the data generation process in detail, as well as the diagnostic flexibility

afforded by CLUTRR.

The short stories in CLUTRR are essentially narrativized renderings of a set of logical facts. In the following sections, we describe how we sample the logical facts that make up a story by generating random kinship graphs and using backward chaining to produce logical reasoning chains.

3.2.1 Graph generation

To generate a kinship graph (say, G) underlying a particular story, we first sample a set of gendered¹ entities and kinship relations using a stochastic generation process. This generation process contains a number of tunable parameters—such as the maximum number of children at each node, the probability of an entity being married to another entity, etc.—and is designed to produce a valid, but possibly incomplete “backbone graph”. For instance, this backbone graph generation process will specify “parent”/“child” relations between entities but does not add “grandparent” relations. After this initial generation process, we recursively apply the logical rules in R to the backbone graph to produce a final graph G that contains the full set of kinship relations between all the entities.²

In the CLUTRR Benchmark, the following kinship relations are used: *son, father, husband, brother, grandson, grandfather, son-in-law, father-in-law, brother-in-law, uncle, nephew, daughter, mother, wife, sister, granddaughter, grandmother, daughter-in-law, mother-in-law, sister-in-law, aunt, niece*.

¹Kinship and gender roles are oversimplified in our data (compared to the real world) to maintain tractability.

²In the context of our data generation process, we distinguish between the knowledge base, R , which contains a finite number of predicates and rules specifying how kinship relations in a family resolve, and a particular kinship graph G , which contains a grounded set of atoms specifying the particular kinship relations that underlie a single story. In other words, R contains the logical rules that govern all the generated stories in CLUTRR, while G contains the grounded facts that underlie a specific story.

$[\text{grand}, X, Y] \vdash [[\text{child}, X, Z], [\text{child}, Z, Y]],$
 $[\text{grand}, X, Y] \vdash [[\text{so}, X, Z], [\text{grand}, Z, Y]],$
 $[\text{grand}, X, Y] \vdash [[\text{grand}, X, Z], [\text{sibling}, Z, Y]],$
 $[\text{inv-grand}, X, Y] \vdash [[\text{inv-child}, X, Z], [\text{inv-child}, Z, Y]],$
 $[\text{inv-grand}, X, Y] \vdash [[\text{sibling}, X, Z], [\text{inv-grand}, Z, Y]],$
 $[\text{child}, X, Y] \vdash [[\text{child}, X, Z], [\text{sibling}, Z, Y]],$
 $[\text{child}, X, Y] \vdash [[\text{so}, X, Z], [\text{child}, Z, Y]],$
 $[\text{inv-child}, X, Y] \vdash [[\text{sibling}, X, Z], [\text{inv-child}, Z, Y]],$
 $[\text{inv-child}, X, Y] \vdash [[\text{child}, X, Z], [\text{inv-grand}, Z, Y]],$
 $[\text{sibling}, X, Y] \vdash [[\text{child}, X, Z], [\text{inv-un}, Z, Y]],$
 $[\text{sibling}, X, Y] \vdash [[\text{inv-child}, X, Z], [\text{child}, Z, Y]]$
 $[\text{sibling}, X, Y] \vdash [[\text{sibling}, X, Z], [\text{sibling}, Z, Y]],$
 $[\text{in-law}, X, Y] \vdash [[\text{child}, X, Z], [\text{so}, Z, Y]],$
 $[\text{inv-in-law}, X, Y] \vdash [[\text{so}, X, Z], [\text{inv-child}, Z, Y]],$
 $[\text{un}, X, Y] \vdash [[\text{sibling}, X, Z], [\text{child}, Z, Y]],$
 $[\text{inv-un}, X, Y] \vdash [[\text{inv-child}, X, Z], [\text{sibling}, Z, Y]],$

We used a small, tractable, and logically sound KB of rules as mentioned above. We carefully select this set of deterministic rules to avoid ambiguity in the resolution. We use gender-neutral predicates and resolve the gender of the predicate in the head \mathcal{H} of a clause \mathcal{C} by deducing the gender of the second constant. We have two types of predicates, *vertical* predicates (parent-child relations) and *horizontal* predicates (sibling or significant other). We denote all the vertical predicates by its *child-to-parent* relation and append the prefix `inv-` to the predicates for the corresponding *parent-to-child* relation. For example, `grandfatherOf` is denoted by the gender-neutral predicate $[\text{inv-grand}, X, Y]$, where the gender is determined by the gender of Y .

3.2.2 Backward chaining

The resulting graph G provides the *background knowledge* for a specific story, as each edge in this graph can be treated as a grounded predicate (i.e., fact) between two entities. From this graph G , we sample the facts that make up the story, as well as the target fact that we seek to predict: First, we (uniformly) sample a target relation \mathcal{H}_c , which is the fact that we want to predict from the story. Then, from this target relation \mathcal{H}_c , we run a simple variation of the backward chaining [Gallaire and Minker, 1978] algorithm for k iterations starting from \mathcal{H}_c , where at each iteration we uniformly sample a subgoal to resolve and then uniformly sample a KB rule that resolves this subgoal. Crucially, unlike traditional backward chaining, we do not stop the algorithm when a proof is obtained; instead, we run for a fixed number of iterations k in order to sample a set of k facts \mathcal{B}_c that imply the target relation \mathcal{H}_c .

3.2.3 Adding natural language

So far, we have described the process of generating a conjunctive logical clause $\mathcal{C} = (\mathcal{H}_c \vdash \mathcal{B}_c)$, where $\mathcal{H}_c = [\alpha^*]$ is the target fact (i.e., relation) we seek to predict and $\mathcal{B}_c = [\alpha_1, \dots, \alpha_k]$ is the set of supporting facts that imply the target relation. We now describe how we convert this logical representation to natural language through crowdsourcing.

Paraphrasing using Amazon Mechanical Turk

We use Amazon Mechanical Turk (AMT), an online platform for collecting annotations from crowd-workers³. The platform supports a mechanism to quickly annotate large amounts of data by paying anonymous workers for their effort. In our work, the crowd-workers are shown a set of facts \mathcal{B}_c corresponding to a story and then they

³<https://www.mturk.com/>

are asked to paraphrase these facts into a narrative. Since workers are given a set of facts \mathcal{B}_C to work from, they are able to combine and split multiple facts across separate sentences and construct diverse narratives (Figure 3.3).

We use ParlAI [Miller et al., 2017] Mturk interface to collect paraphrases from the users. Specifically, given a set of facts, we ask the users to paraphrase the facts into a story. The users (*turkers*) are free to construct any story they like as long as they mention all the entities and all the relations among them. We also provide the head \mathcal{H} of the clause as an *inferred* relation and specifically instruct the users to *not* mention it in the paraphrased story. In order to evaluate the paraphrased stories, we ask the turkers to peer review a story paraphrased by a different turker. Since there are two tasks - paraphrasing a story and rating a story - we choose to pay 0.5\$ for each annotation. A sample task description in our MTurk interface is as follows:

In this task, you will need to write a short, simple story based on a few facts. **It is crucial that the story mentions each of the given facts at least once.** The story does not need to be complicated! It just needs to be grammatical and mention the required facts.

After writing the story, you will be asked to evaluate the quality of a generated story (based on a different set of facts). **It is crucial that you check whether the generated story mentions each of the required facts.**

Example of good and bad stories: Good Example

Facts to Mention

- John is the father of Sylvia.
- Sylvia has a brother Patrick.

Implied Fact: John is the father of Patrick.

Written story

John is the proud father of the lovely Sylvia. Sylvia has a love-hate relationship with her brother Patrick.

Bad Example

Facts to Mention

- Vincent is the son of Tim.
- Martha is the wife of Tim.

Implied Fact : Martha is Vincent's mother.

Written story

Vincent is married at Tim and his mother is Martha.

The reason the above story is bad:

- This story is bad because it is nonsense / ungrammatical.
- This story is bad because it does not mention the proper facts.
- This story is bad because it reveals the implied fact.

A sample of the AMT interface is shown in Figure 3.2. To ensure that the turkers are providing high-quality annotations without revealing the inferred fact, we also launch another task to ask the turkers to rate three annotations to be either good or bad which are provided by a set of *different* turkers. We pay 0.2\$ for each HIT consisting of three reviews. This helped to remove logical and grammatical inconsistencies to a large extent. Based on the reviews, 79% of the collected paraphrases passed the peer-review sanity check where all the reviewers agree on the quality. This subset of the placeholders is used in the benchmark. A sample of programmatically generated dataset for clause length of $k = 2$ to $k = 6$ is provided in the tables 3.3 to 3.7.

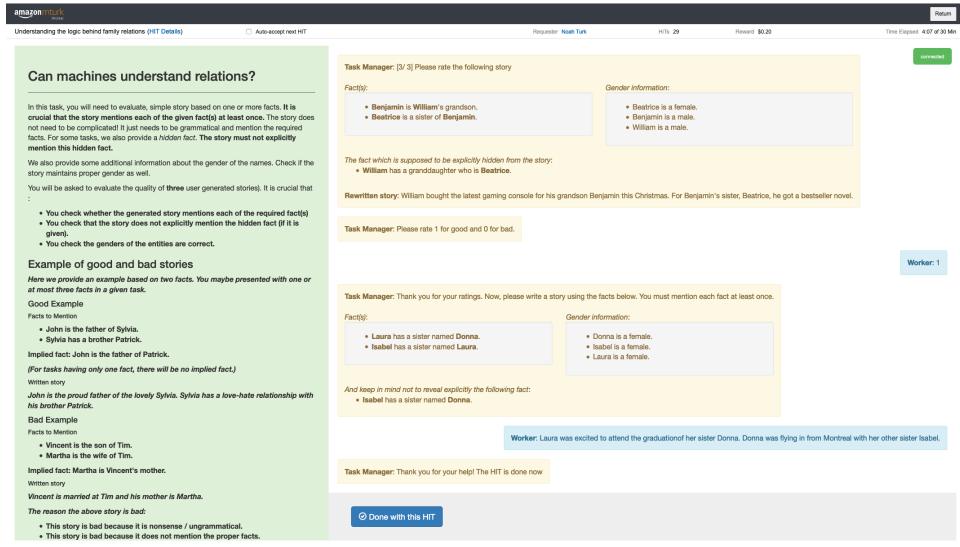


Figure 3.2 Amazon Mechanical Turk interface built using ParlAI which was used to collect data as well as peer reviews.

Reusability and composition

One challenge for data collection via AMT is that the number of possible stories generated by CLUTRR grows combinatorially as the number of supporting facts increases, i.e., as $k = |\mathcal{B}_C|$ grows. This combinatorial explosion for large k —combined with the difficulty of maintaining the quality of the crowd-sourced paraphrasing for long stories—makes it infeasible to obtain a large number of paraphrased examples for $k > 3$. To circumvent this issue and increase the flexibility of our benchmark, we reuse and compose AMT paraphrases to generate longer stories. In particular, we collected paraphrases for stories containing $k = 1, 2, 3$ supporting facts and then replaced the entities from these collected stories with placeholders in order to re-use them to generate longer semi-synthetic stories. An example of a story generated by stitching together two shorter paraphrases is provided below:

[Frank] went to the park with his father, [Brett]. [Frank] called his brother [Boyd] on the phone. He wanted to go out for some beers. [Boyd] went to the baseball game with his son [Jim].

Q: What is [Brett] and [Jim]'s relationship?

Thus, instead of simply collecting paraphrases for a fixed number of stories, we instead obtain a diverse collection of natural language templates that can be programmatically recombined to generate stories with various properties.

3.2.4 AMT Template statistics

Number of Paraphrases	# clauses	
$k = 1$	1,868	20
$k = 2$	1,890	58
$k = 3$	2,258	236
Total	6,016	
Unique Word Count	3,797	
Jaccard Word Overlap	Unigrams	0.201
	Bigrams	0.0385

Table 3.1 Statistics of the AMT paraphrases. Jaccard word overlap is calculated within the templates of each individual clause of length k .

At the time of submission, we have collected 6,016 unique paraphrases with an average of 19 paraphrases for every possible logical clause of length $k = 1, 2, 3$. Table 3.1 contains summary statistics of the collected paraphrases. Overall, we found high linguistic diversity in the collected paraphrases. For instance, the average Jaccard overlap in unigrams between pairs paraphrases corresponding to the same logical clause was only 0.201 and only 0.0385 for bigrams.

3.2.5 Human performance

To get a sense of the data quality and difficulty involved in CLUTRR, we asked human annotators to solve the task for random examples of length $k = 2, 3, \dots, 6$. (Table 3.2)

Relation Length	Human Performance		Reported Difficulty
	Time Limited	Unlimited Time	
2	0.848	1	1.488 +- 1.25
3	0.773	1	2.41 +- 1.33
4	0.477	1	3.81 +- 1.46
5	0.424	1	3.78 +- 0.96
6	0.406	1	4.46 +- 0.87

Table 3.2 Human performance accuracies on CLUTRR dataset. Humans are provided the Clean-Generalization version of the dataset, and we test on two scenarios: when a human is given limited time to solve the task, and when a human is given unlimited time to solve the task. Regardless of time, our evaluators provide a score of difficulty of individual puzzles.

We perform the evaluation in two scenarios: first a time-limited scenario where we ask AMT Turkers to solve the puzzle in a fixed time. Turkers were provided a maximum time of 30 mins, but they solved the puzzles in an average of 1 minute 23 seconds. Secondly, we use another set of expert evaluators who are given ample time to solve the tasks. Not surprisingly, if a human being is given ample time (experts took an average of 6 minutes per puzzle) and a pen and a paper to aid in the reasoning, they get all the relations correct. However, if an evaluator is short of time, they might miss important details on the relations and perform poorly. Thus, our tasks require *active attention*.

We found that time-constrained AMT annotators performed well (i.e., > 70%) accuracy for $k \leq 3$ but struggled with examples involving longer stories, achieving 40-50% accuracy for $k > 3$. However, trained annotators with unlimited time were able to solve 100% of the examples, highlighting the fact that this task requires attention and involved reasoning, even for humans.

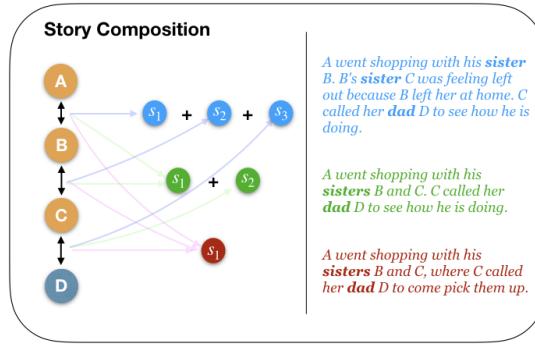


Figure 3.3 Illustration of how a set of facts can split and combine in various ways across sentences.

3.2.6 Representing the question and entities

The AMT paraphrasing approach described above allows us to convert the set of supporting facts \mathcal{B}_c to a natural language story, which can be used to predict the target relation/query \mathcal{H}_c . However, instead of converting the target query, $\mathcal{H}_c = [\alpha^*]$, to a natural language question, we instead opt to represent the target query as a K -way classification task, where the two entities in the target relation are provided as input and the goal is to classify the relation that holds between these two entities. This representation avoids the pitfall of revealing information about the answer in the question [Kaushik and Lipton, 2018].

When generating stories, entity names are randomly drawn from a set of 300 common gendered English names. Thus, depending on each run, the entities are never the same. This ensures that the entity names are simply placeholders and uncorrelated from the task.

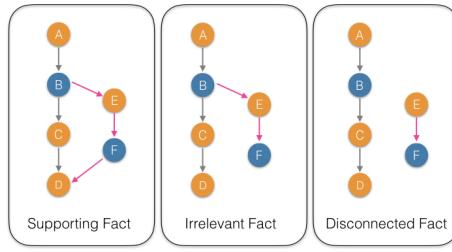


Figure 3.4 Noise generation procedures of CLUTRR.

3.3 Experimental Setups

The modular nature of CLUTRR provides rich diagnostic capabilities for evaluating the robustness and generalization abilities of neural language understanding systems. We highlight some key diagnostic capabilities available via different variations of CLUTRR below. These diagnostic variations correspond to the concrete datasets that we generated in this work, and we describe the results on these datasets in §3.5.

3.3.1 Systematic generalization

Most prominently, CLUTRR allows us to explicitly evaluate a model’s ability for generalizing with the property of systematicity. In particular, we rely on the following hold-out procedures to test systematic generalization:

- During training, we hold out a subset of the collected paraphrases, and we only use this held-out subset of paraphrases when generating the test set. Thus, to succeed on CLUTRR, an NLU system must exhibit *linguistic generalization* and be robust to linguistic variation at test time.
- We also hold out a subset of the logical clauses during training (for clauses of length $k > 2$).⁴ In other words, during training, the model sees all logical rules but does not

⁴One should not holdout clauses from length $k = 2$ in order to allow models to learn the compositionality of all possible binary predicates.

see all *combinations* of these logical rules. Thus, in addition to linguistic generalization, success on this task also requires *logical generalization*.

- Lastly, as a more extreme form of both logical and linguistic generalization, we consider the setting where the models are trained on stories generated from clauses of length $\leq k$ and evaluated on stories generated from larger clauses of length $> k$. Thus, we explicitly test the ability for models to generalize on examples that require more steps of reasoning than any example they encountered during training.

3.3.2 Robust Reasoning

In addition to evaluating systematic generalization, the modular setup of CLUTRR also allows us to diagnose model robustness by adding *noise facts* to the generated narratives. Due to the controlled semi-synthetic nature of CLUTRR, we are able to provide a precise taxonomy of the kinds of noise facts that can be added (Figure 3.4). In order to structure this taxonomy, it is important to recall that any set of supporting facts \mathcal{B}_c generated by CLUTRR can be interpreted as a path, p_c , in the corresponding kinship graph G (Figure 3.1). Based on this interpretation, we view adding noise facts from the perspective of sampling three different types of noise paths, p_n , from the kinship graph G :

- *Irrelevant facts*: We add a path p_n , which has exactly one shared end-point with p_c . In this way, this is a *distractor* path, which contains facts that are connected to one of the entities in the target relation, \mathcal{H}_c , but do not provide any information that could be used to help answer the query.
- *Supporting facts*: We add a path p_n , whose two end-points are on the path p_c . The facts on this path p_n are noise because they are not needed to answer the query, but they are supporting facts because they can, in principle, be used to construct alternative (longer) reasoning paths that connect the two target entities.
- *Disconnected facts*: We add paths which neither originate nor end in any entity on p_c .

These disconnected facts involve entities and relations that are completely unrelated to the target query.

3.3.3 Generated Datasets

For all experiments, we generated datasets with 10-15k training examples. In many experiments, we report training and testing results on stories with different clause lengths k . (For brevity, we use the phrase “clause length” throughout this section to refer to the value $k = |\mathcal{B}_c|$, i.e., the number of steps of reasoning that are required to predict the target query.) In all cases, the training set contains 5000 train stories per k value, and, during testing, all experiments use 100 test stories per k value. All experiments were run 10 times with different randomly generated stories, and means and standard errors over these 10 runs are reported. As discussed above, during training we holdout 20% of the paraphrases, as well as 10% of the possible logical clauses.

Table 3.3 Snapshot of puzzles in the dataset for $k=2$

Puzzle	Question	Gender	Answer
<i>Charles's son Christopher entered rehab for the ninth time at the age of thirty. Randolph had a nephew called Christopher who had n't seen for a number of years.</i>	Randolph is the _____ of Charles	Charles:male, Christopher:male, Randolph:male	
<i>Randolph and his sister Sharon went to the park. Arthur went to the baseball game with his son Randolph</i>	Sharon is the _____ of Arthur	Arthur:male, Randolph:male, Sharon:female	daughter
<i>Frank went to the park with his father, Brett. Frank called his brother Boyd on the phone. He wanted to go out for some beers.</i>	Brett is the _____ of Boyd	Boyd:male, Frank:male, Brett:male	father

3.4 Evaluated Models

Our primary baselines are neural language understanding models that take unstructured text as input. We consider bidirectional LSTMs [Hochreiter and Schmidhuber, 1997, Cho et al., 2014] (with and without attention), as well as models that aim to incorporate inductive biases towards relational reasoning: Relation Networks (RN) [Santoro

Table 3.4 Snapshot of puzzles in the dataset for k=3

Puzzle	Question	Gender	Answer
<i>Roger was playing baseball with his sons <i>Sam</i> and <i>Leon</i>. <i>Sam</i> had to take a break though because he needed to call his sister <i>Robin</i>.</i>	Leon is the _____ of Robin	Robin:female, Sam:male, Roger:male, Leon:male	brother
<i>Elvira and her daughter <i>Nancy</i> went shopping together last Monday and they bought new shoes for <i>Elvira's</i> kids. <i>Pedro</i> and his sister <i>Allison</i> went to the fair. <i>Pedro's</i> mother, <i>Nancy</i>, was out with friends for the day.</i>	Elvira is the _____ of Allison	Allison:female, Pedro:male, Nancy:female, Elvira:female	grandmother
<i>Roger met up with his sister <i>Nancy</i> and her daughter <i>Cynthia</i> at the mall to go shopping together. <i>Cynthia's</i> brother <i>Pedro</i> was going to be the star in the new show.</i>	Pedro is the _____ of Roger	Roger:male, Nancy:female, Cynthia:female, Pedro:male	nephew

Table 3.5 Snapshot of puzzles in the dataset for k=4

Puzzle	Question	Gender	Answer
<i>Celina has been visiting her sister, <i>Fran</i> all week. <i>Fran</i> is also the daughter of <i>Bethany</i>. <i>Ronald</i> loves visiting his aunt <i>Bethany</i> over the weekends. <i>Samuel's</i> son <i>Ronald</i> entered rehab for the ninth time at the age of thirty.</i>	Celina is the _____ of Samuel	Samuel:male, Ronald:male, Bethany:female, Fran:female, Celina:female	niece
<i>Celina adores her daughter <i>Bethany</i>. <i>Bethany</i> loves her very much, too. <i>Jackie</i> called her mother <i>Bethany</i> to let her know she will be back home soon. <i>Thomas</i> was helping his daughter <i>Fran</i> with her homework at home. Afterwards, <i>Fran</i> and her sister <i>Jackie</i> played Xbox together.</i>	Celina is the _____ of Thomas	Thomas:male, Fran:female, Jackie:female, Bethany:female, Celina:female	daughter
<i>Raquel is <i>Samuel</i>'s daughter and they go shopping at least twice a week together. <i>Kenneth</i> and her mom, <i>Theresa</i>, had a big fight. <i>Theresa's</i> son, <i>Ronald</i>, refused to get involved. <i>Ronald</i> was having an argument with her sister, <i>Raquel</i>.</i>	Samuel is the _____ of Kenneth	Kenneth:male, Theresa:female, Ronald:male, Raquel:female, Samuel:male	father

Table 3.6 Snapshot of puzzles in the dataset for k=5

Puzzle	Question	Gender	Answer
<i>Steven's son is Bradford. Bradford and his father always go fishing together on Sundays and have a great time together. Diane is taking her brother Brad out for a late dinner. Kristin, Brad's mother, is home with a cold. Diane's father Elmer, and his brother Steven, all got into the rental car to start the long cross-country roadtrip they had been planning.</i>	Bradford is the _____ of Kristin	Kristin:female, Brad:male, Diane:female, Elmer:male, Steven:male, Bradford:male	nephew
<i>Elmer went on a roadtrip with his youngest child, Brad. Lena and her sister Diane are going to a restaurant for lunch. Lena's brother Brad is going to meet them there with his father Elmer. Brad can't stand his unfriendly aunt Lizzie.</i>	Lizzie is the _____ of Diane	Diane:female, Lena:female, Brad:male, Elmer:male, Lizzie:female	aunt
<i>Ira took his niece April fishing Saturday. They caught a couple small fish. Ronald was enjoying spending time with his parents, Damion and Claudine. Damion's other son, Dennis, wanted to come visit too. Dennis often goes out for lunch with his sister, April.</i>	Ira is the _____ of Claudine	Claudine:female, Ronald:male, Damion:male, Dennis:male, April:female, Ira:male	brother

Table 3.7 Snapshot of puzzles in the dataset for k=6

Puzzle	Question	Gender	Answer
<i>Mario wanted to get a good gift for his sister, Marianne. Jean and her sister Darlene were going to a party held by Jean's mom, Marianne. Darlene invited her brother Roy to come, too, but he was too busy. Teri and her father, Mario, had an argument over the weekend. However, they made up by Monday. Agnes wants to make a special meal for her daughter Teri's birthday.</i>	Roy is the _____ of Agnes	Agnes:female, Teri:female, Mario:male,	Marianne:female, nephew Jean:female, Darlene:female, Roy:male
<i>Robert's aunt, Marianne, asked Robert to mow the lawn for her. Robert said he could n't because he had a bad back. William's parents, Brian and Marianne, threw him a surprise party for his birthday. Brian's daughter Jean made a mental note to be out of town for her birthday! Agnes's biggest accomplishment is raising her son Robert. Jean is looking for a good gift for her sister Darlene.</i>	Darlene is the _____ of Agnes	Agnes:female, Robert:male, Marianne:female,	William:male, niece Brian:male, Jean:female, Darlene:female
<i>Sharon and her brother Mario went shopping. Teri, Mario's daughter, came too. Agnes, Annie's mother, is unhappy with Robert. She feels her son is cruel to Annie's sister Teri, and she wants Robert to be nicer. Robert's sister, Nicole, participated in the dance contest.</i>	Nicole is the _____ of Sharon	Sharon:female, Mario:male, Teri:female,	Annie:female, niece Agnes:female, Robert:male, Nicole:female

et al., 2017], Relational Recurrent Networks (RMC) [Santoro et al., 2018] and Compositional Memory Attention Network (MAC) [Hudson and Manning, 2018]. We also use the large pre-trained language model, BERT [Devlin et al., 2018], as well as a modified version of BERT having a trainable LSTM encoder on top of the pretrained BERT embeddings. All of these models (except BERT) were re-implemented in PyTorch 1.0 [Paszke et al., 2017] and adapted to work with the CLUTRR benchmark.

Since the underlying relations in the stories generated by CLUTRR inherently form a graph, we also experiment with a Graph Attention Network (GAT) [Veličković et al., 2018]. Rather than taking the textual stories as input, the GAT baseline receives a structured graph representation of the facts that underlie the story.

Entity and query representations. We use the various baseline models to encode the natural language story (or graph) into a fixed-dimensional embedding. With the exception of the BERT models, we do not use pre-trained word embeddings and learn the word embeddings from scratch using end-to-end backpropagation. An important note, however, is that we perform Cloze-style anonymization [Hermann et al., 2015] of the entities (i.e., names) in the stories, where each entity name is replaced by a `@entity-k` placeholder, which is randomly sampled from a small, fixed pool of placeholder tokens. The embeddings for these placeholders are randomly initialized and fixed during training.

To make a prediction about a target query given a story, we concatenate the embedding of the story (generated by the baseline model) with the embeddings of the two target entities and we feed this concatenated embedding to a 2-layer feed-forward neural network with a softmax prediction layer.

3.4.1 Hyperparameters

For all models, the common hyperparameters used are: Embedding dimension: 100 (except BERT based models), Optimizer: Adam, Learning rate: 0.001, Number of

epochs: 100, Number of runs: 10. Specific model-based hyperparameters are given as follows:

- **Bidirectional LSTM** [Hochreiter and Schmidhuber, 1997, Cho et al., 2014]: LSTM hidden dimension: 100, # layers: 2, Classifier MLP hidden dimension: 200
- **Relation Networks** [Santoro et al., 2017]: $f_{\theta_1} : 256$, $f_{\theta_2} : 64$, $g_{\theta} : 64$
- **Compositional Memory Attention Network (MAC)** [Hudson and Manning, 2018]: # Iterations: 6, shareQuestion: True, Dropout - Memory, Read and Write: 0.2
- **Relational Recurrent Networks** [Santoro et al., 2018]: Memory slots: 2, Head size: 192, Number of heads: 4, Number of blocks : 1, forget bias : 1, input bias: 0, gate style: unit, key size: 64, # Attention layers: 3, Dropout: 0
- **BERT** [Devlin et al., 2018]: Layers : 12, Fixed pretrained embeddings from bert-base-uncased using Pytorch HuggingFace BERT repository⁵, Word dimension: 768, appended with a two-layer MLP for final prediction.
- **BERT-LSTM**: Same parameters as above, with a two-layer unidirectional LSTM encoder on top of BERT word embeddings.
- **GAT** [Veličković et al., 2018]: Node dimension: 100, Message dimension: 100, Edge dimension: 20, number of rounds: 3

3.5 Results

We evaluate several NLU systems on the proposed CLUTRR benchmark to surface the relative strengths and shortcomings of these models in the context of inductive reasoning and combinatorial generalization.⁶ We aim to answer the following key questions:

⁵<https://github.com/huggingface/pytorch-pretrained-BERT>

⁶Code to reproduce all the results in this section are available at <https://github.com/facebookresearch/clutrr/>.

- (Q1) How do state-of-the-art NLU models compare in terms of systematic generalization? Can these models generalize to stories with unseen combinations of logical rules?
- (Q2) How does the performance of neural language understanding models compare to a graph neural network that has full access to graph structure underlying the stories?
- (Q3) How robust are these models to the addition of noise facts to a given story?

3.5.1 Systematic Generalization

We begin by using CLUTRR to evaluate the ability of the baseline models to perform systematic generalization (Q1). In this setting, we consider two training regimes: in the first regime, we train all models with clauses of length $k = 2, 3$, and in the second regime, we train with clauses of length $k = 2, 3, 4$. We then test the generalization of these models on test clauses of length $k = 2, \dots, 10$.

Figure 3.5 illustrates the performance of different models on this generalization task. We observe that the GAT model is able to perform near-perfectly on the held-out logical clauses of length $k = 3$, with the BERT-LSTM being the top-performer among the text-based models but still significantly below the GAT. Not surprisingly, the performance of all models degrades monotonically as we increase the length of the test clauses, which highlights the challenge of “zero-shot” systematic generalization Lake and Baroni [2018], Sodhani et al. [2018]. However, as expected, all models improve on their generalization performance when trained on $k = 2, 3, 4$ rather than just $k = 2, 3$ (Figure 3.5, right). The GAT, in particular, achieves the biggest gain by this expanded training.

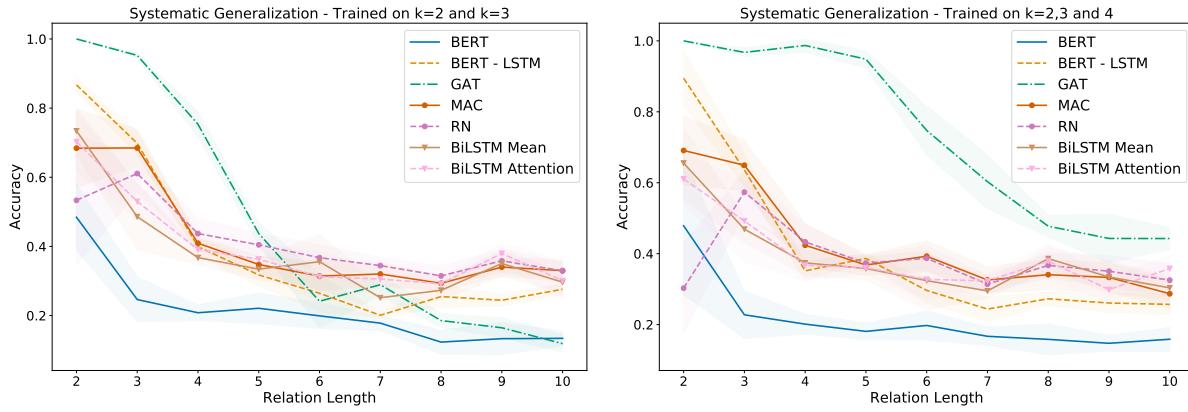


Figure 3.5 Systematic generalization performance of different models when trained on clauses of length $k = 2, 3$ (Left) and $k = 2, 3, 4$ (Right).

3.5.2 The benefit of structure

The empirical results on systematic generalization also provide insight into how the text-based NLU systems compare against the graph-based GAT model that has full access to the logical graph structure underlying the stories (**Q2**). Indeed, the relatively strong performance of the GAT model (Figure 3.5) suggests that the language-based models fail to learn a robust mapping from the natural language narratives to the underlying logical facts.

To further confirm this trend, we ran experiments with modified train and test splits for the text-based models, where the same set of natural language paraphrases were used to construct the narratives in both the train and test splits (Figure 3.6). In this simplified setting, the text-based models must still learn to reason about held-out logical patterns, but the difficulty of parsing the natural language is essentially removed, as the same natural language paraphrases are used during testing and training. We found that the text-based models were competitive with the GAT model in this simplified setting, confirming that the poor performance of the text-based models on the main task is driven by the difficulty of parsing the unseen natural language narratives.

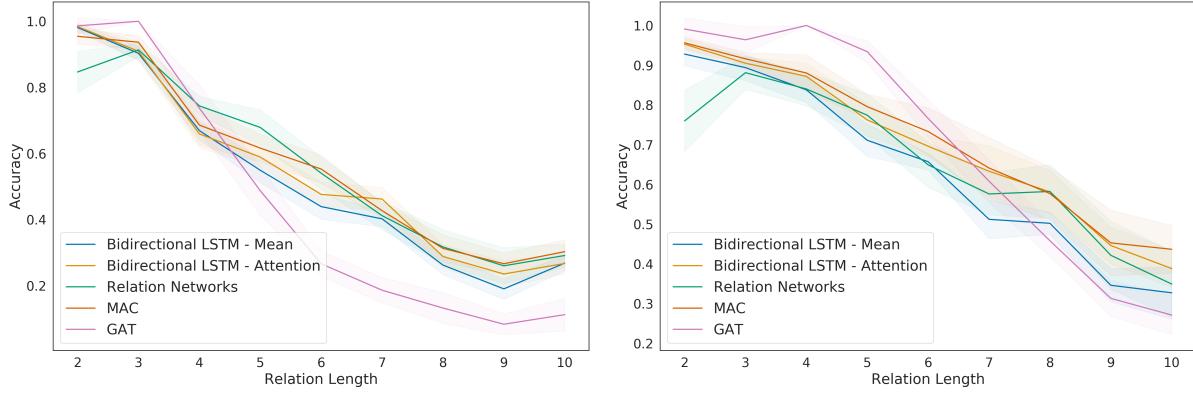


Figure 3.6 Systematic Generalizability of different models on CLUTRR-Gen task (having 20% less placeholders and without training and testing placeholder split), when **Left:** trained with $k = 2$ and $k = 3$ and **Right:** trained with $k = 2, 3$ and 4

Models		Unstructured models (no graph)						Structured model (with graph)
Training	Testing	BiLSTM - Attention	BiLSTM - Mean	RN	MAC	BERT	BERT-LSTM	GAT
Clean	Clean	0.58 ± 0.05	0.53 ± 0.05	0.49 ± 0.06	0.63 ± 0.08	0.37 ± 0.06	0.67 ± 0.03	1.0 ± 0.0
	Supporting	0.76 ± 0.02	0.64 ± 0.22	0.58 ± 0.06	0.71 ± 0.07	0.28 ± 0.1	0.66 ± 0.06	0.24 ± 0.2
	Irrelevant	0.7 ± 0.15	0.76 ± 0.02	0.59 ± 0.06	0.69 ± 0.05	0.24 ± 0.08	0.55 ± 0.03	0.51 ± 0.15
	Disconnected	0.49 ± 0.05	0.45 ± 0.05	0.5 ± 0.06	0.59 ± 0.05	0.24 ± 0.08	0.5 ± 0.06	0.8 ± 0.17
Supporting	Supporting	0.67 ± 0.06	0.66 ± 0.07	0.68 ± 0.05	0.65 ± 0.04	0.32 ± 0.09	0.57 ± 0.04	0.98 ± 0.01
Irrelevant	Irrelevant	0.51 ± 0.06	0.52 ± 0.06	0.5 ± 0.04	0.56 ± 0.04	0.25 ± 0.06	0.53 ± 0.06	0.93 ± 0.01
Disconnected	Disconnected	0.57 ± 0.07	0.57 ± 0.06	0.45 ± 0.11	0.4 ± 0.1	0.17 ± 0.05	0.47 ± 0.06	0.96 ± 0.01
Average		0.61 ± 0.08	0.59 ± 0.08	0.54 ± 0.07	0.61 ± 0.06	0.30 ± 0.07	0.56 ± 0.05	0.77 ± 0.09

Table 3.8 Testing the robustness of the various models when training and testing on stories containing various types of noise facts. The types of noise facts (supporting, irrelevant, and disconnected) are defined in Section .

3.5.3 Robust Reasoning

Finally, we use CLUTRR to systematically evaluate how various baseline neural language understanding systems cope with noise (**Q3**). In all the experiments we provide a combination of $k = 2$ and $k = 3$ length clauses in training and testing, with noise facts being added to the train and/or test set depending on the setting (Table 3.8). We use the different types of noise facts defined in Section 3.3.2..

Overall, we find that the GAT baseline outperforms the unstructured text-based

models across most testing scenarios (Table 3.8), which showcases the benefit of a structured feature space for robust reasoning. When training on clean data and testing on noisy data, we observe two interesting trends that highlight the benefits and shortcomings of the various model classes:

1. All the text-based models excluding BERT actually perform better when testing on examples that have *supporting* or *irrelevant* facts added. This suggests that these models actually benefit from having more content related to the entities in the story. Even though this content is not strictly useful or needed for the reasoning task, it may provide some linguistic cues (e.g., about entity genders) that the models exploit. In contrast, the BERT-based models do not benefit from the inclusion of this extra content, which is perhaps due to the fact that they are already built upon a strong language model (e.g., that already adequately captures entity genders.)
2. The GAT model performs poorly when *supporting* facts are added but has no performance drop when *disconnected* facts are added. This suggests that the GAT model is sensitive to changes that introduce cycles in the underlying graph structure but is robust to the addition of noise that is disconnected from the target entities.

Learning from noisy data

Moreover, when we trained on noisy examples, we found that only the GAT model was able to consistently improve its performance (Table 3.8). We notice that the GAT model, having access to the true underlying graph of the puzzles, perform better across different testing scenarios when trained with the noisy data. As the *Supporting facts* contains cycles, it is difficult for GAT to generalize for a dataset with cycles when it is trained on a dataset without cycles. However, when trained with cycles, GAT learns to attend to *all* the paths leading to the correct answer. This effect is disastrous when GAT is tested on *Irrelevant facts* which contains dangling paths as GAT still tries to attend to all the paths. Training on *Irrelevant facts* proved to be most beneficial to GAT, as the

Models		Unstructured models (no graph)					Structured model (with graph)	
Training	Testing	BiLSTM - Attention	BiLSTM - Mean	RN	MAC	BERT	BERT-LSTM	GAT
Supporting	Clean	0.38 \pm 0.04	0.32 \pm 0.04	0.4 \pm 0.09	0.45 \pm 0.03	0.19 \pm 0.06	0.39 \pm 0.06	0.92 \pm 0.17
	Supporting	0.67 \pm 0.06	0.66 \pm 0.07	0.68 \pm 0.05	0.65 \pm 0.04	0.32 \pm 0.09	0.57 \pm 0.04	0.98 \pm 0.01
	Irrelevant	0.44 \pm 0.03	0.39 \pm 0.03	0.51 \pm 0.08	0.46 \pm 0.09	0.2 \pm 0.06	0.36 \pm 0.05	0.5 \pm 0.23
	Disconnected	0.31 \pm 0.21	0.25 \pm 0.16	0.47 \pm 0.08	0.41 \pm 0.06	0.2 \pm 0.08	0.32 \pm 0.04	0.92 \pm 0.05
Irrelevant	Clean	0.57 \pm 0.05	0.56 \pm 0.05	0.46 \pm 0.13	0.67 \pm 0.05	0.24 \pm 0.06	0.46 \pm 0.08	0.92 \pm 0.0
	Supporting	0.38 \pm 0.22	0.31 \pm 0.16	0.61 \pm 0.07	0.61 \pm 0.04	0.27 \pm 0.06	0.46 \pm 0.04	0.77 \pm 0.12
	Irrelevant	0.51 \pm 0.06	0.52 \pm 0.06	0.5 \pm 0.04	0.56 \pm 0.04	0.25 \pm 0.06	0.53 \pm 0.06	0.93 \pm 0.01
	Disconnected	0.44 \pm 0.26	0.54 \pm 0.27	0.55 \pm 0.05	0.61 \pm 0.06	0.26 \pm 0.03	0.45 \pm 0.08	0.85 \pm 0.25
Disconnected	Clean	0.45 \pm 0.02	0.47 \pm 0.03	0.53 \pm 0.09	0.5 \pm 0.06	0.22 \pm 0.09	0.44 \pm 0.05	0.75 \pm 0.07
	Supporting	0.47 \pm 0.03	0.46 \pm 0.05	0.54 \pm 0.03	0.58 \pm 0.06	0.22 \pm 0.06	0.38 \pm 0.08	0.78 \pm 0.12
	Irrelevant	0.47 \pm 0.05	0.48 \pm 0.03	0.52 \pm 0.04	0.51 \pm 0.05	0.17 \pm 0.04	0.38 \pm 0.05	0.56 \pm 0.26
	Disconnected	0.57 \pm 0.07	0.57 \pm 0.06	0.45 \pm 0.11	0.4 \pm 0.1	0.17 \pm 0.05	0.47 \pm 0.06	0.96 \pm 0.01
Average		0.47 \pm 0.08	0.46 \pm 0.08	0.52 \pm 0.07	0.53 \pm 0.06	0.23 \pm 0.07	0.43 \pm 0.05	0.82 \pm 0.09

Table 3.9 Testing the robustness of the various models when trained various types of noisy facts and evaluated on other noisy / clean facts.

model now perfectly attends to *only relevant paths*. Since *Disconnected facts* contains disconnected paths, the message passing function of the graph is unable to forward any information from the disjoint cliques, thereby having superior testing scores throughout several scenarios.

Again, these results highlights the performance gap between the unstructured text-based models and GAT for solving the CLUTRR task.

Learning with synthetic placeholders

In order to further understand the effect of language placeholders on robustness, we performed another set of experiments where we use bABI Weston et al. [2015] style simple placeholders (Table 3.10). We observe a marked increase in performance of all NLU models, where they significantly decrease the gap between their performance with that of GAT, even outperforming GAT on various settings. This shows the significance of using paraphrased placeholders in devising the complexity of the dataset.

Models		Unstructured models (no graph)					Structured model (with graph)	
Training	Testing	BiLSTM - Attention	BiLSTM - Mean	RN	MAC	BERT	BERT-LSTM	GAT
Supporting	Clean	0.96 \pm 0.01	0.97 \pm 0.01	0.88 \pm 0.05	0.94 \pm 0.02	0.48 \pm 0.08	0.57 \pm 0.08	0.92 \pm 0.17
	Supporting	0.96 \pm 0.03	0.96 \pm 0.03	0.97 \pm 0.01	0.97 \pm 0.01	0.75 \pm 0.07	0.88 \pm 0.05	0.98 \pm 0.01
	Irrelevant	0.92 \pm 0.02	0.93 \pm 0.01	0.9 \pm 0.03	0.91 \pm 0.01	0.56 \pm 0.04	0.54 \pm 0.06	0.5 \pm 0.23
	Disconnected	0.8 \pm 0.04	0.83 \pm 0.04	0.76 \pm 0.08	0.86 \pm 0.04	0.27 \pm 0.06	0.42 \pm 0.08	0.92 \pm 0.05
Irrelevant	Clean	0.63 \pm 0.02	0.61 \pm 0.07	0.85 \pm 0.09	0.8 \pm 0.07	0.53 \pm 0.09	0.44 \pm 0.06	0.92 \pm 0.0
	Supporting	0.66 \pm 0.03	0.64 \pm 0.04	0.69 \pm 0.06	0.76 \pm 0.06	0.42 \pm 0.08	0.43 \pm 0.08	0.77 \pm 0.12
	Irrelevant	0.89 \pm 0.04	0.86 \pm 0.1	0.74 \pm 0.11	0.78 \pm 0.06	0.61 \pm 0.1	0.83 \pm 0.06	0.93 \pm 0.01
	Disconnected	0.64 \pm 0.02	0.62 \pm 0.05	0.72 \pm 0.05	0.73 \pm 0.04	0.41 \pm 0.04	0.61 \pm 0.05	0.85 \pm 0.25
Disconnected	Clean	0.9 \pm 0.05	0.82 \pm 0.12	0.94 \pm 0.02	0.93 \pm 0.04	0.68 \pm 0.07	0.64 \pm 0.02	0.75 \pm 0.07
	Supporting	0.87 \pm 0.04	0.82 \pm 0.05	0.85 \pm 0.03	0.88 \pm 0.04	0.54 \pm 0.08	0.5 \pm 0.05	0.78 \pm 0.12
	Irrelevant	0.87 \pm 0.03	0.85 \pm 0.03	0.83 \pm 0.03	0.87 \pm 0.02	0.59 \pm 0.09	0.58 \pm 0.05	0.56 \pm 0.26
	Disconnected	0.91 \pm 0.04	0.91 \pm 0.03	0.8 \pm 0.17	0.71 \pm 0.11	0.49 \pm 0.1	0.79 \pm 0.1	0.96 \pm 0.01
Average		0.83 \pm 0.08	0.82 \pm 0.08	0.83 \pm 0.07	0.84 \pm 0.06	0.58 \pm 0.07	0.60 \pm 0.05	0.82 \pm 0.09

Table 3.10 Testing the robustness on toy placeholders of the various models when trained various types of noisy facts and evaluated on other noisy / clean facts.

3.6 Related Work

To design the CLUTRR dataset, we draw inspiration from the classic work on inductive logic programming (ILP), a long line of reading comprehension benchmarks in NLP, as well as work combining language and knowledge graphs.

3.6.1 Reading comprehension benchmarks

Many datasets have been proposed to test the reading comprehension ability of NLP systems. This includes the SQuAD Rajpurkar et al. [2016a], NewsQA Trischler et al. [2016], and MCTest Richardson et al. [2013] benchmarks that focus on factual questions; the SNLI Bowman et al. [2015b] and MultiNLI Williams et al. [2018d] benchmarks for sentence understanding; and the bABI tasks Weston et al. [2015], to name a few. Our primary contribution to this line of work is the development of a carefully designed *diagnostic* benchmark to evaluate model robustness and systematic generalization in the context of NLU.

3.6.2 Systematic generalization

A growing body of literature has demonstrated that NLU models tend to exploit statistical artifacts in datasets and lack true generalization capabilities Jia and Liang [2017], Gururangan et al. [2018b], Kaushik and Lipton [2018], Lake and Baroni [2018]. These critical examinations have dovetailed with similar studies on visual question answering [Agrawal et al., 2016, Bahdanau et al., 2019, Johnson et al., 2017]. CLUTRR, contributes to this growing area by introducing a principled and flexible benchmark to evaluate systematic generalization in the context of language understanding—with our notion of systematic generalization being grounded in classic work on inductive logic programming (ILP) Quinlan [1990].

3.6.3 Question-answering with knowledge graphs

Our work is also related to the domain of question answering and reasoning in knowledge graphs [Das et al., 2018, Xiong et al., 2018, Hamilton et al., 2018, Wang et al., 2018, Xiong et al., 2017, Welbl et al., 2018, Kartsaklis et al., 2018], where either the model is provided with a knowledge graph to perform inference over or where the model must infer a knowledge graph from the text itself. However, unlike previous benchmarks in this domain—which are generally *transductive* and focus on leveraging and extracting knowledge graphs as a source of background knowledge about a fixed set of entities—CLUTRR requires *inductive logical reasoning*, where every example requires reasoning over a new set of previously unseen entities.

3.7 Discussion

In this paper we introduced the CLUTRR benchmark suite to test the systematic generalization and inductive reasoning capabilities of NLU systems. We demonstrated the diagnostic capabilities of CLUTRR and found that existing NLU systems exhibit rel-

atively poor robustness and systematic generalization capabilities—especially when compared to a graph neural network that works directly with symbolic input. Concretely, using CLUTRR we were able to make the following key insights about the reasoning capability of modern neural networks:

- **Neural language models are unable to reason when tested with systematicity.**

We saw in §3.5.1 that the performance of all NLU models drastically degrade when we test on instances which require systematicity - the knowledge of combination of existing parts - to solve the task. While all models had access to all possible rules (by ingesting a combination of relations in the training data), all models are notably worse when tested with longer chain of reasoning than the ones trained upon. This shortcoming could be due to overly associating to certain patterns seen during training, or learning to solve the task by taking shortcuts - associating some combination of tokens for certain relations [Gururangan et al., 2018b].

- **Models are not robust in their language understanding.** When evaluated with enabling (supporting) and distractor information (noise), we observe models to display conflicting results. While supporting information is indeed useful for certain classes of models (§3.5.3), irrelevant and distracting information also seems to aide in the reasoning process, which is not a systematic behaviour. Furthermore, when trained with noise, majority of the NLU models are unable to discern between the correct and the incorrect information. These results indicate a potential surface form realization issue.

- **The key hurdle behind systematic generalization is the natural language itself.**

Finally, we observe overwhelmingly that when a model which is only provided a graph, stripped of the natural language layer, the model is able to reason with surprising ability. The graph model, GAT, does not have to extract the relevant

information from a given free-form text. This makes it easier for the model to generalize more effectively, even in the scenarios when the model is tasked to learn from distractor (noisy) information.

These results highlight the gap that remains between machine reasoning models that work with unstructured text and models that are given access to more structured input. It appears the key hindrance for a neural model for effective generalization and reasoning is the access to proper surface forms. These results raises questions on the syntax processing capabilities of NLU models, and call for more in-depth investigation on the same. In fact, in the following chapters of this thesis, I will discuss my works on further studying the notions of syntax encoding in NLU models using the tool of systematicity.

3.8 Follow-up findings in the community

Chapter 4

Quantifying syntactic generalization using word order

Of late, large scale pre-trained Transformer-based [Vaswani et al., 2017] models—such as RoBERTa [Liu et al., 2019a], BART [Lewis et al., 2020], and GPT-2 and -3 [Radford et al., 2019, Brown et al., 2020]—have exceeded recurrent neural networks’ performance on many NLU tasks [Wang et al., 2018, 2019]. Several papers have even suggested that Transformers pretrained on a language modeling (LM) objective can capture syntactic information [Hewitt and Manning, 2019a, Jawahar et al., 2019a, Warstadt and Bowman, 2020, Wu et al., 2020], with their self-attention layers being capable of surprisingly effective learning Rogers et al. [2020b]. In the preceding chapter, we observed that NLU models, including BERT, are unable to reason systematicity, primarily due to their lack of understanding the surface forms of the given task. Thus, in this chapter, we question the claim that state-of-the-art NLU models “know syntax”.

Since there are many ways to investigate “syntax”, we must be clear on what we mean by the term. Knowing the syntax of a sentence means being sensitive to the *order of the words* in that sentence (among other things). Humans are sensitive to word order, so clearly, “language is not merely a bag of words” [Harris, 1954a, p.156]. More-

over, it is easier for us to identify or recall words presented in canonical orders than in disordered, ungrammatical sentences; this phenomenon is called the “*sentence superiority effect*” (Cattell 1886, Scheerer 1981, Toyota 2001, Baddeley et al. 2009, Snell and Grainger 2017, 2019, Wen et al. 2019, i.a.). This effect also finds some neurobiological support from work showing ordered text activates portions of the temporal lobe more than unordered word lists [Bemis and Pylkkänen, 2013, Pylkkänen et al., 2014]. In our estimation then, if one wants to claim that a model “knows syntax”, then they should minimally show that the model is sensitive to word order (at least for e.g. English or Mandarin Chinese).

Generally, knowing the syntax of a sentence is taken to be a prerequisite for understanding what that sentence means [Heim and Kratzer, 1998]. Models should have to know the syntax first then, if performing any particular NLU task that genuinely requires a humanlike understanding of meaning (cf. Bender and Koller 2020). Thus, if our models are as good at NLU as our current evaluation methods suggest, we should expect them to be sensitive to word order. In this chapter, I discuss our paper Sinha et al. [2021c] where we use a suite of permutation metrics to find the models are not sensitive to word order.

We focus here on textual entailment, one of the hallmark tasks used to measure how well models understand language [Condoravdi et al., 2003, Dagan et al., 2005a]. This task, often also called Natural Language Inference (NLI; Bowman et al. 2015a, i.a.), typically consists of two sentences: a premise and a hypothesis. The objective is to predict whether the premise entails the hypothesis, contradicts it, or is neutral with respect to it. We find rampant word order insensitivity in purportedly high performing NLI models. For nearly all premise-hypothesis pairs, **there are many permuted examples that fool the models** into providing the correct prediction. In case of MNLI, for example, the current state-of-the-art of 90.5% can be increased to 98.7% merely by permuting the word order of test set examples. We even find drastically increased cross-dataset

Premise	Hypothesis	Predicted Label
Boats in daily use lie within feet of the fashionable bars and restaurants.	There are boats close to bars and restaurants.	E
restaurants and use feet of fashionable lie the in Boats within bars daily .	bars restaurants are There and to close boats .	E
He and his associates weren't operating at the level of metaphor.	He and his associates were operating at the level of the metaphor.	C
his at and metaphor the of were He operating associates n't level .	his the and metaphor level the were He at associates operating of .	C

Table 4.1 Examples from the MNLI Matched development set. Both the original example and the permuted one elicit the same classification label (entailment and contradiction respectively) from RoBERTa (large). A simple demo is provided in an associated Google Colab notebook.

generalization when we reorder words. This is not just a matter of chance—we show that the model output probabilities are significantly different from uniform. A sample of the model outputs with permuted examples is shown in Table 4.1.

We verify our findings with three popular English NLI datasets—SNLI [Bowman et al., 2015a], MultiNLI [Williams et al., 2018c] and ANLI [Nie et al., 2020])—and one Chinese one, OCNLI Hu et al. [2020a]. It is thus less likely that our findings result from some quirk of English or a particular tokenization strategy. We also observe the effect for various transformer architectures pre-trained on language modeling (RoBERTa [Liu et al., 2019a], BART [Lewis et al., 2020], DistilBERT [Sanh et al., 2020]), and non-transformers, including a ConvNet [Zhao et al., 2015], an InferSent model [Conneau

et al., 2017], and a BiLSTM [Collobert and Weston, 2008].

Thus, in this chapter I discuss our contributions in Sinha et al. [2021c], which are as follows: (i) we propose a suite of metrics (*Permutation Acceptance*) for measuring model insensitivity to word order (§4.2), (ii) we construct multiple permuted test datasets for measuring NLI model performance at a large scale (§4.4), (iii) we show that NLI models focus on words more than word order, but can partially reconstruct syntactic information from words alone (§4.5.1), (iv) we show the problem persists on out-of-domain data, (v) we show that humans struggle with UnNatural Language Inference, underscoring the non-humanlikeness of SOTA models (§4.5.2), (vi) finally, we explore a simple maximum entropy-based method (§4.5.3) to encourage models not to accept permuted examples.

4.1 Technical Background

In this chapter, we investigate the task of Natural Language Inference (NLI), also known as Recognizing Textual Entailment (RTE). NLI task consists of inferring the logical relation between two sentences, typically known as the *premise* and the *hypothesis*. The logical relations that can exist between these two sentences can be one of three types: *entailment* if the premise *entails* the hypothesis, *contradiction* if it is the opposite, and *neutral* if the sentences have non overlapping meaning. Historically, this logical relation formulation is derived from *natural logic* [MacCartney and Manning, 2007], which consisted of seven set-theoretic relations between any given pair of sentences. We use the formulation prescribed by Bowman et al. [2015a], which is the simplified formulation consisting of the three standard relations.

Linguists generally take syntactic structure to be necessary for humans to know what sentences mean. Many also find the NLI task to a very promising approximation of human natural language understanding, in part because it is rooted in the tradition

of logical entailment. In the spirit of propositional logic, sentence meaning is taken to be truth-conditional [Frege, 1948, Montague, 1970, Chierchia and McConnell-Ginet, 1990, Heim and Kratzer, 1998]. That is to say that understanding a sentence is equivalent to knowing the actual conditions of the world under which the sentences would be (judged) true [Wittgenstein, 1922]. If grammatical sentences are required for sentential inference, as per a truth conditional approach [Montague, 1970], then permuted sentences should be meaningless. Put another way, the meanings of highly permuted sentences (if they exist) are not propositions, and thus those sentences don't have truth conditions. Only from their truth conditions of sentences can we tell if a sentence entails another. In short, the textual entailment task is technically undefined in our “unnatural” setting.

Since existing definitions don't immediately extend to unnatural word orders, we outline several hypothetical *systematic* ways that a model might perform, had it been sensitive to word order. We hypothesize two models that operate on the first principles of NLI, and one that doesn't. In the first case, Model A deems permuted sentences meaningless (devoid of truth values), as formal semantic theories of human language would predict. Thus, it assigns “*neutral*” to every permuted example. Next, Model B does not deem permuted sentences meaningless, and attempts to understand them. Humans find understanding permuted sentences difficult (see our human evaluations in §4.5.2). Model B could also similarly struggle to decipher the meaning, and just equally sample labels for each example (i.e., assigns equal probability mass to the outcome of each label). Finally, we hypothesize a non-systematic model, Model C, which attempts to treat permuted sentences as though they weren't permuted at all. This model could operate similarly as bag-of-words (BOW), and thus always assign the same label to the permuted examples as it would to the un-permuted examples. If the model failed to assign the original gold label to the original unpermuted examples, it will also fail to assign the original gold label to its permutations; it will never get higher

accuracy on permuted examples than on unpermuted ones.

We find in our experiments that the state-of-the-art Transformer-based NLI models (as well as pre-Transformer class of models) do not perform like any of the above hypothetical models. They perform closest to Model C, but are, in some cases, actually able to achieve *higher* accuracy on permuted examples. To better quantitatively describe this behaviour, we introduce our suite of **Permutation Acceptance** metrics that enable us to quantify how accepting models are of permuted sentences.

4.2 Experimental Setup

4.2.1 Constructing the permuted dataset.

For a given dataset D having splits D_{train} and D_{test} , we first train an NLI model M on D_{train} to achieve comparable accuracy to what was reported in the original papers. We then construct a randomized version of D_{test} , which we term as \hat{D}_{test} such that: for each example $(p_i, h_i, y_i) \in D_{\text{test}}$ (where p_i and h_i are the premise and hypothesis sentences of the example respectively and y_i is the gold label), we use a permutation operator \mathcal{F} that returns a list (\hat{P}_i, \hat{H}_i) of q permuted sentences (\hat{p}_i and \hat{h}_i), where q is a hyperparameter. \mathcal{F} essentially permutes all positions of the words in a given sentence (i.e., either in premise or hypothesis) with the restriction that *no words maintain their original position*. In our initial setting, we do not explicitly control the placement of the words relative to their original neighbors, but we analyze clumping effects in §4.4. \hat{D}_{test} now consists of $|D_{\text{test}}| \times q$ examples, with q different permutations of hypothesis and premise for each original test example pair. If a sentence S (e.g., h_i) contains w words, then the total number of available permutations of S are $(w - 1)!$, thus making the output of \mathcal{F} a list of $\binom{(w-1)!}{q}$ permutations in this case. For us, the space of possible outputs is larger, since we permute p_i and h_i separately (and ignore examples for which any $|S| \leq 5$).

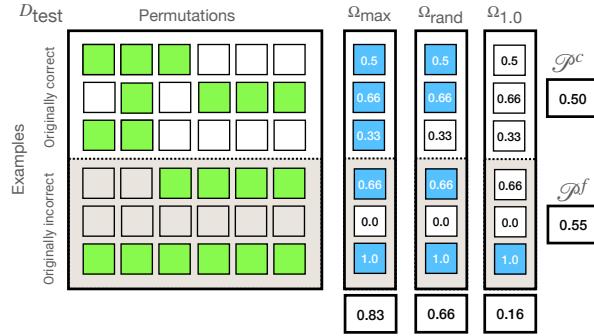


Figure 4.1 Graphical representation of the Permutation Acceptance class of metrics. Given a sample test set D_{test} with six examples, three of which originally predicted correctly (model predicts gold label), three incorrectly (model fails to predict gold label), with $n = 6$ permutations, $\Omega_{\max}, \Omega_{\text{rand}}, \Omega_{1.0}, \mathcal{P}^c$ and \mathcal{P}^f are provided. Green boxes indicate permutations accepted by the model. Blue boxes mark examples that crossed each threshold and were used to compute the corresponding metric.

4.2.2 Defining Permutation Acceptance.

The choice of q naturally allows us to analyze a statistical view of the predictability of a model on the permuted sentences. To that end, we define the following notational conventions. Let \mathcal{A} be the original accuracy of a given model M on a dataset D , and c be the number of examples in a dataset which are marked as correct according to the standard formulation of accuracy for the original dataset (i.e., they are assigned the ground truth label). Typically \mathcal{A} is given by $\frac{c}{|D_{\text{test}}|}$ or $\frac{c}{|D_{\text{dev}}|}$.

Let $\Pr_M(\hat{P}_i, \hat{H}_i)_{\text{cor}}$ then be the percentage of q permutations of an example (p_i, h_i) assigned the ground truth label y_i by M :

$$\Pr_M(\hat{P}_i, \hat{H}_i)_{\text{cor}} = \frac{1}{q} \sum_{(\hat{p}_j \in \hat{P}_i, \hat{h}_j \in \hat{H}_i)} ((M(\hat{p}_j, \hat{h}_j) = y_i) \rightarrow 1) \quad (4.1)$$

To get an overall summary score, we let Ω_x be the percentage of examples $(p_i, h_i) \in D_{\text{test}}$ for which $\Pr_M(\hat{P}_i, \hat{H}_i)_{\text{cor}}$ exceeds a predetermined threshold $0 < x < 1$. Concretely,

a given example will count as correct according to Ω_x if more than x percent of its permutations (\hat{P}_i and \hat{H}_i) are assigned y_i by the model M . Mathematically,

$$\Omega_x = \frac{1}{|D_{\text{test}}|} \sum_{(p_i, h_i) \in D_{\text{test}}} ((\Pr_M(\hat{P}_i, \hat{H}_i)_{\text{cor}} > x) \rightarrow 1). \quad (4.2)$$

There are two specific cases of Ω_x that we are most interested in. First, we define Ω_{\max} or the **Maximum Accuracy**, where $x = 1/|D_{\text{test}}|$. In short, Ω_{\max} gives the percentage of examples $(p_i, h_i) \in D_{\text{test}}$ for which there is *at least one* permutation (\hat{p}_j, \hat{h}_j) that model M assigns the gold label y_i ¹. Second, we define Ω_{rand} , or **Random Baseline Accuracy**, where $x = 1/m$ or chance probability (for balanced m -way classification, where $m = 3$ in NLI). This metric is less stringent than Ω_{\max} , as it counts an example if at least *one third* of its permutations are assigned the gold label (hence provides a lower-bound relaxation). See Figure 4.1 for a graphical representation of Ω_x .

We also define D^f to be the list of examples originally marked incorrect according to \mathcal{A} , but are now deemed correct according Ω_{\max} . D^c is the list of examples originally marked correct according to \mathcal{A} . Thus, we should expect $D^f < D^c$ for models that have high accuracy. Additionally, we define \mathcal{P}^c and \mathcal{P}^f , as the dataset average percentage of permutations which predicted the gold label, when the examples were originally correct (D^c) and when the examples were originally incorrect (D^f) as per \mathcal{A} (hence, flipped) respectively.

$$\mathcal{P}^c = \frac{1}{|D^c|} \sum_{i=0}^{|D^c|} M(\hat{P}_i, \hat{H}_i)_{\text{cor}} \quad (4.3)$$

\mathcal{P}^f is defined similarly by replacing D^c by D^f . Note that for a classic BOW model, $\mathcal{P}^c = 100$ and $\mathcal{P}^f = 0$, because it would rely on the words alone (not their order) to make its classification decision. Since permuting removes no words, BOW models

¹Theoretically, $\Omega_{\max} \rightarrow 1$ if the number of permutations q is large. Thus, in our experiments we set $q = 100$.

should come to the same decisions for permuted examples as for the originals.

4.3 Evaluated Models

We run our experiments on two types of models: **(a)** Transformer-based models and **(b)** Non-Transformer Models. In **(a)**, we investigate the state-of-the-art pre-trained models such as RoBERTa-Large Liu et al. [2019a], BART-Large Lewis et al. [2020] and DistilBERT Sanh et al. [2020]. For **(b)** we consider several recurrent and convolution based neural networks, such as InferSent Conneau et al. [2017], Bidirectional LSTM Collobert and Weston [2008] and ConvNet Zhao et al. [2015]. We train all models on MNLI, and evaluate on in-distribution (SNLI and MNLI) and out-of-distribution datasets (ANLI). We independently verify results of **(a)** using both our fine-tuned model using HuggingFace Transformers Wolf et al. [2020] and pre-trained checkpoints from FairSeq Ott et al. [2019] (using PyTorch Model Hub). For **(b)**, we use the InferSent codebase. We sample $q = 100$ permutations for each example in D_{test} , and use 100 seeds for each of those permutations to ensure full reproducibility. We drop examples from test sets where we are unable to compute *all unique* randomizations, typically these are examples with sentences of length of less than 6 tokens.²

4.4 Results

4.4.1 Models accept many permuted examples.

We find Ω_{\max} is very high for models trained and evaluated on MNLI (in-domain generalization), reaching 98.7% on MNLI dev. and test sets (in RoBERTa, compared to \mathcal{A} of 90.6% (Table 4.2). Recall, human accuracy is approximately 92% on MNLI dev., Nangia and Bowman 2019). This shows that there exists at least one permutation (usually

²Code, data, and model checkpoints are available at <https://github.com/facebookresearch/unlu>.

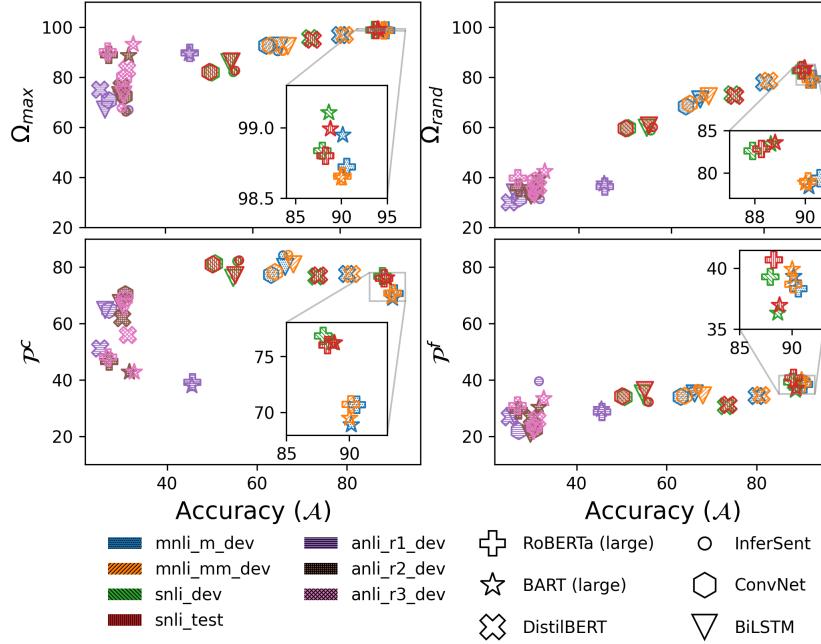


Figure 4.2 Comparison of Ω_{\max} , Ω_{rand} , \mathcal{P}^c and \mathcal{P}^f with the model accuracy \mathcal{A} on multiple datasets, where all models are trained on the MNLI corpus Williams et al. [2018c].

many more) for almost all examples in D_{test} such that model M predicts the gold label. We also observe high Ω_{rand} at 79.4%, showing that there are many examples for which the models outperform even a random baseline in accepting permuted sentences. We provide an example of the behaviour in Table 4.1.

Evaluating out-of-domain generalization with ANLI dataset splits resulted in an Ω_{\max} value that is notably higher than \mathcal{A} (89.7% Ω_{\max} for RoBERTa compared to 45.6% \mathcal{A}). As a consequence, we encounter many *flips*, i.e., examples where the model is unable to predict the gold label, but at least one permutation of that example is able to. However, recall this analysis expects us to know the gold label upfront, so this test can be thought of as running a word-order probe test on the model until the model predicts the gold label (or give up by exhausting our set of q permutations). For out-of-

domain generalization, Ω_{rand} decreases considerably (36.4% Ω_{rand} on A1), which means fewer permutations are accepted by the model. Next, recall that a classic bag-of-words model would have $\mathcal{P}^c = 100$ and $\mathcal{P}^f = 0$. No model performs strictly like a classic bag of words although they do perform somewhat BOW-like ($\mathcal{P}^c >> \mathcal{P}^f$ for all test splits, Figure 4.2). We find this BOW-likeness to be higher for certain non-Transformer models, (InferSent) as they exhibit higher \mathcal{P}^c (84.2% for InferSent compared to 70.7% for RoBERTa on MNLI).

Investigating other Ω values

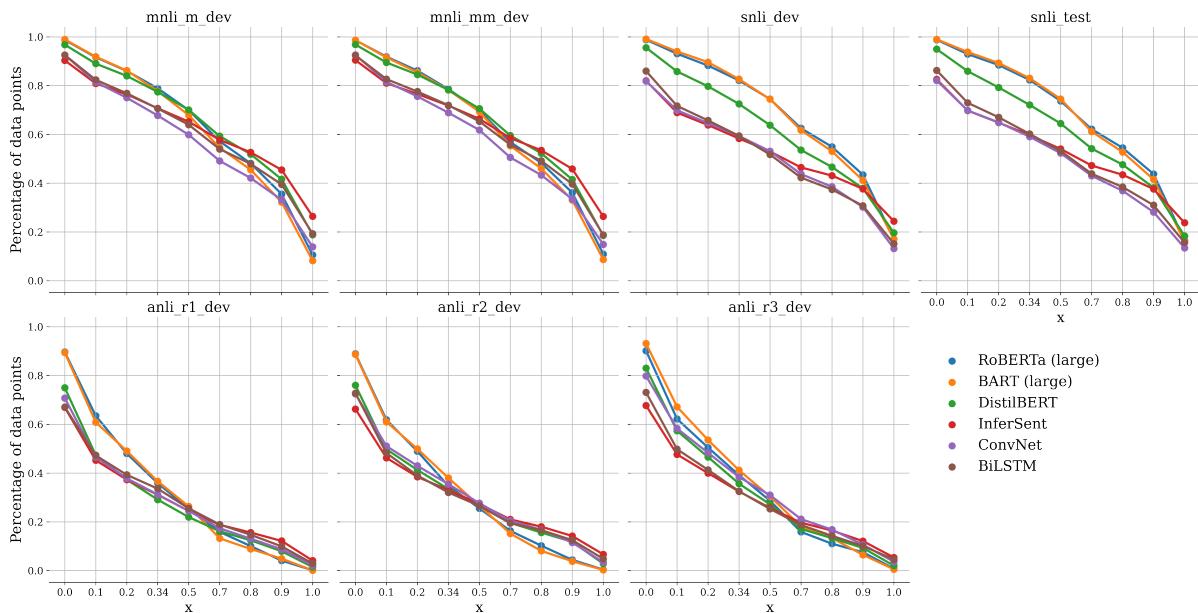


Figure 4.3 Ω_x threshold for all datasets with varying x and computing the percentage of examples that fall within the threshold. The top row consists of in-distribution datasets (MNLI, SNLI) and the bottom row contains out-of-distribution datasets (ANLI)

We defined two variations of Ω_x , Ω_{\max} and Ω_{rand} , but theoretically it is possible to define any arbitrary threshold percentage x to evaluate the unnatural language inference mechanisms of different models. In Figure 4.3 we show the effect of different

Model	Eval. Dataset	\mathcal{A}	Ω_{\max}	\mathcal{P}^c	\mathcal{P}^f	Ω_{rand}
RoBERTa-Large	MNLI_m_dev	0.906	0.987	0.707	0.383	0.794
	MNLI_mm_dev	0.901	0.987	0.707	0.387	0.790
	SNLI_dev	0.879	0.988	0.768	0.393	0.826
	SNLI_test	0.883	0.988	0.760	0.407	0.828
	A1*	0.456	0.897	0.392	0.286	0.364
	A2*	0.271	0.889	0.465	0.292	0.359
	A3*	0.268	0.902	0.480	0.308	0.397
	Mean	0.652	0.948	0.611	0.351	0.623
BART-Large	MNLI_m_dev	0.902	0.989	0.689	0.393	0.784
	MNLI_mm_dev	0.900	0.986	0.695	0.399	0.788
	SNLI_dev	0.886	0.991	0.762	0.363	0.834
	SNLI_test	0.888	0.990	0.762	0.370	0.836
	A1*	0.455	0.894	0.379	0.295	0.374
	A2*	0.316	0.887	0.428	0.303	0.397
	A3*	0.327	0.931	0.428	0.333	0.424
	Mean	0.668	0.953	0.592	0.351	0.634
DistilBERT	MNLI_m_dev	0.800	0.968	0.775	0.343	0.779
	MNLI_mm_dev	0.811	0.968	0.775	0.346	0.786
	SNLI_dev	0.732	0.956	0.767	0.307	0.731
	SNLI_test	0.738	0.950	0.770	0.312	0.725
	A1*	0.251	0.750	0.511	0.267	0.300
	A2*	0.300	0.760	0.619	0.265	0.343
	A3*	0.312	0.830	0.559	0.259	0.363
	Mean	0.564	0.883	0.682	0.300	0.575
InferSent	MNLI_m_dev	0.658	0.904	0.842	0.359	0.712
	MNLI_mm_dev	0.669	0.905	0.844	0.368	0.723
	SNLI_dev	0.556	0.820	0.821	0.323	0.587
	SNLI_test	0.560	0.826	0.824	0.321	0.600
	A1*	0.316	0.669	0.425	0.395	0.313
	A2*	0.310	0.662	0.689	0.249	0.330
	A3*	0.300	0.677	0.675	0.236	0.332
	Mean	0.481	0.780	0.731	0.322	0.514
ConvNet	MNLI_m_dev	0.631	0.926	0.773	0.340	0.684
	MNLI_mm_dev	0.640	0.926	0.782	0.343	0.694
	SNLI_dev	0.506	0.819	0.813	0.339	0.597
	SNLI_test	0.501	0.821	0.809	0.341	0.596
	A1*	0.271	0.708	0.648	0.218	0.316
	A2*	0.307	0.725	0.703	0.224	0.356
	A3*	0.306	0.798	0.688	0.234	0.388
	Mean	0.452	0.817	0.745	0.291	0.519
BiLSTM	MNLI_m_dev	0.662	0.925	0.800	0.351	0.711
	MNLI_mm_dev	0.681	0.924	0.809	0.344	0.724
	SNLI_dev	0.547	0.860	0.762	0.351	0.598
	SNLI_test	0.552	0.862	0.771	0.363	0.607
	A1*	0.262	0.671	0.648	0.271	0.340
	A2*	0.297	0.728	0.672	0.209	0.328
	A3*	0.304	0.731	0.656	0.219	0.331
	Mean	0.472	0.814	0.731	0.301	0.520

Table 4.2 Statistics for Transformer-based models trained on MNLI corpus Williams et al. [2018c]. The highest values are bolded (red indicates the model most insensitive to permutation) per metric and per model class (Transformers and non-Transformers). A1*, A2* and A3* refer to the ANLI dev. sets [Nie et al., 2020].

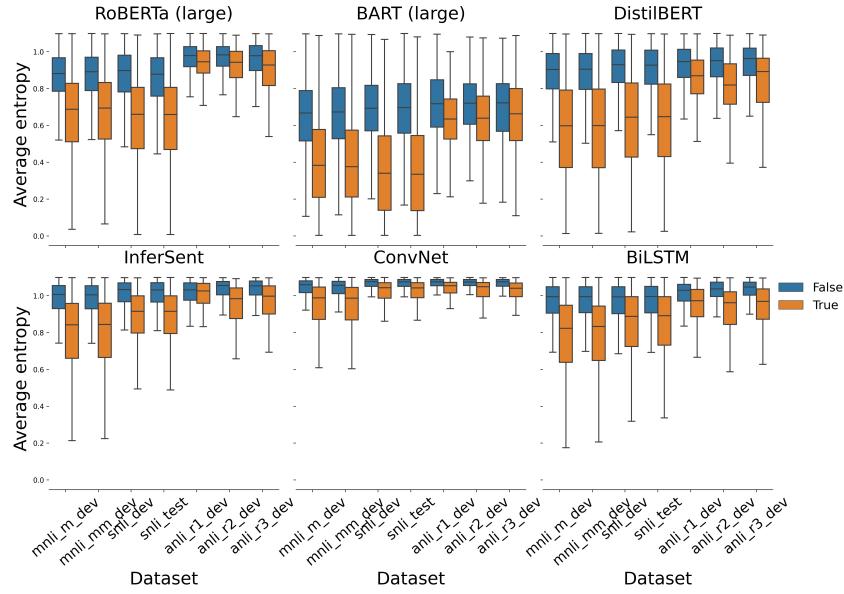


Figure 4.4 Average entropy of model confidences on permutations that yielded the correct results for Transformer-based models (top) and Non-Transformer-based models (bottom). Results are shown for D^c (orange) and D^f (blue). The boxes show the quartiles of the entropy distributions.

thresholds, including Ω_{\max} where $x = 1/|D_{\text{test}}|$ and Ω_{rand} where $x = 0.34$. We observe for in-distribution datasets (top row, MNLI and SNLI splits), in the extreme setting when $x = 1.0$, there are more than 10% of examples available, and more than 25% in case of InferSent and DistilBERT. For out-of-distribution datasets (bottom row, ANLI splits) we observe a much lower trend, suggesting generalization itself is the bottleneck in permuted sentence understanding.

4.4.2 Models are very confident.

The phenomenon we observe would be of less concern if the correct label prediction was just an outcome of chance, which could occur when the entropy of the log probabilities of the model output is high (suggesting uniform probabilities on entailment, neutral and contradiction labels, recall Model B from §4.1). We first investigate the

model probabilities for the Transformer-based models on the permutations that lead to the correct answer in Figure 4.4. We find overwhelming evidence that model confidences on in-distribution datasets (MNLI, SNLI) are highly skewed, resulting in low entropy, and it varies among different model types. BART proves to be the most skewed Transformer-based model. This skewness is not a property of model capacity, as we observe DistilBERT log probabilities to have similar skewness as RoBERTa (large) model, while exhibiting lower \mathcal{A} , Ω_{\max} , and Ω_{rand} .

For non-Transformers whose accuracy \mathcal{A} is lower, the Ω_{\max} achieved by these models are also predictably lower. We observe roughly the same relative performance in the terms of Ω_{\max} (Figure 4.2 and Table 4.2) and Average entropy (Figure 4.4). However, while comparing the averaged entropy of the model predictions, it is clear that there is some benefit to being a worse model—non-Transformer models are not as overconfident on randomized sentences as Transformers are. High confidence of Transformer models can be attributed to the *overthinking* phenomenon commonly observed in deep neural networks Kaya et al. [2019] and BERT-based models Zhou et al. [2020].

4.4.3 Similar artifacts in Chinese NLU.

Model	\mathcal{A}	Ω_{\max}	\mathcal{P}^c	\mathcal{P}^f	Ω_{rand}
RoBERTa-Large	0.784	0.988	0.726	0.339	0.773
InferSent	0.573	0.931	0.771	0.265	0.615
ConvNet	0.407	0.752	0.808	0.199	0.426
BiLSTM	0.566	0.963	0.701	0.271	0.611

Table 4.3 Results on evaluation on OCNLI Dev set. All models are trained on OCNLI corpus Hu et al. [2020a]. Bold marks the highest value per metric (red shows the model is insensitive to permutation).

We extended the experiments to the Original Chinese NLI dataset [Hu et al., 2020a, OCNLI], and re-used the pre-trained RoBERTa-Large and InferSent (non-Transformer)

models on OCNLI. Our findings are similar to the English results (Table 4.3), thereby suggesting that the phenomenon is not just an artifact of English text or tokenization.

4.4.4 Other Results.

We investigated the effect of sentence length (which correlates with number of possible permutations), and hypothesis-only randomization (models exhibit similar phenomenon even when only hypothesis is permuted). In terms of sentence length, we observe that shorter sentences in general have a somewhat higher probability of acceptance for examples which was originally predicted correctly—since shorter sentences have fewer unique permutations (Figure 4.5). However, for the examples which were originally incorrect, the trend is not present.

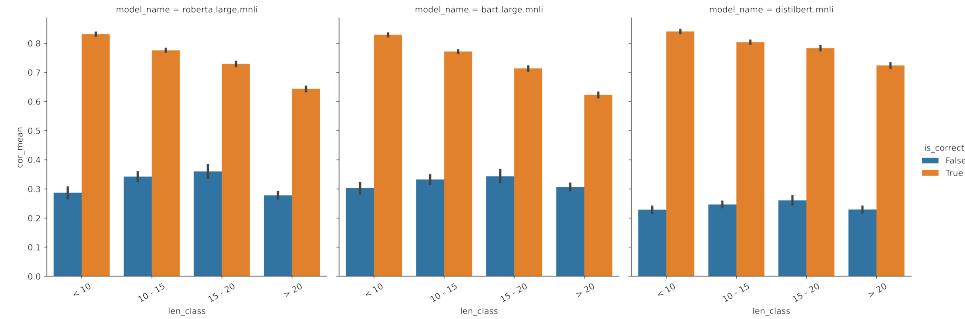


Figure 4.5 Length and Permutation Acceptance by Transformer-based models.

In recent years, the impact of the hypothesis sentence [Gururangan et al., 2018a, Tsuchiya, 2018, Poliak et al., 2018] on NLI classification has been a topic of much interest. As we define in §4.1, logical entailment can only be defined for pairs of propositions. We investigated one effect where we randomize only the hypothesis sentences while keeping the premise intact. Figure 4.6 and Figure 4.7 shows that the Ω_{\max} value is almost the same for the two schemes; randomizing the hypothesis alone also leads the model to accept many permutations.

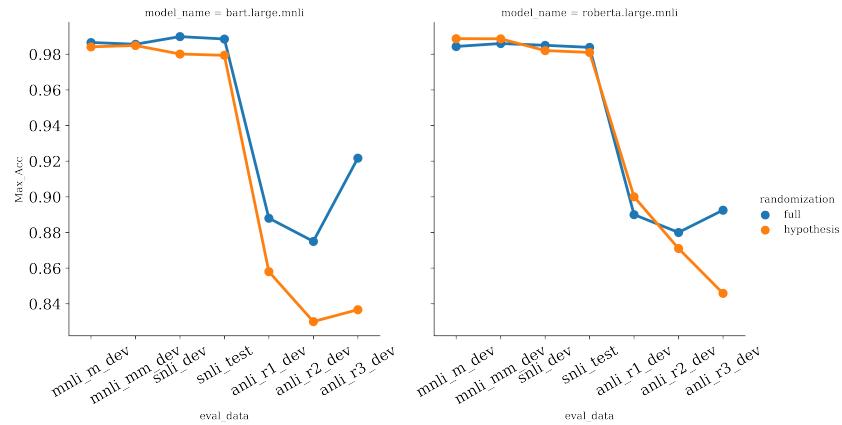


Figure 4.6 Comparing the effect between randomizing both premise and hypothesis and only hypothesis on two Transformer-based models, RoBERTa and BART. Here, we observe the difference of Ω_{\max} is marginal in in-distribution datasets (SNLI, MNLI), while hypothesis-only randomization is worse for out-of-distribution datasets (ANLI).

4.5 Analysis

4.5.1 Analyzing Syntactic Structure Associated with Tokens

A natural question to ask following our findings: what is it about particular permutations that leads models to accept them? Since the permutation operation is drastic and only rarely preserves local word relations, we first investigate whether there exists a relationship between Permutation Acceptance scores and local word order preservation. Concretely, we compare bi-gram word overlap (BLEU-2) with the percentage of permutations that are deemed correct (Figure 4.8).³ Although the probability of a permuted sentence to be predicted correctly does appear to track BLEU-2 score (Figure 4.8), the percentage of examples which were assigned the gold label by the

³We observe, due to our permutation process, the maximum BLEU-3 and BLEU-4 scores are negligibly low (< 0.2 BLEU-3 and < 0.1 BLEU-4), already calling into question the hypothesis that n-grams are the sole explanation for our finding. Because of this, we only compare BLEU-2 scores. Detailed experiments on specially constructed permutations that cover the entire range of BLEU-3 and BLEU-4 is provided in ??.

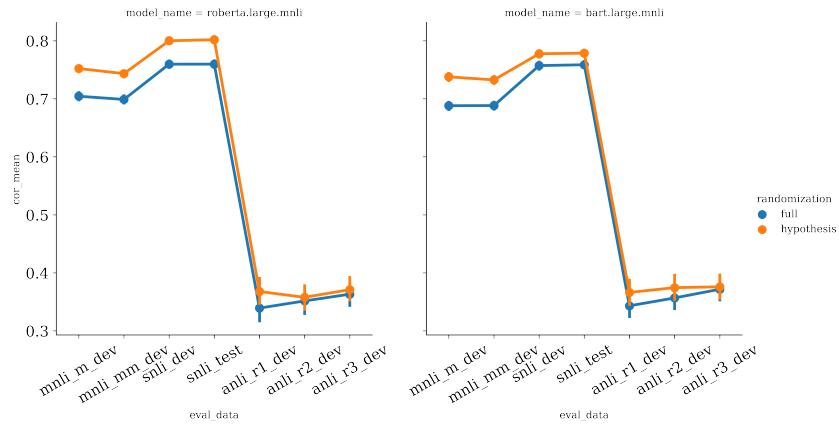


Figure 4.7 Comparing the effect between randomizing both premise and hypothesis and only hypothesis on two Transformer-based models, RoBERTa and BART. In this figure, we compare the mean number of permutations which elicited correct response, and naturally the hypothesis-only randomization causes more percentage of randomizations to be correct.

Transformer-based models is still higher than we would expect from permutations with lower BLEU-2 (66% for the lowest BLEU-2 range of 0 – 0.15), suggesting preserved relative word order alone cannot explain the high permutation acceptance rates.

Thus, we find that local order preservation does correlate with Permutation Acceptance, but it doesn't fully explain the high Permutation Acceptance scores. We now further ask whether Ω is related to a more abstract measure of local word relations, i.e., part-of-speech (POS) neighborhood.

Many syntactic formalisms, like Lexical Functional Grammar [Kaplan and Bresnan, 1995, Bresnan et al., 2015, LFG], Head-drive Phrase Structure Grammar [Pollard and Sag, 1994, HPSG] or Lexicalized Tree Adjoining Grammar [Schabes et al., 1988, Abeille, 1990, LTAG], are “lexicalized”, i.e., individual words or morphemes bear syntactic features telling us which other words they can combine with. For example, “buy” could be associated with (at least) two lexicalized syntactic structures, one containing two noun phrases (as in Kim bought cheese), and another with three (as in Lee bought Logan

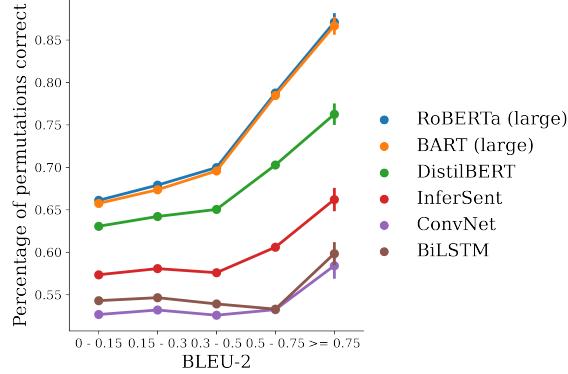


Figure 4.8 BLEU-2 score versus acceptability of permuted sentences across all test datasets. RoBERTa and BART performance is similar but differs considerably from the performance of non-Transformer-based models, such as InferSent and ConvNet.

cheese). We speculate that our NLI models might accept permuted examples at high rates, because they are (perhaps noisily) reconstructing the original sentence from abstract, word-anchored information about common neighbors.

To test this, we POS-tagged D_{train} using 17 Universal Part-of-Speech tags (using spaCy, Honnibal et al. 2020b). For each $w_i \in S_i$, we compute the occurrence probability of POS tags on tokens in the *neighborhood* of w_i . The neighborhood is specified by the radius r (a symmetrical window r tokens from $w_i \in S_i$ to the left and right). We denote this sentence level probability of neighbor POS tags for a word w_i as $\psi_{\{w_i, S_i\}}^r \in \mathcal{R}^{17}$. Sentence-level word POS neighbor scores can be averaged across D_{train} to get a type level score $\psi_{\{w_i, D_{\text{train}}\}}^r \in \mathcal{R}^{17}, \forall w_i \in D_{\text{train}}$. Then, for a sentence $S_i \in D_{\text{test}}$, for each word $w_i \in S_i$, we compute a **POS mini-tree overlap score**:

$$\beta_{\{w_i, S_i\}}^k = \frac{1}{k} |\operatorname{argmax}_k \psi_{\{w_i, D_{\text{train}}\}}^r \cap \operatorname{argmax}_k \psi_{\{w_i, S_i\}}^r| \quad (4.4)$$

Concretely, $\beta_{\{w_i, S_i\}}^k$ computes the overlap of top- k POS tags in the neighborhood of

a word w_i in S with that of the train statistic. If a word has the same mini-tree in a given sentence as it has in the training set, then the overlap would be 1. For a given sentence S_i , the aggregate $\beta_{\{S_i\}}^k$ is defined by the average of the overlap scores of all its words: $\beta_{\{S_i\}}^k = \frac{1}{|S_i|} \sum_{w_i \in S_i} \beta_{\{w_i, S_i\}}^k$, and we call it a POS minitree *signature*. We can also compute the POS minitree signature of a permuted sentence \hat{S}_i to have $\beta_{\{\hat{S}_i\}}^k$. If the permuted sentence POS signature comes close to that of the true sentence, then their ratio (i.e., $\beta_{\{\hat{S}_i\}}^k / \beta_{\{S_i\}}^k$) will be close to 1. Also, since POS signature is computed with respect to the train distribution, a ratio of > 1 indicates that the permuted sentence is closer to the overall train statistic than to the original unpermuted sentence in terms of POS signature. If high overlap with the training distribution correlates with percentage of permutations deemed correct, then our models treat words as if they project syntactic minitrees. Figure 4.9 provides a snapshot a word "river" from the test set and shows how the POS signature distribution of the word in a particular example match with that of aggregated training statistic. In practice, we select the top k POS tags for the word in the test signature as well as the train, and calculate their overlap. When comparing the model performance with permuted sentences, we compute a ratio between the original test overlap score and an overlap score calculated instead from the permuted test. In the Figure 4.9, 'river' would have a POS tag minitree score of 0.75.

We investigate the relationship with percentage of permuted sentences accepted with $\beta_{\{\hat{S}_i\}}^k / \beta_{\{S_i\}}^k$ in Figure 4.10. We observe that the POS Tag Minitree hypothesis holds for Transformer-based models, RoBERTa, BART and DistilBERT, where the percentage of accepted pairs increase as the sentences have higher overlap with the un-permuted sentence in terms of POS signature. For non-Transformer models such as InferSent, ConvNet, and BiLSTM models, the POS signature ratio to percentage of correct permutation remains the same or decreases, suggesting that the reasoning process employed by these models does not preserve local abstract syntax structure (i.e., POS neighbor relations).

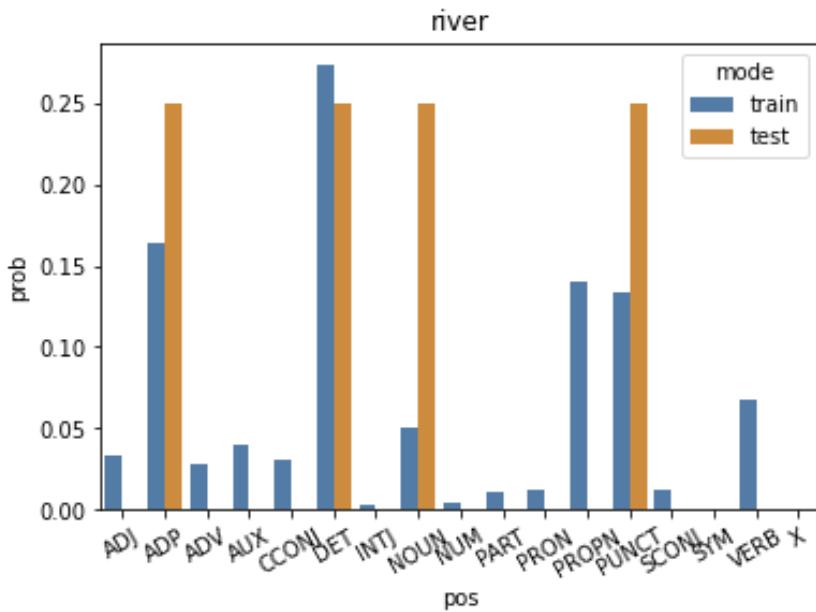


Figure 4.9 Example POS signature for the word ‘river’, calculated with a radius of 2. Probability of each neighbor POS tag is provided. Orange examples come from the permuted test set, and blue come from the original training data.

4.5.2 Human Evaluation

We expect humans to struggle with UNLI, given our intuitions and the sentence superiority findings (but see Mollica et al. 2020). To test this, we presented two experts in NLI (one a linguist) with permuted sentence pairs to label.⁴ Concretely, we draw equal number of examples from MNLI Matched dev set (100 examples where RoBERTa predicts the gold label, D^c and 100 examples where it fails to do so, D^f), and then permute these examples using \mathcal{F} . The experts were given no additional information (recall that it is common knowledge that NLI is a roughly balanced 3-way classification task). Unbeknownst to the experts, all permuted sentences in the sample were actually accepted

⁴Concurrent work by Gupta et al. [2021] found that untrained crowdworkers accept NLI examples that have been subjected to different kinds of perturbations at roughly most frequent class levels—i.e., only 35% of the time.

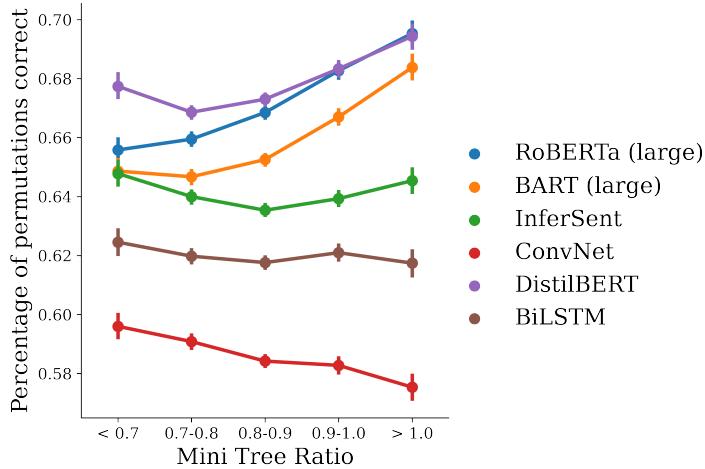


Figure 4.10 POS Tag Mini Tree overlap score and percentage of permutations which the models assigned the gold-label.

Evaluator	Accuracy	Macro F1	Acc on D^c	Acc on D^f
X	0.581 ± 0.068	0.454	0.649 ± 0.102	0.515 ± 0.089
Y	0.378 ± 0.064	0.378	0.411 ± 0.098	0.349 ± 0.087

Table 4.4 Human (expert) evaluation on 200 permuted examples from the MNLI matched development set. Half of the permuted pairs contained shorter sentences and the other, longer ones. All permuted examples were assigned the gold label by RoBERTa-Large.

by the RoBERTa (large) model (trained on MNLI dataset). We observe that the experts performed much worse than RoBERTa (Table 4.4), although their accuracy was a bit higher than random. We also find that for both experts, accuracy on permutations from D^c was higher than on D^f , which verifies findings that showed high word overlap can give hints about the ground truth label [Dasgupta et al., 2018, Poliak et al., 2018, Gururangan et al., 2018a, Naik et al., 2019].

4.5.3 Training by Maximizing Entropy

We propose an initial attempt to mitigate the effect of correct prediction on permuted examples. As we observe in §4.4.2, model entropy on permuted examples is significantly lower than expected. Neural networks tend to output higher confidence than random for even unknown inputs Gandhi and Lake [2020], which might be an underlying cause of the high Permutation Acceptance.

An ideal model would be ambivalent about randomized ungrammatical sentences. Thus, we train NLI models baking in the principle of mutual exclusivity [Gandhi and Lake, 2020] by maximizing model entropy. Concretely, we fine-tune RoBERTa on MNLI while maximizing the entropy (\mathcal{H}) on a subset of n randomized examples $((\hat{p}_i, \hat{r}_i)$, for each example (p, h) in MNLI. We modify the loss function as follows:

$$\mathcal{L} = \operatorname{argmin}_{\theta} \sum_{((p,h),y)} y \log(p(y|(p,h);\theta)) + \sum_{i=1}^n \mathcal{H}(y|(\hat{p}_i, \hat{h}_i);\theta) \quad (4.5)$$

Using this maximum entropy method ($n = 1$), we find that the model improves considerably with respect to its robustness to randomized sentences, all while taking no hit to accuracy (Table 4.5). We observe that no model reaches a Ω_{\max} score close to 0, suggesting further room to explore other methods for decreasing models’ Permutation Acceptance. Similar approaches have also proven useful [Gupta et al., 2021] for other tasks as well.

4.6 Related Work

Researchers in NLP have realized the importance of syntactic structure in neural networks going back to Tabor [1994]. An early hand annotation effort on PASCAL RTE [Dagan et al., 2006] suggested that “syntactic information alone was sufficient to make a judgment” for roughly one third of examples [Vanderwende and Dolan, 2005]. Anecdotally, large generative language models like GPT-2 or -3 exhibit a seemingly human-

Eval Dataset	\mathcal{A} (V)	\mathcal{A} (ME)	Ω_{\max} (V)	Ω_{\max} (ME)
MNLI_m_dev	0.905	0.908	0.984	0.328
MNLI_mm_dev	0.901	0.903	0.985	0.329
SNLI_test	0.882	0.888	0.983	0.329
SNLI_dev	0.879	0.887	0.984	0.333
ANLI_r1_dev	0.456	0.470	0.890	0.333
ANLI_r2_dev	0.271	0.258	0.880	0.333
ANLI_r3_dev	0.268	0.243	0.892	0.334

Table 4.5 NLI Accuracy (\mathcal{A}) and Permutation Acceptance metrics (Ω_{\max}) of RoBERTa when trained on MNLI dataset using vanilla (V) and Maximum Random Entropy (ME) method.

like ability to generate fluent and grammatical text [Goldberg, 2019a, Wolf, 2019b]. However, the jury is still out as to whether transformers genuinely acquire syntax.

4.6.1 Models appear to have acquired syntax.

When researchers have peeked inside Transformer LM’s pretrained representations, familiar syntactic structure [Hewitt and Manning, 2019a, Jawahar et al., 2019a, Lin et al., 2019, Warstadt and Bowman, 2020, Wu et al., 2020], or a familiar order of linguistic operations [Jawahar et al., 2019a, Tenney et al., 2019], has appeared. There is also evidence, notably from agreement attraction phenomena [Linzen et al., 2016] that transformer-based models pretrained on LM do acquire some knowledge of natural language syntax [Gulordava et al., 2018a, Chrupała and Alishahi, 2019, Jawahar et al., 2019a, Lin et al., 2019, Manning et al., 2020a, Hawkins et al., 2020, Linzen and Baroni, 2021]. Results from other phenomena [Warstadt and Bowman, 2020] such as NPI licensing [Warstadt et al., 2019a] lend additional support. The claim that LMs acquire some syntactic knowledge has been made not only for transformers, but also for convolutional neural nets [Bernard and Lappin, 2017], and RNNs [Gulordava et al.,

2018a, van Schijndel and Linzen, 2018, Wilcox et al., 2018, Zhang and Bowman, 2018, Prasad et al., 2019, Ravfogel et al., 2019]—although there are many caveats (e.g., Ravfogel et al. 2018, White et al. 2018, Davis and van Schijndel 2020, Chaves 2020, Da Costa and Chaves 2020, Kodner and Gupta 2020).

4.6.2 Models appear to struggle with syntax.

Several works have cast doubt on the extent to which NLI models in particular know syntax (although each work adopts a slightly different idea of what “knowing syntax” entails). For example, McCoy et al. [2019] argued that the knowledge acquired by models trained on NLI (for at least some popular datasets) is actually not as syntactically sophisticated as it might have initially seemed; some transformer models rely mainly on simpler, non-humanlike heuristics. In general, transformer LM performance has been found to be patchy and variable across linguistic phenomena [Dasgupta et al., 2018, Naik et al., 2018, An et al., 2019, Ravichander et al., 2019, Jeretic et al., 2020]. This is especially true for syntactic phenomena [Marvin and Linzen, 2018, Hu et al., 2020b, Gauthier et al., 2020, McCoy et al., 2020, Warstadt et al., 2020a], where transformers are, for some phenomena and settings, worse than RNNs [van Schijndel et al., 2019]. From another angle, many have explored architectural approaches for increasing a network’s sensitivity to syntactic structure [Chen et al., 2017, Li et al., 2020]. Williams et al. [2018a] showed that learning jointly to perform NLI and to parse resulted in parse trees that match no popular syntactic formalisms. Furthermore, models trained explicitly to differentiate acceptable sentences from unacceptable ones (i.e., one of the most common syntactic tests used by linguists) have, to date, come nowhere near human performance [Warstadt et al., 2019c].

4.6.3 Insensitivity to Perturbation.

Most relatedly, several concurrent works [Pham et al., 2020a, Alleman et al., 2021, Gupta et al., 2021, Sinha et al., 2021b, Parthasarathi et al., 2021] investigated the effect of word order permutations on transformer NNs. Pham et al. [2020a] is very nearly a proper subset of our work except for investigating additional tasks (i.e. from the GLUE benchmark of Wang et al. 2018) and performing a by-layer-analysis. Gupta et al. [2021] also relies on the GLUE benchmark, but additionally investigates other types of “destructive” perturbations. Our contribution differs from these works in that we additionally include the following: we (i) outline theoretically-informed predictions for how models *should be expected* to react to permuted input (we outline a few options), (ii) show that permuting can “flip” an incorrect prediction to a correct one, (iii) show that the problem isn’t specific to Transformers, (iv) show that the problem persists on out of domain data, (v) offer a suite of flexible metrics, and (vi) analyze *why* models might be accepting permutations (BLEU and POS-tag neighborhood analysis). Finally, we replicate our findings in another language. While our work (and Pham et al., Gupta et al.) only permutes data during fine-tuning and/or evaluation, recently Sinha et al. explored the sensitivity during pre-training, and found that models trained on n-gram permuted sentences perform remarkably close to regular MLM pre-training. In the context of generation, Parthasarathi et al. [2021] crafted linguistically relevant perturbations (on the basis of part-of-speech tagging and dependency parsing) to evaluate whether permutation hinders automatic machine translation models. Relatedly, but not for translation, Alleman et al. [2021] investigated a smaller inventory of perturbations with emphasis on phrasal boundaries and the effects of n-gram perturbations on different layers in the network.

4.6.4 NLI Models are very sensitive to words.

NLI models often over-attend to particular words to predict the correct answer [Gururangan et al., 2018a, Clark et al., 2019]. Wallace et al. [2019] show that some short sequences of non-human-readable text can fool many NLU models, including NLI models trained on SNLI, into predicting a specific label. In fact, Ettinger [2020] observed that for one of three test sets, BERT loses some accuracy in word-perturbed sentences, but that there exists a subset of examples for which BERT’s accuracy remains intact. If performance isn’t affected (or if permutation helps, as we find it does in some cases), it suggests that these state-of-the-art models actually perform somewhat similarly to bag-of-words models Blei et al. [2003], Mikolov et al. [2013a].

4.7 Discussion

In this chapter, we observe that state-of-the-art models do not rely on sentence structure the way we think they should: NLI models (Transformer-based models, RNNs, and ConvNets) are largely insensitive to permutations of word order that corrupt the original syntax. This raises questions about the extent to which such systems understand “syntax”, and highlights the unnatural language understanding processes they employ. To summarize, our primary observations from this chapter are:

- **NLU models can still perform the task even if the word orders are scrambled**
We observed overwhelming evidence in §4.4.1 that state-of-the-art NLU models tend to perform the task comparably even on word order scrambled text, which has no inherent semantic meaning.
- **Certain permutations allow NLU models to flip classification labels, leading to better task scores.** We find that certain permutations of the order of words in a given input sentence pair can trigger the model to change the classification

labels from the baseline, leading to large performance gains in the scores of the NLI task. For instance, in the examples which the models find difficult to predict, a permutation of the word order can elicit the model to assign the gold label.

- **NLU models display rudimentary understanding of syntax, as evident by the preservation of abstract parts-of-speech neighborhood information.** We do find that models seem to have learned some syntactic information as is evidenced by a correlation between preservation of abstract POS neighborhood information and rate of acceptance by models, but these results do not discount the high rates of Permutation Acceptance, and require further verification.

Given these findings, and coupled with the observation that humans cannot perform UNLI at all well, the high rate of permutation acceptance that we observe leads us to conclude that current models do not yet “know syntax” in the fully systematic and humanlike way we would like them to. This study leads us to further investigate the training dynamics employed by these large language models. In the next chapter, we will investigate the training data dependency of one such model family in detail, to shed more light on the sentence processing pipelines employed by these models.

4.8 Follow-up findings in the community

Chapter 5

Probing syntax understanding through distributional hypothesis

The field of natural language processing (NLP) has become dominated by the pretrain-and-finetune paradigm, where we first obtain a good parametric *prior* in order to subsequently model downstream tasks accurately. In particular, masked language model (MLM) pre-training, as epitomized by BERT [Devlin et al., 2019], has proven wildly successful, although the precise reason for this success has remained unclear. On one hand, we can view BERT as the newest in a long line of NLP techniques Deerwester et al. [1990], Landauer and Dumais [1997], Collobert and Weston [2008], Mikolov et al. [2013b], Peters et al. [2018a] that exploit the well-known distributional hypothesis Harris [1954b]. On the other hand, it has been claimed that BERT “redisCOVERS the classical NLP pipeline” Tenney et al. [2019], suggesting that it has learned “the types of syntactic and semantic abstractions traditionally believed necessary for language processing” rather than “simply modeling complex co-occurrence statistics” (*ibid.* p.1).

In this chapter, we aim to uncover how much of MLM’s success comes from learning simple distributional information, as opposed to grammatical abstractions [Tenney et al., 2019, Manning et al., 2020b]. Thus, in this chapter I discuss our work [Sinha

et al., 2021a], where we disentangle these two hypotheses by measuring the effect of removing word order information during pre-training: any sophisticated (English) NLP pipeline would presumably depend on the syntactic information conveyed by the order of words. We find that surprisingly most of MLM’s high performance can in fact be explained by the “distributional prior” rather than its ability to replicate the classical NLP pipeline.

Concretely, we pre-train MLMs (RoBERTa, Liu et al. 2019b) on various corpora with permuted word order while preserving some degree of distributional information, and examine their downstream performance. We also experiment with training MLMs without positional embeddings, making them entirely order agnostic, and with training on a corpus sampled from the source corpus’s unigram distribution. We then evaluate these “permuted” models in a wide range of settings and compare with regularly-pre-trained models.

We demonstrate that pre-training on permuted data has surprisingly little effect on downstream task performance after fine-tuning (on non-shuffled training data). In our previous chapter we observed that MLMs are quite robust to permuting downstream test data (§4.4.1) and even do quite well using permuted “unnatural” downstream train data Sinha et al. [2021d], Gupta et al. [2021]. In this chapter, we show that downstream performance for “unnatural language pre-training” is much closer to standard MLM pre-training than one might expect.

In an effort to shed light on these findings, we experiment with various probing tasks. We verify via non-parametric probes that the permutations do in fact make the model worse at syntax-dependent tasks. However, just like on the downstream fine-tuning tasks, permuted models perform well on parametric syntactic probes, in some cases almost matching the unpermuted model’s performance, which is quite surprising given how important word order is crosslinguistically (Greenberg 1963, Dryer 1992, Cinque 1999, i.a.).

Our results can be interpreted in different ways. One could argue that our downstream and probing tasks are flawed, and that we need to examine models with examples that truly test strong generalization and compositionality. Alternatively, one could argue that prior works have overstated the dependence of human language understanding on word order, and that human language understanding depends less on the structure of the sentence and more on the structure of the *world*, which can be inferred to a large extent from distributional information. This work is meant to deepen our understanding of MLM pre-training and, through this, move us closer to finding out what is actually required for adequately modelling natural language.

5.1 Technical Background

5.2 Experimental Setup

5.2.1 Sentence word order permutation

To investigate to what extent the performance of MLM pre-training is a consequence of distributional information, we construct a training corpus devoid of natural word order but preserving local distributional information. We construct word order-randomized versions of the BookWiki corpus (the Toronto Books Corpus, Zhu et al. 2015, plus English Wikipedia) from Liu et al. [2019b], following the setup described in §4.2. Concretely, given a sentence S containing N words, we permute the sentence using a seeded random function \mathcal{F}_1 such that no word can remain in its original position. In total, there exist $(N - 1)!$ possible permutations of a given sentence. We randomly sample a single permutation per sentence, to keep the total dataset size similar to the original.

We extend the permutation function \mathcal{F}_1 to a function \mathcal{F}_n that preserves n -gram information. Specifically, given a sentence S of length N and n -gram value n , we sample

a starting position i for possible contiguous n -grams $\in \{0, N - n\}$ and convert the span $S[i, i + n]$ to a single token, to form \hat{S} , of length $\hat{N} = N - (n + 1)$. We continue this process repeatedly (without using the previously created n-grams) until there exists no starting position for selecting a contiguous n-gram in \hat{S} . For example, given a sentence of length $N = 6$, \mathcal{F}_4 will first convert one span of 4 tokens into a word, to have \hat{S} consisting of three tokens (one conjoined token of 4 contiguous words, and two left-over words). Then, the resulting sentence \hat{S} is permuted using \mathcal{F}_1 . We train RoBERTa models on four permutation variants of BookWiki corpus, $\mathcal{M}_1, \mathcal{M}_2, \mathcal{M}_3, \mathcal{M}_4$ for each n -gram value $\in \{1, 2, 3, 4\}$.

We provide pseudo-code for \mathcal{F}_i in Algorithm 1. Following the formulation in §4.2, we do not explicitly control whether the permuted words maintain any of their original neighbors. Thus, a certain amount of extra n-grams are expected to co-occur, purely as a product of random shuffling. We quantify the amount of such shuffling on a sample of 1 million sentences drawn from the BookWiki random corpus, and present the BLEU-2, BLEU-3 and BLEU-4 scores in Table 5.1. We provide a sample snapshot of the generated data in Table 5.18.

	BLEU-2	BLEU-3	BLEU-4
\mathcal{M}_1	0.493 +/- 0.12	0.177 +/- 0.16	0.040 +/- 0.11
\mathcal{M}_2	0.754 +/- 0.07	0.432 +/- 0.18	0.226 +/- 0.19
\mathcal{M}_3	0.824 +/- 0.06	0.650 +/- 0.09	0.405 +/- 0.20
\mathcal{M}_4	0.811 +/- 0.08	0.671 +/- 0.11	0.553 +/- 0.12

Table 5.1 BLEU-2,3,4 scores (mean and std dev) on a sample of 1M sentences drawn from the corpus used to train $\mathcal{M}_1, \mathcal{M}_2, \mathcal{M}_3$ and \mathcal{M}_4 compared to \mathcal{M}_N .

5.2.2 Corpus word order bootstrap resample

The above permutations preserve higher order distributional information by keeping words from the same sentence together. However, we need a baseline to under-

Algorithm 1 SentenceRandomizer

```

1: procedure  $\mathcal{F}(S, t, n)$                                  $\triangleright$  Randomize a sentence  $S$  with seed  $t$  and n grams  $n$ 
2:    $W$  = tokenize the words in  $S$ 
3:   Set the seed to  $t$ 
4:   if  $n > 1$  then
5:     while True do
6:        $K$  = Sample all possible starting points from  $[0, |W| - n]$ 
7:       Ignore the starting points in  $K$  which overlap with conjoined tokens       $\triangleright$ 
      Conjoined tokens consists of joined unigrams
8:       if  $|K| \geq 1$  then
9:         Sample one position  $p \in K$ 
10:         $g$  = Extract the n-gram  $W[p : p + n]$ 
11:        Delete  $W[p + 1 : p + n]$ 
12:         $W[p] =$  Convert  $g$  to a conjoined token
13:       else
14:         Break from While loop
15:     while True do
16:        $\hat{W}$  = randomly shuffle tokens in  $W$ 
17:        $r = \sum(\hat{W}[i] = W[i])$      $\triangleright$  Count number of positions where the token remains in its
      original position
18:       if  $r = 0$  then Break out of While loop
19:      $\hat{S} =$  join the tokens in  $\hat{W}$ 
20:   Return  $\hat{S}$ 

```

stand how a model would perform without such co-occurrence information. We construct a baseline, \mathcal{M}_{UG} , that captures word/subword information, without access to co-occurrence statistics. To construct \mathcal{M}_{UG} , we sample unigrams from BookWiki according to their frequencies, while also treating named entities as unigrams. We leverage Spacy [Honnibal et al., 2020a]¹ to extract unigrams and named entities from the corpus, and construct \mathcal{M}_{UG} by drawing words from this set according to their frequency. This allows us to construct \mathcal{M}_{UG} such that it has exactly the same size as BookWiki but without any distributional (i.e. co-occurrence) information beyond the unigram frequency distribution. Our hypothesis is that any model pre-trained on this data will perform poorly, but it should provide a baseline for the limits on learning language of the inductive bias of the model in isolation.

5.2.3 Further baselines

To investigate what happens if a model has absolutely no notion of word order, we also experiment with pre-training RoBERTa on the original corpus without positional embeddings. Concretely, we modify the RoBERTa architecture to remove the positional embeddings from the computation graph, and then proceed to pre-train on the natural order BookWiki corpus. We denote this model \mathcal{M}_{NP} . Finally, we consider a randomly initialized RoBERTa model \mathcal{M}_{RI} to observe the extent we can learn from each task with only the model’s base inductive bias.

5.3 Evaluated Models & Tasks

We use the RoBERTa (base) Liu et al. [2019b] MLM architecture, due to its relative computational efficiency and good downstream task performance. We expect that other variants of MLMs would provide similar insights, given their similar characteristics.

¹<https://spacy.io/>

In all of our experiments, we use the original 16GB BookWiki corpus.² We denote the model trained on the original, un-modified BookWiki corpus as \mathcal{M}_N (for “natural”). We use two types of word order randomization methods: permuting words at the sentence level, and resampling words at the corpus level.

5.3.1 Pre-training details

Each model $\in \{\mathcal{M}_N, \mathcal{M}_1, \mathcal{M}_2, \mathcal{M}_3, \mathcal{M}_4, \mathcal{M}_{UG}, \mathcal{M}_{NP}\}$ is a RoBERTa-base model (12 layers, hidden size of 768, 12 attention heads, 125M parameters), trained for 100k updates using 8k batch-size, 20k warmup steps, and 0.0006 peak learning rate. These are identical hyperparameters to Liu et al. [2019b], except for the number of warmup steps which we changed to 20k for improved training stability. Each model was trained using 64 GPUs for up to 72 hours each. We train three seeds for each data configuration. We use FairSeq Ott et al. [2019] for the pre-training and fine-tuning experiments. We use the Wiki 103 validation and test set to validate and test the array of pre-trained models, as validation on this small dataset is quick, effective, and reproducible for comparison among publicly available datasets (Figure 5.1). We observe that perplexity monotonically increases from \mathcal{M}_N , through $\mathcal{M}_4 - \mathcal{M}_1$, to \mathcal{M}_{UG} , and finally \mathcal{M}_{NP} .

5.3.2 Fine-tuning tasks

We evaluate downstream performance using the General Language Understanding and Evaluation (GLUE) benchmark, the Paraphrase Adversaries from Word Scrambling (PAWS) dataset, and various parametric and non-parametric tasks (see ??).

GLUE. The GLUE Wang et al. [2018] benchmark is a collection of 9 datasets for evaluating natural language understanding systems, of which we use Corpus of Linguistic Acceptability [CoLA, Warstadt et al., 2019b], Stanford Sentiment Treebank [SST, Socher

²We release the pre-trained RoBERTa models used in our experiments through the FairSeq repository: https://github.com/pytorch/fairseq/tree/master/examples/shuffled_word_order.

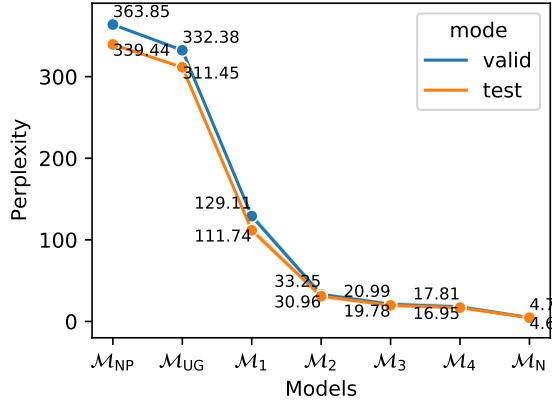


Figure 5.1 Perplexity of various models on Wiki 103 valid and test sets.

et al., 2013], Microsoft Research Paragraph Corpus [MRPC, Dolan and Brockett, 2005], Quora Question Pairs (QQP)³, Multi-Genre NLI [MNLI, Williams et al., 2018c], Question NLI [QNLI, Rajpurkar et al., 2016b, Demszky et al., 2018], Recognizing Textual Entailment [RTE, Dagan et al., 2005b, Haim et al., 2006, Giampiccolo et al., 2007, Bentivogli et al., 2009]. Pham et al. [2020b] show the word order insensitivity of several GLUE tasks (QQP, SST-2), evaluated on public regularly pre-trained checkpoints.

PAWS. The PAWS task Zhang et al. [2019] consists of predicting whether a given pair of sentences are paraphrases. This dataset contains both paraphrase and non-paraphrase pairs with high lexical overlap, which are generated by controlled word swapping and back translation. Since even a small word swap and perturbation can drastically modify the meaning of the sentence, we hypothesize the randomized pre-trained models will struggle to attain a high performance on PAWS.

Fine-tuning details. We use the same fine-tuning methodology used by Liu et al. [2019b], where we run hyperparameter search over the learning rates $\{1 \times 10^{-5}, 2 \times 10^{-5}, 3 \times 10^{-5}\}$ and batch sizes $\{16, 32\}$ for each model. For the best hyperparam con-

³<http://data.quora.com/First-Quora-Dataset-Release-Question-Pairs>

Model	QNLI	RTE	QQP	SST-2	MRPC	PAWS	MNLI-m/mm	CoLA
\mathcal{M}_N	92.45 +/- 0.2	73.62 +/- 3.1	91.25 +/- 0.1	93.75 +/- 0.4	89.09 +/- 0.9	94.49 +/- 0.2	86.08 +/- 0.2 / 85.4 +/- 0.2	52.45 +/- 21
\mathcal{M}_4	91.65 +/- 0.1	70.94 +/- 1.2	91.39 +/- 0.1	92.46 +/- 0.3	86.90 +/- 0.3	94.26 +/- 0.2	83.79 +/- 0.2 / 83.94 +/- 0.3	35.25 +/- 32
\mathcal{M}_3	91.56 +/- 0.4	69.75 +/- 2.8	91.22 +/- 0.1	91.97 +/- 0.5	86.22 +/- 0.8	94.03 +/- 0.1	83.83 +/- 0.2 / 83.71 +/- 0.1	40.78 +/- 23
\mathcal{M}_2	90.51 +/- 0.1	70.00 +/- 2.5	91.33 +/- 0.0	91.78 +/- 0.3	85.90 +/- 1.2	93.53 +/- 0.3	83.45 +/- 0.3 / 83.54 +/- 0.3	50.83 +/- 5.8
\mathcal{M}_1	89.05 +/- 0.2	68.48 +/- 2.5	91.01 +/- 0.0	90.41 +/- 0.4	86.06 +/- 0.8	89.69 +/- 0.6	82.64 +/- 0.1 / 82.67 +/- 0.2	31.08 +/- 10
\mathcal{M}_{NP}	77.59 +/- 0.3	54.78 +/- 2.2	87.78 +/- 0.4	83.21 +/- 0.6	72.78 +/- 1.6	57.22 +/- 1.2	63.35 +/- 0.4 / 63.63 +/- 0.2	2.37 +/- 3.2
\mathcal{M}_{UG}	66.94 +/- 9.2	53.70 +/- 1.0	85.57 +/- 0.1	83.17 +/- 1.5	70.57 +/- 0.7	58.59 +/- 0.3	71.93 +/- 0.2 / 71.33 +/- 0.5	0.92 +/- 2.1
\mathcal{M}_{RI}	62.17 +/- 0.4	52.97 +/- 0.2	81.53 +/- 0.2	82.0 +/- 0.7	70.32 +/- 1.5	56.62 +/- 0.0	65.70 +/- 0.2 / 65.75 +/- 0.3	8.06 +/- 1.6

Table 5.2 GLUE and PAWS-Wiki dev set results on different RoBERTa (base) models trained on variants of the BookWiki corpus (with mean and std). The top row is the original model, the middle half contains our primary models under investigation, and the bottom half contains the baselines.

figurations of each model, we fine-tune with 5 different seeds and report the mean and standard deviation for each setting. \mathcal{M}_{NP} is fine-tuned without positional embeddings, matching the way it was pre-trained.

5.4 Results

5.4.1 Downstream task results

In this section, we present the downstream task performance of the models defined in §5.2. For evaluation, we report Matthews correlation for CoLA and accuracy for all other tasks.

Word order permuted pre-training

In our first set of experiments, we finetune the pre-trained models on the GLUE and PAWS tasks. We report the results in Table 5.2.⁴ First, we observe that the model without access to distributional or word order information, \mathcal{M}_{UG} (unigram) performs

⁴The \mathcal{M}_N results are not directly comparable with that of publicly released `roberta-base` model by Liu et al. [2019b], as that uses the significantly larger 160GB corpus, and is trained for 500K updates. For computational reasons, we restrict our experiments to the 16GB BookWiki corpus and 100K updates, mirroring the RoBERTa ablations.

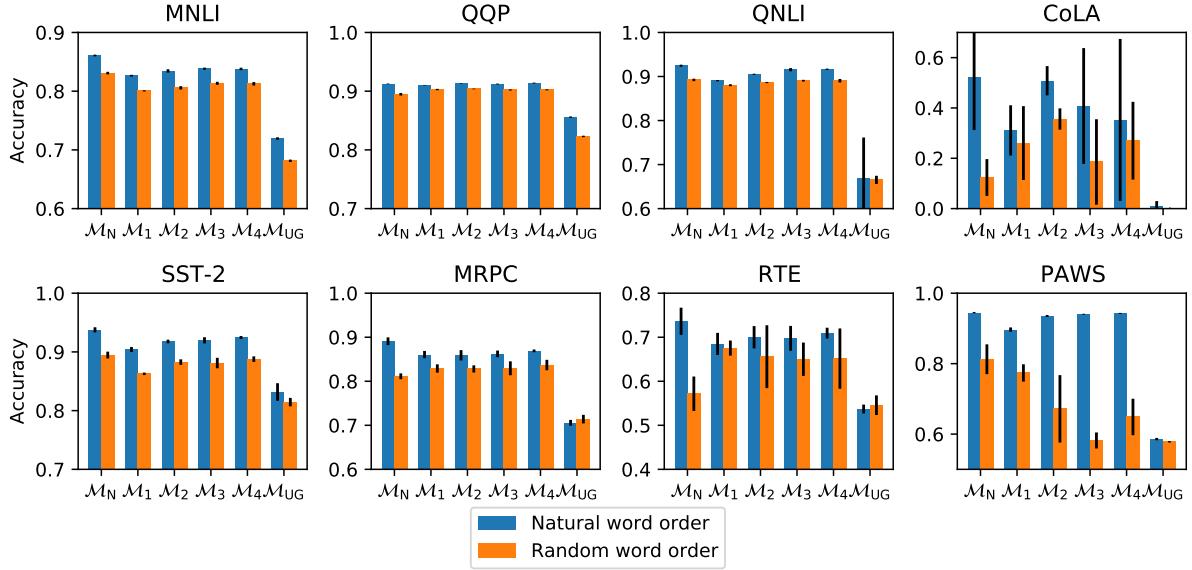


Figure 5.2 GLUE & PAWS task dev performance when finetuned on naturally (blue) and randomly ordered (orange) text, respectively, using pre-trained RoBERTa (base) models trained on different versions of BookWiki corpus.

much worse than M_N overall: M_{UG} is 18 points worse than M_N on average across the accuracy-based tasks in Table 5.2 and has essentially no correlation with human judgments on CoLA. M_{UG} , M_{NP} and M_{RI} perform comparably on most of the tasks, while achieving surprisingly high scores in QQP and SST-2. However, all three models perform significantly worse on GLUE and PAWS, compared to M_N (Table 5.2, bottom half). M_{UG} reaches up to 71.9 on MNLI - possibly due to the fact that M_{UG} has access to (bags of) words and some phrases (from NER) is beneficial for MNLI. For the majority of tasks, the difference between M_{NP} and M_{RI} is small - a pure bag of words model performs comparably to a randomly initialized model.

Next, we observe a significant improvement on all tasks when we give models access to sentence-level distributional information during pre-training. M_1 , the model pre-trained on completely shuffled sentences, is on average only 3.3 points lower than

\mathcal{M}_N on the accuracy-based tasks, and within 0.3 points of \mathcal{M}_N on QQP. Even on PAWS, which was designed to require knowledge of word order, \mathcal{M}_1 is within 5 points of \mathcal{M}_N . Randomizing n -grams instead of words during pre-training results in a (mostly) smooth increase on these tasks: \mathcal{M}_4 , the model pre-trained on shuffled 4-grams, trails \mathcal{M}_N by only 1.3 points on average, and even comes within 0.2 points of \mathcal{M}_N on PAWS. We observe a somewhat different pattern on CoLA, where \mathcal{M}_2 does almost as well as \mathcal{M}_N and outperforms \mathcal{M}_3 and \mathcal{M}_4 , though we also observe very high variance across random seeds for this task. Crucially, we observe that \mathcal{M}_1 outperforms \mathcal{M}_{NP} by a large margin. This shows that positional embeddings are critical for learning, even when the word orders themselves are not natural. Recall, \mathcal{M}_{NP} is fed natural sentences as \mathcal{M}_N while not having the ability to learn positional embeddings. To further quantify the effect of positional embeddings, we also investigated the effect of shuffling the entire context window, to keep the co-occurrence information same as \mathcal{M}_{NP} in §5.5.1. We observed this model to be worse than \mathcal{M}_1 but significantly better than \mathcal{M}_{NP} to support the claim about the importance of positional embeddings while training.

Overall, these results confirm our hypothesis that RoBERTa’s strong performance on downstream tasks can be explained for a large part by the distributional prior.

Word order permuted fine-tuning

There are two possible explanations for the results in §5.4.1: either the tasks do not need word order information to be solved, or any necessary word order information can be acquired during fine-tuning. To examine this question, we permute the word order during fine-tuning as well. Concretely, for each task, we construct a unigram order-randomized version of each example in the fine-tuning training set using \mathcal{F}_1 . We then fine-tune our pre-trained models on this shuffled data and evaluate task performance. For all experiments, we evaluate and perform early stopping on the original, natural word order dev set, in order to conduct a fair evaluation on the exact same optimization

setup for all models.

Our results in Figure 5.2 provide some evidence for both hypotheses. On QQP and QNLI, accuracy decreases only slightly for models fine-tuned on shuffled data. Models can also achieve above 80% accuracy on MNLI, SST-2, and MRPC when fine-tuned on shuffled data, suggesting that purely lexical information is quite useful on its own.

On the other hand, for all datasets besides QQP and QNLI, we see noticeable drops in accuracy when fine-tuning on shuffled data and testing on normal order, both for \mathcal{M}_N and for shuffled models \mathcal{M}_1 through \mathcal{M}_4 . This suggests both that word order information is useful for these tasks, and that shuffled models must be learning to use word order information during fine-tuning.⁵ Having word order during fine-tuning is especially important for achieving high accuracy on CoLA, RTE (cf. Pham et al. 2020b), as well as PAWS, suggesting that these tasks are the most word order reliant. Recent research Yu and Ettinger [2021] raised some questions about potential artefacts inflating performance on PAWS: their swapping-distance cue of appears consistent both with our finding of high PAWS performance for n-gram shuffled models in Table 5.2, and with our PAWS results in Figure 5.2, which suggests that PAWS performance does in fact rely to some extent on natural word order at the fine-tuning stage.

Finally, for CoLA, MRPC, and RTE, performance is higher after fine-tuning on shuffled data for \mathcal{M}_1 than \mathcal{M}_N . We hypothesize that \mathcal{M}_N represents shuffled and non-shuffled sentences very differently, resulting in a domain mismatch problem when fine-tuning on shuffled data but evaluating on non-shuffled data.⁶ Since \mathcal{M}_1 never learns to be sensitive to word order during pre-training or fine-tuning, it does not suffer from that issue. Our results in this section also highlights the issues with these

⁵We perform additional experiments on how the model representations change during fine-tuning for shuffled training using Risannen Data Analysis in §5.5.6.

⁶We further study the domain mismatch problem by evaluating on shuffled data *after* fine-tuning on the shuffled data for models in ???. We observe that models improves their scores on evaluation on shuffled data when the training data source is changed from natural to shuffled - highlighting domain match effect.

datasets, concurrent to the findings that many GLUE tasks does not need sophisticated linguistic knowledge to solve, as models typically tend to exploit the statistical artefacts and spurious correlations during fine-tuning (cf. Gururangan et al. 2018a, Poliak et al. 2018, Tsuchiya 2018, McCoy et al. 2019). However, our results overwhelmingly support the fact that word order does not matter during pre-training, if the model has the opportunity to learn the necessary information about word order during fine-tuning.

5.4.2 Probing results

To investigate how much syntactic information is contained in the MLM representations, we evaluate several probing tasks on our trained models. We consider two classes of probes: *parametric* probes, which make use of learnable parameters, and *non-parametric* probes, which directly examine the language model’s predictions.

Parametric Probing

To probe our models for syntactic, semantic and other linguistic properties, we investigate dependency parsing using Pareto probing Pimentel et al. [2020a] and the probing tasks from Conneau et al. [2018] in SentEval Conneau and Kiela [2018].

Pimentel et al. [2020a] proposed a framework based on Pareto optimality to probe for syntactic information in contextual representations. They suggest that an optimal probe should balance optimal performance on the probing task with the complexity of the probe. Following their setup, we use the “difficult” probe: dependency parsing (DEP). We also investigate the “easy” probes, dependency arc labeling (DAL) and POS tag prediction (POS), results are reported §5.5.8. We probe with Linear and MLP probes, and inspect the task accuracy in terms of Unlabeled Attachment Score (UAS). The dependency parsing probe used in Pimentel et al. [2020a] builds on the Biaffine Dependency Parser [Dozat and Manning, 2017], but with simple MLPs on top of the

Transformer representations.⁷

Training setup. Similar to the setup by Pimentel et al. [2020a], we run 50 random hyperparameter searches on both MLP and Linear probes by uniformly sampling from the number of layers (0-5), dropout (0-0.5), log-uniform hidden size $[2^5, 2^{10}]$. We triple this experiment size by evaluating on three pre-trained models of different seeds for each model configuration. We consider Pimentel et al.’s English dataset, derived from Universal Dependencies EWT (UD EWT) Bies et al. [2012], Silveira et al. [2014] which contains 12,543 training sentences. Additionally, we experiment on the Penn Treebank dataset (PTB), which contains 39,832 training sentences.⁸ We report the mean test accuracy over three seeds for the best dev set accuracy for each task.⁹

Model	UD EWT		PTB	
	MLP	Linear	MLP	Linear
\mathcal{M}_N	80.41 +/- 0.85	66.26 +/- 1.59	86.99 +/- 1.49	66.47 +/- 2.77
\mathcal{M}_4	78.04 +/- 2.06	65.61 +/- 1.99	85.62 +/- 1.09	66.49 +/- 2.02
\mathcal{M}_3	77.80 +/- 3.09	64.89 +/- 2.63	85.89 +/- 1.01	66.11 +/- 1.68
\mathcal{M}_2	78.22 +/- 0.88	64.96 +/- 2.32	84.72 +/- 0.55	64.69 +/- 2.50
\mathcal{M}_1	69.26 +/- 6.00	56.24 +/- 5.05	79.43 +/- 0.96	57.20 +/- 2.76
\mathcal{M}_{UG}	74.15 +/- 0.93	65.69 +/- 7.35	80.07 +/- 0.79	57.28 +/- 1.42

Table 5.3 Unlabeled Attachment Score (UAS) (mean and std) on the dependency parsing task (DEP) on two datasets, UD EWT and PTB, using the Pareto Probing framework Pimentel et al. [2020a].

Results. We observe that the UAS scores follow a similar linear trend as the fine-tuning results in that $\mathcal{M}_1 \approx \mathcal{M}_{UG} < \mathcal{M}_2 < \mathcal{M}_3 < \mathcal{M}_4 < \mathcal{M}_N$ (Table 5.3). Surprisingly, \mathcal{M}_{UG} probing scores seem to be somewhat better than \mathcal{M}_1 (though with large overlap

⁷We experimented with a much stronger, state-of-the-art Second order Tree CRF Neural Dependency Parser Zhang et al. [2020], but did not observe any difference in UAS with different pre-trained models (see §5.5.4)

⁸PTB data [Kitaev et al., 2019] is used from github.com/nikitakit/self-attentive-parser/tree/master/data.

⁹Pimentel et al. [2020a] propose computing the *Pareto Hypervolume* over all hyperparameters in each task. We did not observe a significant difference in the hypervolumes for the models, as reported in ??.

Model	Length (Surface)	WordContent (Surface)	TreeDepth (Syntactic)	TopConstituents (Syntactic)	BigramShift (Syntactic)	Tense (Semantic)	SubjNumber (Semantic)	ObjNumber (Semantic)	OddManOut (Semantic)	CoordInversion (Semantic)
M_N	78.92 +/- 1.91	31.83 +/- 1.75	35.97 +/- 1.38	78.26 +/- 4.08	81.82 +/- 0.55	87.83 +/- 0.51	85.05 +/- 1.23	75.94 +/- 0.68	58.40 +/- 0.33	70.87 +/- 2.46
M_4	92.88 +/- 0.15	57.78 +/- 0.36	40.05 +/- 0.29	72.50 +/- 0.51	76.12 +/- 0.29	88.32 +/- 0.13	85.65 +/- 0.13	82.95 +/- 0.05	58.89 +/- 0.30	61.31 +/- 0.19
M_3	91.52 +/- 0.16	48.81 +/- 0.26	38.63 +/- 0.61	70.29 +/- 0.31	77.36 +/- 0.12	86.74 +/- 0.12	83.83 +/- 0.38	80.99 +/- 0.26	57.01 +/- 0.21	60.00 +/- 0.26
M_2	93.54 +/- 0.29	62.52 +/- 0.21	41.40 +/- 0.32	74.31 +/- 0.29	75.44 +/- 0.14	87.91 +/- 0.35	84.88 +/- 0.11	83.98 +/- 0.14	57.60 +/- 0.36	59.46 +/- 0.37
M_1	88.33 +/- 0.14	64.03 +/- 0.34	40.24 +/- 0.20	70.94 +/- 0.38	58.37 +/- 0.40	87.88 +/- 0.08	83.49 +/- 0.12	83.44 +/- 0.06	56.51 +/- 0.26	56.98 +/- 0.50
M_{UG}	86.69 +/- 0.33	36.60 +/- 0.33	32.53 +/- 0.76	61.54 +/- 0.60	57.42 +/- 0.04	68.45 +/- 0.23	71.25 +/- 0.12	66.63 +/- 0.21	50.06 +/- 0.40	56.26 +/- 0.17

Table 5.4 SentEval Probing Conneau et al. [2018], Conneau and Kiela [2018] results (with mean and std) on different model variants.

in their standard deviations), even though M_{UG} cannot learn information related to either word order or co-occurrence patterns. The performance gap appears to be task- and probe specific. We observe a low performance gap in several scenarios, the lowest being between M_N vs. M_3 / M_4 , for PTB using the both MLP and Linear probes.

SentEval Probes

We also investigate the suite of 10 probing tasks Conneau et al. [2018] available in the SentEval toolkit Conneau and Kiela [2018]. This suite contains a range of semantic, syntactic and surface level tasks. Jawahar et al. [2019b] utilize this set of probing tasks to arrive at the conclusion that “*BERT embeds a rich hierarchy of linguistic signals: surface information at the bottom, syntactic information in the middle, semantic information at the top*”. We re-examine this hypothesis by using the same probing method and comparing against models trained with random word order.

Training setup. We run the probes on the final layer of each of our pre-trained models for three seeds, while keeping the encoder frozen. SentEval trains probes on top of fixed representations individually for each task. We follow the recommended setup and run grid search over the following hyperparams: number of hidden layer dimensions ([0, 50, 100, 200]), dropout ([0, 0.1, 0.2]), 4 epochs, 64 batch size. We select the best performance based on the dev set, and report the test set accuracy.

Results. We provide the results in Table 5.4. The \mathcal{M}_N pre-trained model scores better than the unnatural word order models for only one out of five semantic tasks and in none of the lexical tasks. However, \mathcal{M}_N does score higher for two out of three syntactic tasks. Even for these two syntactic tasks, the gap among \mathcal{M}_{UG} and \mathcal{M}_N is much higher than \mathcal{M}_1 and \mathcal{M}_N . These results show that while natural word order is useful for at least some probing tasks, the distributional prior of randomized models alone is enough to achieve a reasonably high accuracy on syntax sensitive probing.

Non-Parametric Probing

How to probe effectively with parametric probes is a matter of much recent debate [Hall Maudslay et al., 2020, Belinkov, 2021]. From our results so far, it is unclear whether parametric probing meaningfully distinguishes models trained with corrupted word order from those trained with normal orders. Thus, we also investigate non-parametric probes Linzen et al. [2016], Marvin and Linzen [2018], Gulordava et al. [2018b] using the formulation of Goldberg [2019b] and Wolf [2019a].

We consider a set of non-parametric probes that use a range of sentences varying in their linguistic properties. For each, the objective is for a pre-trained model to provide higher probability to a grammatically correct word than to an incorrect one. Since both the correct and incorrect options occupy the same sentential position, we call them “focus words”. Linzen et al. [2016] use sentences from Wikipedia containing present-tense verbs, and compare the probability assigned by the encoder to plural vs. singular forms of the verb; they focus on sentences containing at least one noun between the verb and its subject, known as “agreement attractors.” Gulordava et al. [2018b] instead replace focus words with random substitutes from the same part-of-speech and inflection. Finally, Marvin and Linzen [2018] construct minimal pairs of grammatical and ungrammatical sentences, and compare the model’s probability for the words that differ.

Setup. In our experiments, we mask the focus words in the stimuli and compute the probability of the correct and incorrect token respectively. To handle Byte-Pair Encoding (BPE), we use the WordPiece Wu et al. [2016] tokens prepended with a space. We observe that the Linzen et al. [2016] and Gulordava et al. [2018b] datasets are skewed towards singular focus words, which could disproportionately help weaker models that just happen to assign more probability mass to singular focus words. To counter this, we balance these datasets to have an equal number of singular and plural focus words by upsampling, and report the aggregated and balanced results in Table 5.5 (see §5.5.9 for more detailed results). We verify our experiments by using three pre-trained models with different seeds for each model configuration.

Results. We observe for the Linzen et al. [2016] and Marvin and Linzen [2018] datasets that the gap between the \mathcal{M}_N and randomization models is relatively large. The Gulordava et al. [2018b] dataset shows a smaller gap between \mathcal{M}_N and the randomization models. While some randomization models (e.g., \mathcal{M}_2 , \mathcal{M}_3 , and \mathcal{M}_4) performed quite similarly to \mathcal{M}_N according to the parametric probes, they all are markedly worse than \mathcal{M}_N according to the non-parametric ones. This suggests that non-parametric probes identify certain syntax-related modeling failures that parametric ones do not.

Model	Linzen et al. [2016] *	Gulordava et al. [2018b] *	Marvin and Linzen [2018]
\mathcal{M}_N	91.17 +/- 2.6	68.66 +/- 11.6	88.05 +/- 6.5
\mathcal{M}_4	66.93 +/- 3.2	69.47 +/- 4.9	70.66 +/- 12.5
\mathcal{M}_3	64.60 +/- 2.7	66.10 +/- 5.9	73.82 +/- 15.7
\mathcal{M}_2	61.27 +/- 3.1	60.20 +/- 7.6	73.95 +/- 14.3
\mathcal{M}_1	58.96 +/- 1.8	68.10 +/- 14.4	70.69 +/- 11.6
\mathcal{M}_{UG}	65.36 +/- 7.1	60.88 +/- 24.3	50.10 +/- 0.2

Table 5.5 Mean (and std) non-parametric probing accuracy on different datasets. * indicates rebalanced datasets, see ?? for more details.

Model	QNLI	RTE	QQP	SST-2	MRPC	PAWS	MNLI-m/mm	CoLA
\mathcal{M}_N	92.45 +/- 0.2	73.62 +/- 3.1	91.25 +/- 0.1	93.75 +/- 0.4	89.09 +/- 0.9	94.49 +/- 0.2	86.08 +/- 0.2 / 85.4 +/- 0.2	52.45 +/- 21.2
\mathcal{M}_4	91.65 +/- 0.1	70.94 +/- 1.2	91.39 +/- 0.1	92.46 +/- 0.3	86.90 +/- 0.3	94.26 +/- 0.2	83.79 +/- 0.2 / 83.94 +/- 0.3	35.25 +/- 32.2
\mathcal{M}_3	91.56 +/- 0.4	69.75 +/- 2.8	91.22 +/- 0.1	91.97 +/- 0.5	86.22 +/- 0.8	94.03 +/- 0.1	83.83 +/- 0.2 / 83.71 +/- 0.1	40.78 +/- 23.0
\mathcal{M}_2	90.51 +/- 0.1	70.00 +/- 2.5	91.33 +/- 0.0	91.78 +/- 0.3	85.90 +/- 1.2	93.53 +/- 0.3	83.45 +/- 0.3 / 83.54 +/- 0.3	50.83 +/- 5.80
\mathcal{M}_1	89.05 +/- 0.2	68.48 +/- 2.5	91.01 +/- 0.0	90.41 +/- 0.4	86.06 +/- 0.8	89.69 +/- 0.6	82.64 +/- 0.1 / 82.67 +/- 0.2	31.08 +/- 10.0
\mathcal{M}_{512}	84.97 +/- 0.3	56.09 +/- 0.6	90.15 +/- 0.1	86.11 +/- 0.7	79.41 +/- 0.6	77.3 +/- 12.63	77.58 +/- 0.3 / 77.89 +/- 0.4	12.54 +/- 5.57
\mathcal{M}_{NP}	77.59 +/- 0.3	54.78 +/- 2.2	87.78 +/- 0.4	83.21 +/- 0.6	72.78 +/- 1.6	57.22 +/- 1.2	63.35 +/- 0.4 / 63.63 +/- 0.2	2.37 +/- 3.20
\mathcal{M}_{UF}	77.69 +/- 0.4	53.84 +/- 0.6	85.92 +/- 0.1	84.00 +/- 0.6	71.35 +/- 0.8	58.43 +/- 0.3	72.10 +/- 0.4 / 72.58 +/- 0.4	8.89 +/- 1.40
\mathcal{M}_{UG}	66.94 +/- 9.2	53.70 +/- 1.0	85.57 +/- 0.1	83.17 +/- 1.5	70.57 +/- 0.7	58.59 +/- 0.3	71.93 +/- 0.2 / 71.33 +/- 0.5	0.92 +/- 2.10
\mathcal{M}_{RI}	62.17 +/- 0.4	52.97 +/- 0.2	81.53 +/- 0.2	82.0 +/- 0.7	70.32 +/- 1.5	56.62 +/- 0.0	65.70 +/- 0.2 / 65.75 +/- 0.3	8.06 +/- 1.60

Table 5.6 GLUE and PAWS-Wiki dev set results on different ablations of the RoBERTa (base) models, trained on variants of the BookWiki corpus (with mean and std dev). The top row is the original model, the middle half contains the sentence randomization models, and the bottom half contains the ablations.

5.5 Analysis

5.5.1 Word-order pre-training ablations

We also train further model ablations with low to high distributional priors. Following the construction of the corpus bootstrap resample, we train a model where words are drawn uniformly from BookWiki corpus, thus destroying the natural frequency distribution (\mathcal{M}_{UF}). We further study an ablation for a high distributional prior, \mathcal{M}_{512} , where we shuffle words (unigram) in a buffer created with joining multiple sentences such that maximum token length of the buffer is 512. This ablation—which is similar to the paragraph word shuffle condition in Gauthier and Levy [2019]—will allow us to study the effect of unigram shuffling in a window larger than the one for \mathcal{M}_1 . Buffer size is chosen to be 512 because BERT/RoBERTa is typically trained with that maximum sequence length.

We observe dev set results on the GLUE benchmark of these ablations, along with baselines \mathcal{M}_{UG} , \mathcal{M}_{RI} and \mathcal{M}_{NP} and random shuffles in Table 5.6 and Figure 5.3. We observe that \mathcal{M}_{512} exhibits worse overall scores than \mathcal{M}_1 , however it is still signif-

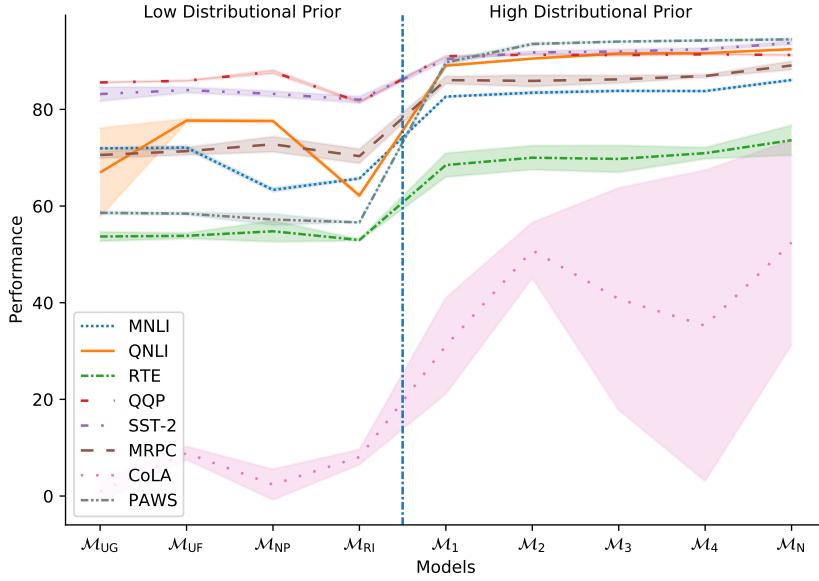


Figure 5.3 GLUE results on various model ablations using BookWiki corpus.

icantly better than \mathcal{M}_{NP} or \mathcal{M}_{UG} baselines. We observe that destroying the natural frequency distribution of words (\mathcal{M}_{UF}) yields comparable or slightly better results compared to random corpus model \mathcal{M}_{UG} . This result shows that merely replicating the natural distribution of words without any context is not useful for the model to learn. These results indicate that at least some form of distributional prior is required for MLM-based models to learn a good downstream representation.

One might argue that the superior results displayed by the unnatural models is due to the ability of RoBERTa to “reconstruct” the natural word order from shuffled sentences. The data generation algorithm, \mathcal{F}_i requires a seed t for every sentence. In our experiments, we had set the same seed for every sentence in the corpus to ensure reproducibility. However, it could be problematic if the sentences of the same length are permuted with the same seed, which could be easier for the model to “reconstruct” the natural word order to learn the necessary syntax. We tested this hypothesis by

Model	RTE	MRPC	SST-2	CoLA	QQP	QNLI	MNLI	PAWS
\mathcal{M}_1	68.48	85.97	90.41	31.07	91.01	89.05	82.64	89.69
\mathcal{M}_1^*	68.41	85.75	90.17	50.14	91.02	89.50	82.92	91.99

Table 5.7 Reconstruction experiments on shuffled word order sentences by fixing the same seed for every sentence (\mathcal{M}_1) and having different seed for different shards of the corpus (\mathcal{M}_1^*). We observe minimal difference in the downstream GLUE and PAWS scores.

constructing a new corpus with different seeds for every sentence in every shard in the corpus (1/5th of BookWiki corpus is typically referred to as a *shard* for computational purposes), to build the model \mathcal{M}_1^* . We observe that there is minimal difference in the raw numbers among \mathcal{M}_1 and \mathcal{M}_1^* for most of the tasks (Table 5.7) (with the exception of CoLA which performs similar to \mathcal{M}_2 possibly due to a difference in initialization). This result consequently proves that even with same seed, it is difficult for the model to just reconstruct the unnatural sentences during pre-training.

5.5.2 Measuring Relative difference

In this section, we further measure the difference in downstream task performance reported in §5.4.1 using as a metric the *relative difference*. Let us denote the downstream task performance as $\mathcal{A}(\mathcal{T}|D)$, where \mathcal{T} is the task and D is the pre-trained model. We primarily aim to evaluate the relative performance gap, i.e. how much the performance differs between our natural and unnatural models. Thus, we define the *Relative Difference* ($\Delta_{\{D\}}(\mathcal{T})$):

$$\Delta_{\{D\}}(\mathcal{T}) = \frac{\mathcal{A}(\mathcal{T}|OR) - \mathcal{A}(\mathcal{T}|D))}{\mathcal{A}(\mathcal{T}|OR) - \mathcal{A}(\mathcal{T}|\emptyset)}, \quad (5.1)$$

where $\mathcal{A}(\mathcal{T}|\emptyset)$ is the random performance on the task \mathcal{T} (0.33 for MNLI, 0 for CoLA, and 0.5 for rest) $\Delta_{\{D\}}(\mathcal{T}) \rightarrow 0$ when the performance of a pre-trained model reaches that

Model	QNLI	RTE	QQP	SST-2	MRPC	CoLA	PAWS	MNLI
\mathcal{M}_1	3.70	7.04	0.26	3.58	3.42	40.74	5.12	3.62
\mathcal{M}_2	2.11	4.95	-0.09	2.12	3.61	3.09	9.06	2.63
\mathcal{M}_3	0.97	5.30	0.03	1.91	3.24	22.25	0.49	2.31
\mathcal{M}_4	0.87	3.67	-0.15	1.39	2.47	32.79	0.25	2.19
\mathcal{M}_{UG}	27.74	27.25	6.26	11.35	20.91	98.24	38.20	16.56
\mathcal{M}_{NP}	16.16	25.77	3.83	11.30	18.42	95.48	39.66	26.10

Table 5.8 $\Delta_{\{D_i\}}(\mathcal{T})$, scaled by a factor of 100 for GLUE and PAWS tasks.

of the pre-trained model trained with natural word order.

We observe the relative difference on the tasks in Table 5.8. CoLA has the largest $\Delta_{\{D\}}(\mathcal{T})$ among all tasks, suggesting the expected high word order reliance. $\Delta_{\{D\}}(\mathcal{T})$ is lowest for QQP.

5.5.3 Fine-tuning with randomized data

We perform additional experiments using the fine-tuned models from §5.4.1. Specifically, we construct unigram randomized train and test sets (denoted as *shuffled*) of a subset of tasks to evaluate whether models fine-tuned on natural or unnatural task data (having natural or unnatural pre-training prior) are able to understand unnatural data during testing. In the previous chapter (§4.4.1), we observed for MNLI there exists at least one permutation for many examples which can be predicted correctly by the model. However, we also observed that every sentence can have many permutations which cannot be predicted correctly as well (§4.7). Similarly in this section we construct 100 permutations for each example in the dev set for each task to capture the overall accuracy.

Concretely, we use \mathcal{M}_N , \mathcal{M}_1 and \mathcal{M}_{UG} as our pre-trained representations (trained with natural, unigram sentence shuffle and corpus shuffle data respectively) and eval-

name	fine-tune-train	fine-tune-eval	MNLI	QNLI	RTE	CoLA	MRPC	SST-2	PAWS
\mathcal{M}_N	natural	natural	86.08 +/- 0.15	92.45 +/- 0.24	73.62 +/- 3.09	52.44 +/- 21.22	89.09 +/- 0.88	93.75 +/- 0.44	94.49 +/- 0.18
	natural	shuffled	68.11 +/- 0.52	81.08 +/- 0.38	56.72 +/- 3.29	4.77 +/- 1.82	75.94 +/- 1.01	80.78 +/- 0.37	62.22 +/- 0.09
	shuffled	natural	82.99 +/- 0.16	89.32 +/- 0.23	57.9 +/- 4.71	0.0 +/- 0.0	79.71 +/- 2.57	89.12 +/- 0.5	72.03 +/- 13.79
	shuffled	shuffled	79.96 +/- 0.1	87.51 +/- 0.09	59.07 +/- 3.2	1.4 +/- 2.17	79.17 +/- 0.35	86.11 +/- 0.5	65.15 +/- 0.48
\mathcal{M}_1	natural	natural	82.64 +/- 0.15	89.05 +/- 0.15	68.48 +/- 2.51	31.07 +/- 9.97	85.97 +/- 0.89	90.41 +/- 0.43	89.69 +/- 0.59
	natural	shuffled	76.67 +/- 0.34	87.21 +/- 0.17	65.8 +/- 6.11	23.06 +/- 5.3	81.84 +/- 0.43	83.94 +/- 0.33	62.86 +/- 0.19
	shuffled	natural	79.87 +/- 0.1	87.81 +/- 0.36	65.65 +/- 2.33	24.53 +/- 13.63	82.51 +/- 0.82	86.45 +/- 0.41	73.34 +/- 6.88
	shuffled	shuffled	79.75 +/- 0.0	88.21 +/- 0.24	64.88 +/- 6.32	22.43 +/- 10.79	82.65 +/- 0.42	86.25 +/- 0.4	63.15 +/- 2.2
\mathcal{M}_{UG}	natural	natural	71.93 +/- 0.21	66.94 +/- 9.21	53.7 +/- 1.01	0.92 +/- 2.06	70.57 +/- 0.66	83.17 +/- 1.5	58.59 +/- 0.33
	natural	shuffled	62.27 +/- 0.57	63.13 +/- 7.13	52.42 +/- 2.77	0.09 +/- 0.21	70.56 +/- 0.33	79.41 +/- 0.63	56.91 +/- 0.16
	shuffled	natural	67.62 +/- 0.3	66.49 +/- 0.49	52.17 +/- 1.26	0.0 +/- 0.0	70.37 +/- 0.93	79.93 +/- 1.01	57.59 +/- 0.29
	shuffled	shuffled	67.02 +/- 0.33	66.24 +/- 0.33	53.44 +/- 0.53	0.08 +/- 0.18	70.28 +/- 0.62	80.05 +/- 0.4	57.38 +/- 0.16

Table 5.9 Fine-tuning evaluation by varying different sources of word order (with mean and std dev). We vary the word order contained in the pre-trained model ($\mathcal{M}_N, \mathcal{M}_1, \mathcal{M}_{UG}$); in fine-tuning training set (natural and shuffled); and in fine-tuning evaluation (natural and shuffled). Here, *shuffled* corresponds to unigram shuffling of words in the input. In case of fine-tune evaluation containing shuffled input, we evaluate on a sample of 100 unigram permutations for each data point in the dev set of the corresponding task.

uate the effect of training and evaluation on natural and unnatural data in Table 5.9. We observe that all models perform poorly on the *shuffled* test set, compared to natural evaluation. However, interestingly, models have a slight advantage with a unigram randomized prior (\mathcal{M}_1), with CoLA having the biggest performance gain. PAWS task suffers the biggest drop in performance (from 94.49 to 62.22) but the lowest gain in \mathcal{M}_1 , confirming our conclusion from §5.4.1 that most of the word order information necessary for PAWS is learned from the task itself.

Furthermore, training on shuffled data surprisingly leads to high performance on natural data for \mathcal{M}_N in case of several tasks, the effect being weakest in case of CoLA and PAWS. This suggests that for tasks other than CoLA and PAWS, spurious correlations are leveraged by the models during fine-tuning (cf. Gururangan et al. 2018a, Poliak et al. 2018, Tsuchiya 2018). We also observe evidence of *domain matching*, where models improve their performance on evaluation on shuffled data when the training data source is changed from natural to shuffled (for \mathcal{M}_N , MNLI shuffled evaluation

improves from 68.11 to 79.96 just by changing the training corpus from natural to shuffled). We observe this behavior consistently for all tasks with all pre-trained representations.

5.5.4 Dependency parsing using Second order Tree CRF Neural Dependency Parser

Model	UD EWT		PTB	
	UAS	LAS	UAS	LAS
	90.92%	87.87%	95.42%	93.75%
\mathcal{M}_N				
\mathcal{M}_1	91.18%	88.19%	95.90%	94.35%
\mathcal{M}_2	91.11%	88.12%	95.74%	94.16%
\mathcal{M}_3	91.05%	87.94%	95.73%	94.14%
\mathcal{M}_4	90.88%	87.78%	95.77%	94.16%
\mathcal{M}_{UG}	90.47%	87.42%	95.81%	94.28%

Table 5.10 Unlabeled Attachment Score (UAS) on Dependency parsing task on two datasets, UD EWT and PTB, using the Second order Tree CRF Neural Dependency Parser Zhang et al. [2020]

We also conduct extensive experiments with Second Order Tree CRF Neural Dependency parser from Zhang et al. [2020], using their provided codebase.¹⁰ We report the results on UD EWT and PTB corpus in Table 5.10. Strangely enough, we find the gap to be even smaller across the different randomization models, even for some cases the performance on R_1 improves over OR . We suspect this result is due to two reasons: **(a)** Due to the presence of the complex Biaffine Dependency parser consisting of multiple LSTMs and individual MLP heads for each dependency arc (left and right), the majority of learning of the task is done by the parser itself; **(b)** Zhang et al. [2020] downsample the BERT representation to 100 dimensions which is then combined with

¹⁰<https://github.com/yzhangcs/parser>

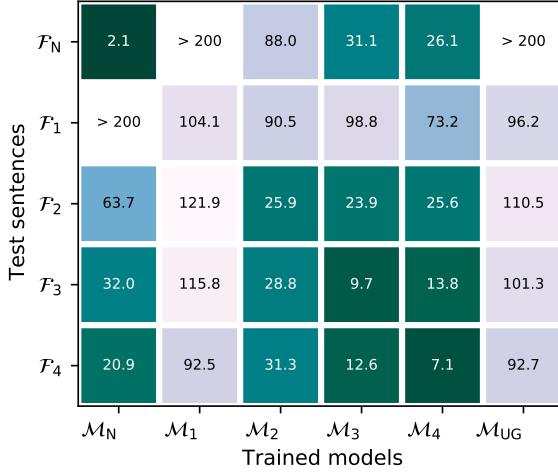


Figure 5.4 BPPL scores per model per test scenario.

the learned LSTM representations, thereby minimizing the impact of the pre-trained representations. Our hypothesis is confirmed by the published results of Zhang et al. [2020] on the Github repository, which shows a minimal gap between models with or without BERT.

5.5.5 Perplexity analysis

We measure perplexity of various pre-trained randomization models on text that is randomized using the same function \mathcal{F} . Conventional language models compute the perplexity of a sentence S by using past tokens ($S_{<t} = (w_1, w_2, \dots, w_{t-1})$) and the application of chain rule ($\sum_{t=1}^{|S|} \log P_{LM}(w_t | S_{t-1})$). However, this formulation is not defined for MLM, as a word is predicted using the entire sentence as a context. Following Salazar et al. [2020], we measure *Pseudo-Perplexity*, i.e., given a sentence S , we compute the log-probability of the missing word in S by iteratively masking out the specific

word, and computing the average log-probability per word in S :

$$\text{PLL}(S) = \frac{1}{|S|} \sum_{w \in S} \log P_{\text{MLM}}(w | S \setminus w; \theta) \quad (5.2)$$

We bootstrap the PLL score of a test corpus T by drawing 100 samples five times with replacement. We also similarly compute the bootstrap perplexity following Salazar et al.:

$$\text{BPLL}_T = \exp\left(-\frac{1}{N} \sum_{S \in W} \text{PLL}(S)\right), \quad (5.3)$$

where W is the combined bootstrap sample containing N sentences drawn with replacement from T . We compute this score on 6 pre-trained models, over four randomization schemes on the bootstrapped sample W (i.e., we use the same n-gram randomization function \mathcal{F}_i). Thus, we obtain a 5x6 matrix of BPLL scores, which we plot in Figure 5.4.

We observe that the pre-trained model \mathcal{M}_N has the lowest perplexity on the sentences with natural word order. Pre-trained models with random word order exhibit significantly higher perplexity than the normal word order sentences (top row). With the exception of \mathcal{M}_1 , the models pre-trained on randomized data ($\mathcal{M}_2, \mathcal{M}_3$ and \mathcal{M}_4) all display the lowest perplexity for their respective $n = 2, 3, 4$ randomizations. These results indicate that the models retain and detect the specific word order for which they were trained.

5.5.6 The usefulness of word order

The results in §5.4.1 suggest that, with proper fine-tuning, an unnaturally trained model can reach a level of performance comparable to that of a naturally pre-trained model. However, we want to understand whether natural word order pre-training offers any advantage during the early stages of fine-tuning. Towards that end, we

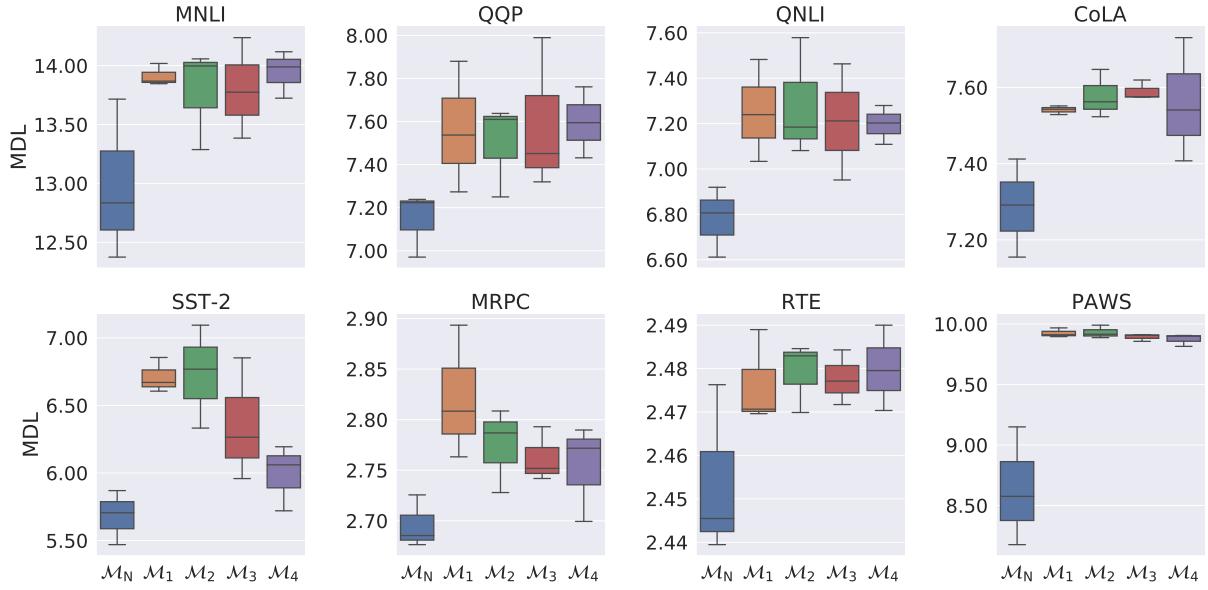


Figure 5.5 Rissanen Data Analysis Perez et al. [2021] on the GLUE benchmark and PAWS datasets. The lower minimum description length (MDL, measured in kilobits), the better the learning ability of the model.

turn to compute the Minimum Description Length [MDL; Rissanen, 1984]. MDL is designed to characterize the complexity of data as the length of the shortest program required to generate it. Thus, the length of the minimum description (in bits) should provide a fair estimate of how much word order is useful for fine-tuning in a few-shot setting. Specifically, we leverage the Rissanen Data Analysis (RDA) framework from Perez et al. [2021] to evaluate the MDL of pre-trained models on our set of downstream tasks. Under mild assumptions, if a pre-trained model θ_1 is useful for solving a particular task T over θ_2 , then the MDL in bits obtained by using θ_1 should be shorter than θ_2 . We follow the experimental setup of Perez et al. to compute the MDL on several tasks using $\theta = \{\mathcal{M}_N, \mathcal{M}_1, \mathcal{M}_2, \mathcal{M}_3, \mathcal{M}_4\}$, over three seeds and on three epochs of training. Concretely, RDA involves sampling 9 blocks of data from the dataset at random, where the size of each block is increased monotonically, training on 8 blocks while evaluating

the model’s loss (or *codelength*) on the ninth. The minimum number of data samples in the smallest block is set at 64, while the largest number of data samples used in the last block is 10,000.

We observe that the value of MDL is consistently lowest for naturally pre-trained data (Figure 5.5). For purportedly word order reliant datasets such as RTE, CoLA and PAWS, the gap between the MDL scores among the natural and unnatural models is high. PAWS, specifically, has the largest advantage in the beginning of optimization, however with more fine-tuning, the model re-learns correct word order (§5.4.1). The present analyses, when taken in conjunction with our main results in §5.4.1, suggest that fine-tuning on large training datasets with complex classifiers in the pursuit of state-of-the-art results has mostly nullified the impact of word order in the pre-trained representations. Few shot Bansal et al. [2020] and few sample Zhang et al. [2021] learning and evaluation could potentially require more word order signal, thereby encouraging the model to leverage its own learned syntax better.

5.5.7 At what point do models learn word order during pre-training?

Results from §5.4.1 beg the question: when, if at all, during pre-training does a model learn the natural word order? We aim to answer that question by comparing downstream task performance of RoBERTa base on intermediate checkpoints with that of the random word order pretrained models. The idea is to find the point during pre-training on natural corpus at which the model exceeds the task performance of the random pre-training model.

Performance on all tasks (Figure 5.6) increases rapidly during the first 20-25 epochs of pre-training. For some tasks, the word order information only helps after 30-50 pre-training epochs.

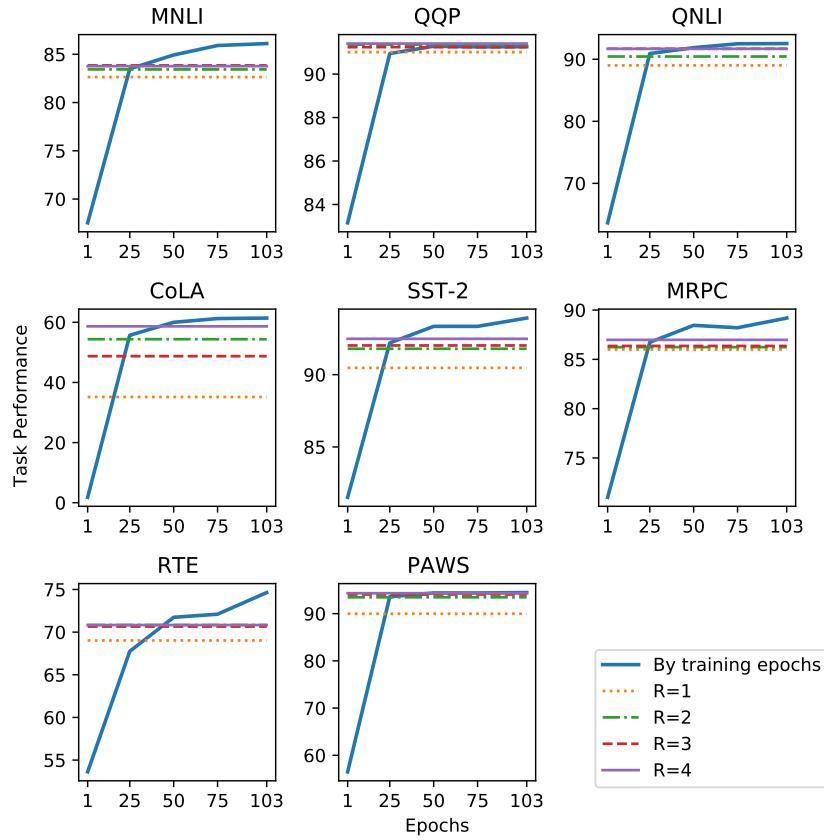


Figure 5.6 Comparison among GLUE task performance from different steps in pre-training of RoBERTa on BookWiki Corpus.

5.5.8 More results from Syntactic Probes

We computed the Pareto Hypervolume on the dependency parsing task [Pimentel et al., 2020a]. The Pareto Hypervolume is computed as the Area Under Curve (AUC) score over all hyperparameter runs, where the models are arranged based on their complexity. We observe minimal differences in the Pareto Hypervolumes (Table 5.13) among \mathcal{M}_N and the randomization models for both datasets.

We also investigated two “easy” tasks, Part-of-Speech tagging (POS) and Dependency Arc Labeling (DAL) from the Pareto Probing framework. For POS (Table 5.11)

Model	UD EWT		PTB	
	MLP	Linear	MLP	Linear
\mathcal{M}_N	93.74 +/- 0.15	88.82 +/- 0.42	97.07 +/- 0.38	93.1 +/- 0.65
\mathcal{M}_1	88.60 +/- 3.43	80.76 +/- 3.38	95.33 +/- 0.37	87.83 +/- 1.86
\mathcal{M}_2	93.39 +/- 0.45	87.58 +/- 1.06	96.96 +/- 0.15	91.80 +/- 0.50
\mathcal{M}_3	92.89 +/- 0.65	86.78 +/- 1.32	97.03 +/- 0.13	91.70 +/- 0.70
\mathcal{M}_4	92.83 +/- 0.61	87.23 +/- 0.77	96.96 +/- 0.12	92.08 +/- 0.39
\mathcal{M}_{UG}	89.10 +/- 0.21	79.75 +/- 0.5	94.12 +/- 0.01	84.15 +/- 0.51

Table 5.11 Accuracy on the part-of-speech labelling task (POS) on two datasets, UD EWT and PTB, using the Pareto Probing framework Pimentel et al. [2020b].

and DAL (Table 5.12), since these tasks are simpler than DEP, the gap between \mathcal{M}_N and unnaturally pre-trained models reduces even more drastically. The gap between \mathcal{M}_N and \mathcal{M}_1 reduces to just 3.5 points on average for PTB in both POS and DAL.

5.5.9 Non parametric probes

Probability difference. In the original formulation [Goldberg, 2019b, Wolf, 2019a], the effectiveness of each stimulus is determined by the accuracy metric, computed as the number of times the probability of the correct focus word is greater than that of the incorrect word ($P(\text{good}) > P(\text{bad})$). We observed that this metric might not be reliable per se, since the probabilities may themselves be extremely low for all tokens, even when focus word probability decreases drastically from \mathcal{M}_N to \mathcal{M}_{UG} . Thus, we also report the mean difference of probabilities, $(\frac{1}{N} \sum_i^N P(\text{good}_i) - P(\text{bad}_i))$, scaled up by a factor of 100 for ease of observation, in ??, ?? and ???. We observe the highest difference between probabilities of the correct and incorrect focus words for the model pre-trained on the natural word order (\mathcal{M}_N). Moreover, with each step from \mathcal{M}_1 to \mathcal{M}_4 ,

Model	UD EWT		PTB	
	MLP	Linear	MLP	Linear
\mathcal{M}_N	89.63 +/- 0.60	84.35 +/- 0.78	93.96 +/- 0.63	88.35 +/- 1.00
\mathcal{M}_1	83.55 +/- 3.31	75.26 +/- 3.08	91.10 +/- 0.38	82.34 +/- 1.37
\mathcal{M}_2	88.57 +/- 0.68	82.05 +/- 1.10	93.27 +/- 0.26	86.88 +/- 0.87
\mathcal{M}_3	88.69 +/- 1.09	82.37 +/- 1.26	93.46 +/- 0.29	87.12 +/- 0.72
\mathcal{M}_4	88.66 +/- 0.76	82.58 +/- 1.04	93.49 +/- 0.33	87.30 +/- 0.79
\mathcal{M}_{UG}	84.93 +/- 0.34	76.30 +/- 0.52	89.98 +/- 0.43	78.59 +/- 0.68

Table 5.12 Accuracy on the dependency arc labelling task (DAL) on two datasets (with mean and std dev), UD EWT and PTB, using the Pareto Prob-ing framework Pimentel et al. [2020a].

the difference between probabilities of correct and incorrect focus words increases, albeit marginally, showing that pre-trained models with fewer n-gram words perturbed capture more word order information. \mathcal{M}_{UG} , the model with the distributional prior ablated, performs the worst, as expected.

Accuracy comparison. We provide the accuracy as measured by Goldberg [2019b], Wolf [2019a] on the probing stimuli in Table 5.14, Table 5.15 and Table 5.16. We also highlight the difference in probability ($P(\text{good}) - P(\text{bad})$) in the table to provide a more accurate picture. All experiments were conducted on three pre-trained seeds for each model in our set of models. However, the low token probabilities in \mathcal{M}_{UG} tend to present unreliable scores. For example, in the case of Gulordava et al. [2018b] stimuli, unnatural models provide better scores compared to the natural model. We also observe for the Linzen et al. [2016] stimuli that the results on model condition 4 (number of attractors) are surprisingly high for \mathcal{M}_{UG} whereas the individual token probabilities are lowest. We believe these inconsistencies stem from extremely low token probabilities themselves.

Model	UD EWT	PTB
\mathcal{M}_N	0.528 +/- 0.01	0.682 +/- 0.01
\mathcal{M}_1	0.489 +/- 0.03	0.648 +/- 0.01
\mathcal{M}_2	0.529 +/- 0.00	0.681 +/- 0.01
\mathcal{M}_3	0.528 +/- 0.02	0.689 +/- 0.01
\mathcal{M}_4	0.525 +/- 0.00	0.683 +/- 0.01
\mathcal{M}_{UG}	0.510 +/- 0.01	0.640 +/- 0.05

Table 5.13 Pareto Hypervolume of dependency parsing task (DEP) on two datasets (with mean and std dev), UD EWT and PTB, using the Pareto Probing framework Pimentel et al. [2020b].

model condition	\mathcal{M}_N	\mathcal{M}_{UG}	\mathcal{M}_1	\mathcal{M}_2	\mathcal{M}_3	\mathcal{M}_4
1	93.45 (0.89) [25.04]	58.87 (0.41) [0.0]	59.96 (1.58) [0.08]	63.63 (0.6) [1.25]	64.7 (1.44) [2.79]	70.47 (1.9) [4.01]
2	92.8 (1.22) [23.8]	63.03 (1.35) [0.01]	58.22 (1.5) [0.09]	61.15 (2.07) [0.82]	63.84 (2.41) [2.09]	64.7 (1.92) [3.07]
3	87.71 (1.34) [22.03]	64.06 (3.52) [0.0]	56.69 (2.98) [0.03]	56.83 (3.63) [0.85]	61.1 (0.32) [2.02]	63.0 (3.36) [2.35]
4	92.67 (0.52) [22.16]	76.33 (1.38) [0.0]	62.33 (7.61) [0.08]	63.17 (9.09) [1.12]	69.42 (1.77) [2.1]	67.67 (7.02) [3.43]

Table 5.14 Linzen et al. [2016] stimuli results in raw accuracy. Values in parenthesis reflect the standard deviation over different seeds of pre-training. Values in square brackets indicate the mean probability difference among correct and incorrect words.

model condition	\mathcal{M}_N	\mathcal{M}_{UG}	\mathcal{M}_1	\mathcal{M}_2	\mathcal{M}_3	\mathcal{M}_4
0	79.42 (5.5) [2.43]	47.83 (3.76) [-0.0]	53.67 (1.38) [0.03]	58.75 (6.38) [0.05]	63.58 (4.11) [0.14]	63.75 (3.28) [0.17]
1	72.83 (4.07) [2.55]	44.5 (0.5) [0.0]	70.83 (5.8) [0.02]	64.83 (1.76) [-0.09]	71.67 (6.71) [0.21]	71.5 (2.65) [0.61]
2	55.56 (0.0) [0.92]	88.89 (11.11) [0.0]	81.48 (12.83) [0.03]	51.85 (6.42) [0.04]	62.96 (6.42) [0.38]	74.07 (16.97) [0.61]

Table 5.15 Gulordava et al. [2018b] stimuli results in raw accuracy. Values in parenthesis reflect the standard deviation over different seeds of pre-training. Values in square brackets indicate the mean probability difference among correct and incorrect words.

Balancing datasets on inflection by upsampling. The stimuli datasets of Linzen et al. [2016] and Gulordava et al. [2018b] turned out to be heavily skewed towards words where singular was the correct inflection (as opposed to plural). This dataset imbalance caused the weak models (such as \mathcal{M}_{UG}) to have surprisingly high scores - the weak models were consistently providing higher probability for the singular inflection (Table 5.17). We upsample for both datasets, balancing the frequency of correct singular and plural inflections. We compute the upsampling number to the next multiple of 100 of the count of original singular inflections. For example, in condition 4 of Linzen et al. [2016] dataset, we upsample both S and P to 300 rows each. This type of balancing via upsampling largely alleviated the inconsistencies we observed, and might prove to be useful when evaluating other models on these datasets in future.

Model condition	\mathcal{M}_N	\mathcal{M}_{UG}	\mathcal{M}_1	\mathcal{M}_2	\mathcal{M}_3	\mathcal{M}_4
AOR	89.98 (1.96) [29.16]	50.0 (0.01) [0.0]	60.17 (1.61) [1.7]	66.61 (7.1) [3.87]	63.57 (2.39) [2.45]	61.26 (4.91) [1.16]
AOR-T	77.4 (7.74) [13.84]	50.0 (0.0) [0.0]	78.88 (0.64) [0.15]	52.17 (2.14) [0.43]	48.85 (3.8) [0.25]	57.06 (3.49) [0.55]
APP	89.94 (4.16) [27.06]	50.01 (0.02) [-0.0]	70.34 (1.9) [0.68]	53.61 (3.3) [0.2]	53.03 (1.75) [0.79]	60.6 (4.41) [0.91]
ARC	85.06 (5.92) [1.2]	50.05 (0.08) [-0.0]	62.39 (1.91) [0.11]	74.57 (5.99) [0.07]	67.55 (3.84) [0.07]	62.88 (3.45) [0.03]
ASR	87.19 (3.58) [26.18]	50.0 (0.0) [-0.0]	78.55 (10.01) [1.91]	81.73 (5.1) [2.91]	62.8 (0.35) [1.38]	67.23 (6.82) [1.48]
IOR	89.83 (3.33) [0.4]	50.55 (0.95) [-0.0]	56.28 (2.66) [0.01]	58.96 (4.28) [0.04]	70.49 (2.2) [0.04]	62.82 (8.51) [0.08]
IOR-T	74.05 (8.26) [0.2]	50.61 (1.05) [-0.0]	52.63 (2.07) [0.01]	57.35 (4.88) [0.01]	61.85 (4.75) [0.01]	55.16 (6.59) [0.03]
ISC	85.87 (9.6) [5.27]	50.0 (0.0) [0.0]	67.85 (2.62) [0.07]	82.66 (9.43) [0.0]	77.69 (4.51) [0.31]	68.65 (5.71) [0.06]
LVC	93.0 (0.75) [26.58]	49.92 (0.14) [-0.0]	70.42 (6.79) [0.1]	87.5 (7.26) [1.44]	85.42 (3.84) [0.63]	81.08 (5.13) [0.66]
SCM	88.6 (3.49) [5.72]	50.0 (0.0) [0.02]	63.73 (7.94) [0.12]	82.12 (0.92) [0.17]	86.44 (3.67) [2.62]	80.27 (2.46) [0.26]
SRX	91.0 (6.07) [2.72]	50.0 (0.0) [0.0]	88.0 (10.11) [0.1]	92.25 (10.27) [0.32]	94.25 (5.02) [0.02]	91.0 (6.5) [3.07]
SVA	95.33 (7.23) [31.35]	50.0 (0.0) [0.02]	86.0 (5.29) [0.57]	85.17 (12.87) [3.17]	94.67 (5.25) [6.15]	88.83 (9.57) [30.57]
SVC	97.54 (1.58) [19.27]	50.0 (0.0) [-0.0]	83.58 (4.58) [0.84]	83.71 (8.78) [1.17]	93.29 (7.4) [6.39]	81.04 (3.66) [11.07]

Table 5.16 Marvin and Linzen [2018] stimuli results in raw accuracy. Values in parenthesis reflect the standard deviation over different seeds of pre-training. Values in square brackets indicate the mean probability difference among correct and incorrect words. Abbreviations: Simple Verb Agreement (SVA), In a sentential complement (SCM), Short VP Coordination (SVC), Long VP Coordination (LVC), Across a prepositional phrase (APP), Across a subject relative clause (ASR), Across an object relative clause (AOR), Across an object relative (no *that*) (AOR-T), In an object relative clause (IOR), In an object relative clause (no *that*) (IOR-T), Simple Reflexive (SRX), In a sentential complement (ISC), Across a relative clause (ARC), Simple NPI (SNP).

Model condition	\mathcal{M}_N	\mathcal{M}_{UG}	\mathcal{M}_1	\mathcal{M}_2	\mathcal{M}_3	\mathcal{M}_4	S/P
1	94.04 (0.8)	62.64 (0.5)	62.18 (1.33)	64.91 (0.14)	65.35 (1.78)	70.88 (1.88)	14011 / 10112
2	93.28 (0.94)	71.24 (0.85)	63.03 (1.69)	62.92 (2.57)	65.25 (3.13)	65.61 (2.35)	3120 / 1312
3	89.1 (0.58)	74.05 (1.85)	62.94 (3.13)	59.18 (3.32)	63.54 (1.72)	63.05 (2.0)	733 / 215
4	90.53 (0.9)	80.03 (0.59)	63.16 (4.83)	63.94 (6.92)	66.41 (3.17)	66.28 (4.64)	206 / 51

Table 5.17 Linzen et al. [2016] stimuli results in raw accuracy on original, unbalanced data. Values in parenthesis reflect the standard deviation. S/P reflects the count of correct singular and plural focus words.

5.6 Related Work

5.6.1 Sensitivity to word order in NLU

Information order has been a topic of research in computational linguistics since Barzilay and Lee [2004] introduced the task of ranking sentence orders as an evaluation for language generation quality, an approach which was subsequently also used to evaluate readability and dialogue coherence [Barzilay and Lapata, 2008, Laban et al., 2021].

More recently, several research groups have investigated information order for words rather than sentences as an evaluation of model humanlikeness. Sinha et al. [2021d] investigate the task of natural language inference (NLI) and find high accuracy on permuted examples for different Transformer and pre-Transformer era models, across English and Chinese datasets Hu et al. [2020a]. Gupta et al. [2021] use targeted permutations on RoBERTa-based models and show word order insensitivity across natural language inference (MNLI), paraphrase detection (QQP) and sentiment analysis tasks (SST-2). Pham et al. [2020b] show insensitivity on a larger set of tasks, including the entire GLUE benchmark, and find that certain tasks in GLUE, such as CoLA and RTE are more sensitive to permutations than others. Ettinger [2020] recently observed that BERT accuracy decreases for some word order perturbed examples, but not for others. In all these prior works, models were given access to normal word order at

5.6.2 Randomization ablations

Random controls have been explored in a variety of prior work. Wieting and Kiela [2019] show that random sentence encoders are surprisingly powerful baselines. Gauthier and Levy [2019] use random sentence reordering to label some tasks as “syntax-light” making them more easily decodeable from images of the brain. Shen et al. [2021] show that entire layers of MLM transformers can be randomly initialized and kept frozen throughout training without detrimental effect and that those layers perform better on some probing tasks than their frozen counterparts. Models have been found to be surprisingly robust to randomizing or cutting syntactic tree structures they were hoped to rely on Scheible and Schütze [2013], Williams et al. [2018b], and randomly permuting attention weights often induces only minimal changes in output Jain and Wallace [2019]. In computer vision, it is well known that certain architectures constitute good “deep image priors” for fine-tuning Ulyanov et al. [2018] or pruning Frankle et al. [2020], and that even randomly wired networks can perform well at image recognition Xie et al. [2019]. Here, we explore randomizing the data, rather than the model, to assess whether certain claims about which phenomena the model has learned are established in fact.

5.6.3 Synthetic pre-training

Kataoka et al. [2020] found that pre-training on synthetically generated fractals for image classification is a very strong prior for subsequent fine-tuning on real image data. In language modeling, Papadimitriou and Jurafsky [2020] train LSTMs [Hochreiter and Schmidhuber, 1997] on non-linguistic data with latent structure such as MIDI music or Java code provides better test performance on downstream tasks than a randomly initialized model. They observe that even when there is no vocabulary overlap among source and target languages, LSTM language models leverage the latent hierarchical structure of the input to obtain better performance than a random, Zipfian corpus of

the same vocabulary.

5.6.4 On the utility of probing tasks

Many recent papers provide compelling evidence that BERT contains a surprising amount of syntax, semantics, and world knowledge Giulianelli et al. [2018], Rogers et al. [2020a], Lakretz et al. [2019], Jumelet et al. [2019, 2021]. Many of these works involve diagnostic classifiers Hupkes et al. [2018] or *parametric* probes, i.e. a function atop learned representations that is optimized to find linguistic information. How well the probe learns a given signal can be seen as a proxy for linguistic knowledge encoded in the representations. However, the community is divided on many aspects of probing [Belinkov, 2021] including how complex probes should be. Many prefer *simple* linear probes over the complex ones Alain and Bengio [2017], Hewitt and Manning [2019b], Hall Maudslay et al. [2020]. However, complex probes with strong representational capacity are able to extract the most information from representations [Voita and Titov, 2020, Pimentel et al., 2020b, Hall Maudslay et al., 2020]. In this chapter, we follow Pimentel et al. [2020a] and use *both* simple (linear) and complex (non-linear) models, as well as “complex” tasks (dependency parsing). As an alternative to parametric probes, stimulus-based *non-parametric* probing Linzen et al. [2016], Jumelet and Hupkes [2018], Marvin and Linzen [2018], Gulordava et al. [2018a], Warstadt et al. [2019a, 2020b,c], Ettinger [2020], Lakretz et al. [2021] has been used to show that even without a learned probe, BERT can predict syntactic properties with high confidence Goldberg [2019b], Wolf [2019a]. We use this class of non-parametric probes to investigate RoBERTa’s ability to learn word order during pre-training.

5.7 Discussion

The assumption that word order information is crucial for any classical NLP pipeline (especially for English) is deeply ingrained in our understanding of syntax itself [Chomsky, 1957]: without order, many linguistic constructs are undefined. Our fine-tuning results in §5.4.1 and parametric probing results in §5.4.2, however, suggests that MLMs do not need to rely much on word order to achieve high accuracy, bringing into question previous claims that they learn a “classical NLP pipeline.”

The assumption that word order information is crucial for any classical NLP pipeline (especially for English) is deeply ingrained in our understanding of syntax itself [Chomsky, 1957]: without order, most linguistic constructs are undefined (e.g. dependency or constituency parses would no longer be syntactic trees, what would sentences be but mere lists of words).

One might ask, though, whether an NLP pipeline would really need natural word order at all: can transformers not simply learn what the correct word order is from unordered text? First, the lower non-parametric probing accuracies of the randomized models indicate that they are not able to accurately reconstruct the original word order (see also §5.5.1). But even if models were able to “unshuffle” the words under our unnatural pre-training set up, they would only be doing so based on distributional information. Models would then abductively learn only the most likely word order. While models might infer a distribution over possible orders and use that information to structure their representations [Papadimitriou et al., 2021], syntax is not about *possible* or even *the most likely* orders: it is about the *actual* order. That is, even if one concludes in the end that Transformers are able to perform word order reconstruction based on distributional information, and recover almost all downstream performance based solely on that, we ought to be a lot more careful when making claims about what our evaluation datasets are telling us.

Thus, our results seem to suggest that we may need to revisit what we mean by

“linguistic structure,” and perhaps subsequently acknowledge that we may not need human-like linguistic abilities for most NLP tasks. Or, our results can be interpreted as evidence that we need to develop more challenging and more comprehensive evaluations, if we genuinely want to measure linguistic abilities, however those are defined, in NLP models.

To summarize, in this chapter, we revisited the hypothesis that masked language modelling’s impressive performance can be explained in part by its ability to learn classical NLP pipelines. We investigated targeted pre-training on sentences with various degrees of randomization in their word order, and observed overwhelmingly that MLM’s success is most likely not due to its ability to discover syntactic and semantic mechanisms necessary for a traditional language processing pipeline during pre-training. Instead, our experiments suggest that MLM’s success can largely be explained by it having learned higher-order distributional statistics that make for a useful prior for subsequent fine-tuning. These results should hopefully encourage the development of better, more challenging tasks that require sophisticated reasoning, and harder probes to narrow down what exact linguistic information is present in the representations learned by our models.

5.8 Follow-up findings in the community



Figure 5.7 The difference in word probabilities for stimuli in Marvin and Linzen [2018]: Simple Verb Agreement (SVA), In a sentential complement (SCM), Short VP Coordination (SVC), Long VP Coordination (LVC), Across a prepositional phrase (APP), Across a subject relative clause (ASR), Across an object relative clause (AOR), Across an object relative (no *that*) (AOR-T), In an object relative clause (IOR), In an object relative clause (no *that*) (IOR-T), Simple Reflexive (SRX), In a sentential complement (ISC), Across a relative clause (ARC), Simple NPI (SNP).

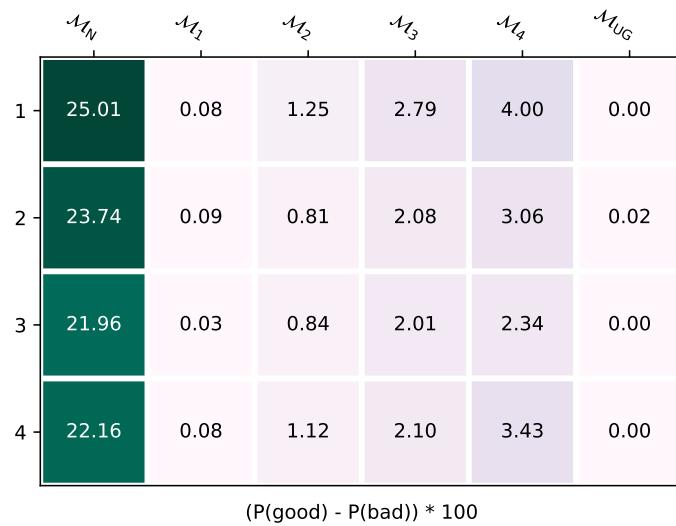


Figure 5.8 Linzen et al. [2016]

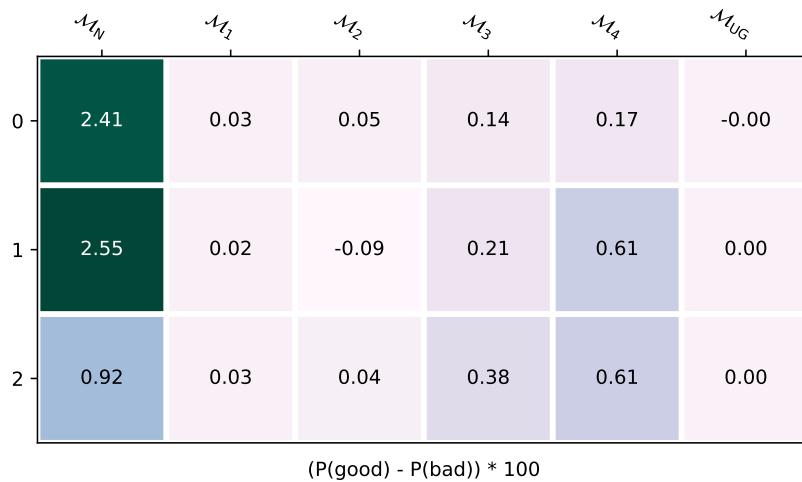


Figure 5.9 Gulordava et al. [2018b]

Chapter 6

Measuring systematic generalization by exploiting absolute positions

6.1 Technical Background

6.2 Systematic understanding of absolute position embeddings

6.3 Related Work

6.4 Experiments

6.5 Discussion

Chapter 7

Conclusion

7.1 Summary

7.2 Limitations

7.3 Future Work

A few years ago, Manning [2015] encouraged NLP to consider “the details of human language, how it is learned, processed, and how it changes, rather than just chasing state-of-the-art numbers on a benchmark task.” We expand upon this view, and suggest one particular future direction: we should train models not only to do well on clean test data, but also to not to overgeneralize to corrupted input.

Bibliography

- Anne Abeille. Lexical and syntactic rules in a Tree Adjoining Grammar. In *28th Annual Meeting of the Association for Computational Linguistics*, pages 292–298, Pittsburgh, Pennsylvania, USA, June 1990. Association for Computational Linguistics. doi: 10.3115/981823.981860. URL <https://www.aclweb.org/anthology/P90-1037>.
- Aishwarya Agrawal, Dhruv Batra, and Devi Parikh. Analyzing the behavior of visual question answering models. *arXiv preprint arXiv:1606.07356*, 2016.
- Guillaume Alain and Yoshua Bengio. Understanding intermediate layers using linear classifier probes. In *ICLR 2017, Workshop Track Proceedings*. OpenReview.net, 2017. URL <https://openreview.net/forum?id=HJ4-rAVt1>.
- Matteo Alleman, Jonathan Mamou, Miguel A Del Rio, Hanlin Tang, Yoon Kim, and SueYeon Chung. Syntactic perturbations reveal representational correlates of hierarchical phrase structure in pretrained language models. *arXiv preprint arXiv:2104.07578*, 2021. URL <https://arxiv.org/abs/2104.07578>.
- Aixiu An, Peng Qian, Ethan Wilcox, and Roger Levy. Representation of constituents in neural language models: Coordination phrase as a case study. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, pages 2888–2899, Hong Kong, China, November 2019. Association for Computational Linguistics.

- tics. doi: 10.18653/v1/D19-1287. URL <https://www.aclweb.org/anthology/D19-1287>.
- Alan D Baddeley, Graham J Hitch, and Richard J Allen. Working memory and binding in sentence recall. *Journal of Memory and Language*, 2009. URL <https://doi.org/10.1016/j.jml.2009.05.004>.
- Dzmitry Bahdanau, Shikhar Murty, Michael Noukhovitch, Thien Huu Nguyen, Harm de Vries, and Aaron Courville. Systematic generalization: What is required and can it be learned? In *International Conference on Learning Representations*, 2019. URL <https://openreview.net/forum?id=HkezXnA9YX>.
- Trapit Bansal, Rishikesh Jha, and Andrew McCallum. Learning to few-shot learn across diverse natural language classification tasks. In *Proceedings of the 28th International Conference on Computational Linguistics*, pages 5108–5123, Barcelona, Spain (Online), December 2020. International Committee on Computational Linguistics. doi: 10.18653/v1/2020.coling-main.448. URL <https://aclanthology.org/2020.coling-main.448>.
- Regina Barzilay and Mirella Lapata. Modeling local coherence: An entity-based approach. *Computational Linguistics*, 34(1):1–34, 2008. doi: 10.1162/coli.2008.34.1.1. URL <https://aclanthology.org/J08-1001>.
- Regina Barzilay and Lillian Lee. Catching the drift: Probabilistic content models, with applications to generation and summarization. In *Proceedings of the Human Language Technology Conference of the North American Chapter of the Association for Computational Linguistics: HLT-NAACL 2004*, pages 113–120, Boston, Massachusetts, USA, May 2 - May 7 2004. Association for Computational Linguistics. URL <https://aclanthology.org/N04-1015>.

Yonatan Belinkov. Probing classifiers: Promises, shortcomings, and alternatives. *arXiv preprint arXiv:2102.12452*, 2021.

Douglas K Bemis and Liina Pylkkänen. Basic linguistic composition recruits the left anterior temporal lobe and left angular gyrus during both listening and reading. *Cerebral Cortex*, 2013.

Emily M. Bender and Alexander Koller. Climbing towards NLU: On meaning, form, and understanding in the age of data. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pages 5185–5198, Online, July 2020. Association for Computational Linguistics. doi: 10.18653/v1/2020.acl-main.463. URL <https://www.aclweb.org/anthology/2020.acl-main.463>.

Luisa Bentivogli, Peter Clark, Ido Dagan, and Danilo Giampiccolo. The fifth PASCAL recognizing textual entailment challenge. In *TAC*, 2009. URL <https://citeseerx.ist.psu.edu/viewdoc/download?doi=10.1.1.232.1231&rep=rep1&type=pdf>.

Jean-Phillipe Bernardy and Shalom Lappin. Using deep neural networks to learn syntactic agreement. In *Linguistic Issues in Language Technology, Volume 15*, 2017. CSLI Publications, 2017. URL <https://www.aclweb.org/anthology/2017.lilt-15.3>.

Ann Bies, Justin Mott, Colin Warner, and Seth Kulick. English web treebank. *Linguistic Data Consortium, Philadelphia, PA*, 2012. URL <https://catalog.ldc.upenn.edu/LDC2012T13>.

David M. Blei, Andrew Y. Ng, and Michael I. Jordan. Latent dirichlet allocation. *JMLR*, 2003. URL <https://www.jmlr.org/papers/volume3/blei03a/blei03a.pdf>.

- Samuel R. Bowman, Gabor Angeli, Christopher Potts, and Christopher D. Manning. A large annotated corpus for learning natural language inference. In *Proceedings of the 2015 Conference on Empirical Methods in Natural Language Processing*, pages 632–642, Lisbon, Portugal, September 2015a. Association for Computational Linguistics. doi: 10.18653/v1/D15-1075. URL <https://www.aclweb.org/anthology/D15-1075>.
- Samuel R Bowman, Gabor Angeli, Christopher Potts, and Christopher D Manning. A large annotated corpus for learning natural language inference. In *Proceedings of the 2015 Conference on Empirical Methods in Natural Language Processing*, pages 632–642, 2015b.
- Joan Bresnan, Ash Asudeh, Ida Toivonen, and Stephen Wechsler. *Lexical-functional syntax*. John Wiley & Sons, 2015.
- Tom Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared D Kaplan, Prafulla Dhariwal, Arvind Neelakantan, Pranav Shyam, Girish Sastry, Amanda Askell, Sandhini Agarwal, Ariel Herbert-Voss, Gretchen Krueger, Tom Henighan, Rewon Child, Aditya Ramesh, Daniel Ziegler, Jeffrey Wu, Clemens Winter, Chris Hesse, Mark Chen, Eric Sigler, Mateusz Litwin, Scott Gray, Benjamin Chess, Jack Clark, Christopher Berner, Sam McCandlish, Alec Radford, Ilya Sutskever, and Dario Amodei. Language models are few-shot learners. In H. Larochelle, M. Ranzato, R. Hadsell, M. F. Balcan, and H. Lin, editors, *Advances in Neural Information Processing Systems*, volume 33, pages 1877–1901. Curran Associates, Inc., 2020. URL <https://proceedings.neurips.cc/paper/2020/file/1457c0d6bfcb4967418bfb8ac142f64a-Paper.pdf>.
- James McKeen Cattell. The time it takes to see and name objects. *Mind*, os-XI(41):63–65, 01 1886. ISSN 0026-4423. doi: 10.1093/mind/os-XI.41.63. URL <https://doi.org/10.1093/mind/os-XI.41.63>.

- Rui Chaves. What don't RNN language models learn about filler-gap dependencies? In *Proceedings of the Society for Computation in Linguistics 2020*, pages 1–11, New York, New York, January 2020. Association for Computational Linguistics. URL <https://www.aclweb.org/anthology/2020.scil-1.1>.
- Qian Chen, Xiaodan Zhu, Zhen-Hua Ling, Si Wei, Hui Jiang, and Diana Inkpen. Enhanced LSTM for natural language inference. In *Proceedings of the 55th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 1657–1668, Vancouver, Canada, July 2017. Association for Computational Linguistics. doi: 10.18653/v1/P17-1152. URL <https://www.aclweb.org/anthology/P17-1152>.
- Gennaro Chierchia and Sally McConnell-Ginet. *Meaning and grammar: An Introduction to Semantics*. Cambridge, Ma: MIT Press, 1990.
- Kyunghyun Cho, Bart van Merriënboer, Caglar Gulcehre, Dzmitry Bahdanau, Fethi Bougares, Holger Schwenk, and Yoshua Bengio. Learning phrase representations using rnn encoder–decoder for statistical machine translation. In *Proceedings of the 2014 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 1724–1734, 2014.
- Noam Chomsky. *Syntactic structures*. Walter de Gruyter, 1957.
- Grzegorz Chrupała and Afra Alishahi. Correlating neural and symbolic representations of language. In *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics*, pages 2952–2962, Florence, Italy, July 2019. Association for Computational Linguistics. doi: 10.18653/v1/P19-1283. URL <https://www.aclweb.org/anthology/P19-1283>.
- Guglielmo Cinque. *Adverbs and functional heads: A cross-linguistic perspective*. Oxford University Press on Demand, 1999.

Kevin Clark, Urvashi Khandelwal, Omer Levy, and Christopher D. Manning. What does BERT look at? an analysis of BERT’s attention. In *Proceedings of the 2019 ACL Workshop BlackboxNLP: Analyzing and Interpreting Neural Networks for NLP*, pages 276–286, Florence, Italy, August 2019. Association for Computational Linguistics. doi: 10.18653/v1/W19-4828. URL <https://www.aclweb.org/anthology/W19-4828>.

Ronan Collobert and Jason Weston. A unified architecture for natural language processing: Deep neural networks with multitask learning. In *ICML*, 2008. URL <https://dl.acm.org/doi/10.1145/1390156.1390177>.

Cleo Condoravdi, Dick Crouch, Valeria de Paiva, Reinhard Stolle, and Daniel G. Bobrow. Entailment, intensionality and text understanding. In *Proceedings of the HLT-NAACL 2003 Workshop on Text Meaning*, pages 38–45, 2003. URL <https://www.aclweb.org/anthology/W03-0906>.

Alexis Conneau and Douwe Kiela. SentEval: An evaluation toolkit for universal sentence representations. In *Proceedings of the Eleventh International Conference on Language Resources and Evaluation (LREC 2018)*, Miyazaki, Japan, May 2018. European Language Resources Association (ELRA). URL <https://aclanthology.org/L18-1269>.

Alexis Conneau, Douwe Kiela, Holger Schwenk, Loïc Barrault, and Antoine Bordes. Supervised learning of universal sentence representations from natural language inference data. In *Proceedings of the 2017 Conference on Empirical Methods in Natural Language Processing*, pages 670–680, Copenhagen, Denmark, September 2017. Association for Computational Linguistics. doi: 10.18653/v1/D17-1070. URL <https://www.aclweb.org/anthology/D17-1070>.

Alexis Conneau, German Kruszewski, Guillaume Lample, Loïc Barrault, and Marco

- Baroni. What you can cram into a single \$&!#* vector: Probing sentence embeddings for linguistic properties. In *Proceedings of the 56th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 2126–2136, Melbourne, Australia, July 2018. Association for Computational Linguistics. doi: 10.18653/v1/P18-1198. URL <https://aclanthology.org/P18-1198>.
- Jillian Da Costa and Rui Chaves. Assessing the ability of transformer-based neural models to represent structurally unbounded dependencies. In *Proceedings of the Society for Computation in Linguistics 2020*, pages 12–21, New York, New York, January 2020. Association for Computational Linguistics. URL <https://www.aclweb.org/anthology/2020.scil-1.2>.
- Ido Dagan, Oren Glickman, and Bernardo Magnini. The pascal recognising textual entailment challenge. In *Machine Learning Challenges Workshop*. Springer, 2005a. URL https://dl.acm.org/doi/10.1007/11736790_9.
- Ido Dagan, Oren Glickman, and Bernardo Magnini. The PASCAL recognising textual entailment challenge. In *Machine Learning Challenges Workshop*, pages 177–190. Springer, 2005b. URL https://link.springer.com/chapter/10.1007/11736790_9.
- Ido Dagan, Oren Glickman, and Bernardo Magnini. The PASCAL recognising textual entailment challenge. In *Machine learning challenges. evaluating predictive uncertainty, visual object classification, and recognising textual entailment*. Springer, 2006.
- Rajarshi Das, Shehzaad Dhuliawala, Manzil Zaheer, Luke Vilnis, Ishan Durugkar, Akshay Krishnamurthy, Alex Smola, and Andrew McCallum. Go for a walk and arrive at the answer: Reasoning over paths in knowledge bases using reinforcement learning. In *International Conference on Learning Representations*, 2018. URL <https://openreview.net/forum?id=Syg-YfWCW>.

- Ishita Dasgupta, Demi Guo, Andreas Stuhlmüller, Samuel J Gershman, and Noah D Goodman. Evaluating compositionality in sentence embeddings. In *Proceedings of Annual Meeting of the Cognitive Science Society*, 2018. URL <https://arxiv.org/abs/1802.04302>.
- Forrest Davis and Marten van Schijndel. Recurrent neural network language models always learn English-like relative clause attachment. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pages 1979–1990, Online, July 2020. Association for Computational Linguistics. doi: 10.18653/v1/2020.acl-main.179. URL <https://www.aclweb.org/anthology/2020.acl-main.179>.
- Scott Deerwester, Susan T Dumais, George W Furnas, Thomas K Landauer, and Richard Harshman. Indexing by latent semantic analysis. *Journal of the American society for information science*, 41(6):391–407, 1990. URL [https://asistdl.onlinelibrary.wiley.com/doi/abs/10.1002/\(SICI\)1097-4571\(199009\)41:6%3C391::AID-ASI1%3E3.0.CO;2-9](https://asistdl.onlinelibrary.wiley.com/doi/abs/10.1002/(SICI)1097-4571(199009)41:6%3C391::AID-ASI1%3E3.0.CO;2-9).
- Dorottya Demszky, Kelvin Guu, and Percy Liang. Transforming question answering datasets into natural language inference datasets. *arXiv preprint arXiv:1809.02922*, 2018. URL <https://arxiv.org/pdf/1809.02922.pdf>.
- Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. Bert: Pre-training of deep bidirectional transformers for language understanding. *arXiv preprint arXiv:1810.04805*, 2018.
- Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. BERT: Pre-training of deep bidirectional transformers for language understanding. In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers)*, pages 4171–4186, Minneapolis, Minnesota, June 2019. Association for Computational

- Linguistics. doi: 10.18653/v1/N19-1423. URL <https://aclanthology.org/N19-1423>.
- William B Dolan and Chris Brockett. Automatically constructing a corpus of sentential paraphrases. In *Proceedings of the Third International Workshop on Paraphrasing (IWP2005)*, 2005. URL <https://www.aclweb.org/anthology/I05-5002.pdf>.
- Timothy Dozat and Christopher D Manning. Deep biaffine attention for neural dependency parsing. In *5th International Conference on Learning Representations, ICLR 2017, Toulon, France, April 24-26, 2017, Conference Track Proceedings*. OpenReview.net, 2017. URL <https://openreview.net/forum?id=Hk95PK9le>.
- Matthew S Dryer. The Greenbergian word order correlations. *Language*, pages 81–138, 1992.
- Allyson Ettinger. What BERT is not: Lessons from a new suite of psycholinguistic diagnostics for language models. *Transactions of the Association for Computational Linguistics*, 8:34–48, 2020. doi: 10.1162/tacl_a_00298. URL <https://www.aclweb.org/anthology/2020.tacl-1.3>.
- Jerry A Fodor and Zenon W Pylyshyn. Connectionism and cognitive architecture: A critical analysis. *Cognition*, 28(1-2):3–71, 1988.
- Jonathan Frankle, David J Schwab, and Ari S Morcos. Training batchnorm and only batchnorm: On the expressive power of random features in cnns. *arXiv preprint arXiv:2003.00152*, 2020. URL <https://arxiv.org/abs/2003.00152>.
- Gottlob Frege. Sense and reference. *The philosophical review*, 1948.
- Herve Gallaire and Jack Minker. *Logic and Data Bases*. Perseus Publishing, 1978.

Kanishk Gandhi and Brenden M Lake. Mutual exclusivity as a challenge for deep neural networks. In H. Larochelle, M. Ranzato, R. Hadsell, M. F. Balcan, and H. Lin, editors, *Advances in Neural Information Processing Systems*, volume 33, pages 14182–14192. Curran Associates, Inc., 2020. URL <https://proceedings.neurips.cc/paper/2020/file/a378383b89e6719e15cd1aa45478627c-Paper.pdf>.

Jon Gauthier and Roger Levy. Linking artificial and human neural representations of language. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, pages 529–539, Hong Kong, China, November 2019. Association for Computational Linguistics. doi: 10.18653/v1/D19-1050. URL <https://aclanthology.org/D19-1050>.

Jon Gauthier, Jennifer Hu, Ethan Wilcox, Peng Qian, and Roger Levy. SyntaxGym: An online platform for targeted evaluation of language models. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics: System Demonstrations*, pages 70–76, Online, July 2020. Association for Computational Linguistics. doi: 10.18653/v1/2020.acl-demos.10. URL <https://www.aclweb.org/anthology/2020.acl-demos.10>.

Danilo Giampiccolo, Bernardo Magnini, Ido Dagan, and William B Dolan. The third PASCAL recognizing textual entailment challenge. In *Proceedings of the ACL-PASCAL workshop on textual entailment and paraphrasing*, pages 1–9, 2007. URL <https://www.aclweb.org/anthology/W07-1401.pdf>.

Mario Julianelli, Jack Harding, Florian Mohnert, Dieuwke Hupkes, and Willem Zuidema. Under the hood: Using diagnostic classifiers to investigate and improve how language models track agreement information. In *Proceedings of the 2018 EMNLP Workshop BlackboxNLP: Analyzing and Interpreting Neural Networks for NLP*, pages 240–248, Brussels, Belgium, November 2018. Association for Computational

- Linguistics. doi: 10.18653/v1/W18-5426. URL <https://aclanthology.org/W18-5426>.
- Goran Glavaš and Ivan Vulić. Is supervised syntactic parsing beneficial for language understanding tasks? an empirical investigation. In *Proceedings of the 16th Conference of the European Chapter of the Association for Computational Linguistics: Main Volume*, pages 3090–3104, Online, April 2021. Association for Computational Linguistics. URL <https://aclanthology.org/2021.eacl-main.270>.
- Yoav Goldberg. Assessing BERT’s syntactic abilities. *arXiv preprint arXiv:1901.05287*, 2019a. URL <https://arxiv.org/abs/1901.05287>.
- Yoav Goldberg. Assessing BERT’s Syntactic Abilities. *CoRR*, page 4, 2019b. URL <https://arxiv.org/pdf/1901.05287.pdf>.
- Joseph Greenberg. Some universals of grammar with particular reference to the order of meaningful elements. In J. Greenberg, ed., *Universals of Language*. 73-113. Cambridge, MA., 1963. URL <https://wals.info/refdb/record/Greenberg-1963>.
- Kristina Gulordava, Piotr Bojanowski, Edouard Grave, Tal Linzen, and Marco Baroni. Colorless green recurrent networks dream hierarchically. In *Proceedings of the 2018 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long Papers)*, pages 1195–1205, New Orleans, Louisiana, June 2018a. Association for Computational Linguistics. doi: 10.18653/v1/N18-1108. URL <https://www.aclweb.org/anthology/N18-1108>.
- Kristina Gulordava, Piotr Bojanowski, Edouard Grave, Tal Linzen, and Marco Baroni. Colorless green recurrent networks dream hierarchically. *arXiv:1803.11138 [cs]*, March 2018b. URL <http://arxiv.org/abs/1803.11138>.
- Ashim Gupta, Giorgi Kvernadze, and Vivek Srikumar. BERT & family eat word salad:

Experiments with text understanding. *AAAI*, 2021. URL <https://arxiv.org/abs/2101.03453>.

Suchin Gururangan, Swabha Swayamdipta, Omer Levy, Roy Schwartz, Samuel Bowman, and Noah A. Smith. Annotation artifacts in natural language inference data. In *Proceedings of the 2018 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 2 (Short Papers)*, pages 107–112, New Orleans, Louisiana, June 2018a. Association for Computational Linguistics. doi: 10.18653/v1/N18-2017. URL <https://www.aclweb.org/anthology/N18-2017>.

Suchin Gururangan, Swabha Swayamdipta, Omer Levy, Roy Schwartz, Samuel Bowman, and Noah A. Smith. Annotation artifacts in natural language inference data. In *Proceedings of the 2018 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 2 (Short Papers)*, pages 107–112, 2018b.

R Bar Haim, Ido Dagan, Bill Dolan, Lisa Ferro, Danilo Giampiccolo, Bernardo Magnini, and Idan Szpektor. The second PASCAL recognising textual entailment challenge. In *Proceedings of the Second PASCAL Challenges Workshop on Recognising Textual Entailment*, 2006. URL <http://citeseerx.ist.psu.edu/viewdoc/download?doi=10.1.1.60.8552&rep=rep1&type=pdf>.

Rowan Hall Maudslay, Josef Valvoda, Tiago Pimentel, Adina Williams, and Ryan Cotterell. A tale of a probe and a parser. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pages 7389–7395, Online, July 2020. Association for Computational Linguistics. doi: 10.18653/v1/2020.acl-main.659. URL <https://aclanthology.org/2020.acl-main.659>.

Will Hamilton, Payal Bajaj, Marinka Zitnik, Dan Jurafsky, and Jure Leskovec.

- Embedding logical queries on knowledge graphs. In S. Bengio, H. Wallach, H. Larochelle, K. Grauman, N. Cesa-Bianchi, and R. Garnett, editors, *Advances in Neural Information Processing Systems 31*, pages 2026–2037. Curran Associates, Inc., 2018. URL <http://papers.nips.cc/paper/7473-embedding-logical-queries-on-knowledge-graphs.pdf>.
- Zellig S Harris. Distributional structure. *Word*, 1954a.
- Zellig S Harris. Distributional structure. *Word*, 10(2-3):146–162, 1954b. URL <https://www.tandfonline.com/doi/pdf/10.1080/00437956.1954.11659520>.
- Robert Hawkins, Takateru Yamakoshi, Thomas Griffiths, and Adele Goldberg. Investigating representations of verb bias in neural language models. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 4653–4663, Online, November 2020. Association for Computational Linguistics. doi: 10.18653/v1/2020.emnlp-main.376. URL <https://www.aclweb.org/anthology/2020.emnlp-main.376>.
- Irene Heim and Angelika Kratzer. *Semantics in generative grammar*. Blackwell Oxford, 1998.
- Karl Moritz Hermann, Tomas Kociský, Edward Grefenstette, Lasse Espeholt, Will Kay, Mustafa Suleyman, and Phil Blunsom. Teaching machines to read and comprehend. In *Advances in Neural Information Processing Systems*, pages 1693–1701, 2015.
- John Hewitt and Christopher D. Manning. A structural probe for finding syntax in word representations. In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers)*, pages 4129–4138, Minneapolis, Minnesota, June 2019a. Association for Computational Linguistics. doi: 10.18653/v1/N19-1419. URL <https://www.aclweb.org/anthology/N19-1419>.

John Hewitt and Christopher D Manning. A structural probe for finding syntax in word representations. In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers)*, pages 4129–4138, 2019b. URL <https://www.aclweb.org/anthology/N19-1419.pdf>.

Sepp Hochreiter and Jürgen Schmidhuber. Long short-term memory. *Neural computation*, 9(8):1735–1780, 1997.

Matthew Honnibal, Ines Montani, Sofie Van Landeghem, and Adriane Boyd. spaCy: Industrial-strength Natural Language Processing in Python, 2020a. URL <https://doi.org/10.5281/zenodo.1212303>.

Matthew Honnibal, Ines Montani, Sofie Van Landeghem, and Adriane Boyd. spaCy: Industrial-strength Natural Language Processing in Python, 2020b. URL <https://doi.org/10.5281/zenodo.1212303>.

Hai Hu, Kyle Richardson, Liang Xu, Lu Li, Sandra Kübler, and Lawrence Moss. OC-NLI: Original Chinese Natural Language Inference. In *Findings of the Association for Computational Linguistics: EMNLP 2020*, pages 3512–3526, Online, November 2020a. Association for Computational Linguistics. doi: 10.18653/v1/2020.findings-emnlp.314. URL <https://www.aclweb.org/anthology/2020.findings-emnlp.314>.

Jennifer Hu, Jon Gauthier, Peng Qian, Ethan Wilcox, and Roger Levy. A systematic assessment of syntactic generalization in neural language models. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pages 1725–1744, Online, July 2020b. Association for Computational Linguistics. doi: 10.18653/v1/2020.acl-main.158. URL <https://www.aclweb.org/anthology/2020.acl-main.158>.

Drew Arad Hudson and Christopher D. Manning. Compositional attention networks for machine reasoning. In *International Conference on Learning Representations*, 2018. URL <https://openreview.net/forum?id=S1Euwz-Rb>.

Dieuwke Hupkes, Sara Veldhoen, and Willem Zuidema. Visualisation and ‘diagnostic classifiers’ reveal how recurrent and recursive neural networks process hierarchical structure. *Journal of Artificial Intelligence Research*, 61:907–926, 2018.

Sarthak Jain and Byron C. Wallace. Attention is not Explanation. In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers)*, pages 3543–3556, Minneapolis, Minnesota, June 2019. Association for Computational Linguistics. doi: 10.18653/v1/N19-1357. URL <https://aclanthology.org/N19-1357>.

Ganesh Jawahar, Benoît Sagot, and Djamel Seddah. What does BERT learn about the structure of language? In *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics*, pages 3651–3657, Florence, Italy, July 2019a. Association for Computational Linguistics. doi: 10.18653/v1/P19-1356. URL <https://www.aclweb.org/anthology/P19-1356>.

Ganesh Jawahar, Benoît Sagot, and Djamel Seddah. What Does BERT Learn about the Structure of Language? In *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics*, pages 3651–3657, Florence, Italy, 2019b. Association for Computational Linguistics. doi: 10.18653/v1/P19-1356.

Paloma Jeretic, Alex Warstadt, Suvrat Bhooshan, and Adina Williams. Are natural language inference models IMPPRESsive? Learning IMPlicature and PRESupposition. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pages 8690–8705, Online, July 2020. Association for Computational Linguistics. doi:

- 10.18653/v1/2020.acl-main.768. URL <https://www.aclweb.org/anthology/2020.acl-main.768>.
- Robin Jia and Percy Liang. Adversarial examples for evaluating reading comprehension systems. In *Proceedings of the 2017 Conference on Empirical Methods in Natural Language Processing*, pages 2021–2031, 2017.
- Justin Johnson, Bharath Hariharan, Laurens van der Maaten, Li Fei-Fei, C Lawrence Zitnick, and Ross Girshick. Clevr: A diagnostic dataset for compositional language and elementary visual reasoning. In *Computer Vision and Pattern Recognition (CVPR), 2017 IEEE Conference on*, pages 1988–1997. IEEE, 2017.
- Jaap Jumelet and Dieuwke Hupkes. Do language models understand anything? on the ability of LSTMs to understand negative polarity items. In *Proceedings of the 2018 EMNLP Workshop BlackboxNLP: Analyzing and Interpreting Neural Networks for NLP*, pages 222–231, Brussels, Belgium, November 2018. Association for Computational Linguistics. doi: 10.18653/v1/W18-5424. URL <https://aclanthology.org/W18-5424>.
- Jaap Jumelet, Willem Zuidema, and Dieuwke Hupkes. Analysing neural language models: Contextual decomposition reveals default reasoning in number and gender assignment. In *Proceedings of the 23rd Conference on Computational Natural Language Learning (CoNLL)*, pages 1–11, Hong Kong, China, November 2019. Association for Computational Linguistics. doi: 10.18653/v1/K19-1001. URL <https://aclanthology.org/K19-1001>.
- Jaap Jumelet, Milica Denic, Jakub Szymanik, Dieuwke Hupkes, and Shane Steinert-Threlkeld. Language models use monotonicity to assess NPI licensing. In *Findings of the Association for Computational Linguistics: ACL-IJCNLP 2021*, pages 4958–4969, Online, August 2021. Association for Computational Linguistics. doi:

- 10.18653/v1/2021.findings-acl.439. URL <https://aclanthology.org/2021.findings-acl.439>.
- Ronald M Kaplan and Joan Bresnan. Formal system for grammatical representation. *Formal Issues in Lexical-Functional Grammar*, 1995.
- Dimitri Kartsaklis, Mohammad Taher Pilehvar, and Nigel Collier. Mapping text to knowledge graph entities using multi-sense lstms. *arXiv preprint arXiv:1808.07724*, 2018.
- Hirokatsu Kataoka, Kazushige Okayasu, Asato Matsumoto, Eisuke Yamagata, Ryosuke Yamada, Nakamasa Inoue, Akio Nakamura, and Yutaka Satoh. Pre-training without Natural Images. In *Proceedings of the Asian Conference on Computer Vision*, 2020. URL https://openaccess.thecvf.com/content/ACCV2020/papers/Kataoka_Pre-training_without_Natural_Images_ACCV_2020_paper.pdf.
- Divyansh Kaushik and Zachary C Lipton. How much reading does reading comprehension require? a critical investigation of popular benchmarks. In *Proceedings of the 2018 Conference on Empirical Methods in Natural Language Processing*, pages 5010–5015, 2018.
- Yigitcan Kaya, Sanghyun Hong, and Tudor Dumitras. Shallow-Deep Networks: Understanding and mitigating network overthinking. In *Proceedings of the 2019 International Conference on Machine Learning (ICML)*, Long Beach, CA, Jun 2019. URL <https://arxiv.org/abs/1810.07052>.
- Nikita Kitaev, Steven Cao, and Dan Klein. Multilingual constituency parsing with self-attention and pre-training. In *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics*, pages 3499–3505, Florence, Italy, July 2019. As-

- society for Computational Linguistics. doi: 10.18653/v1/P19-1340. URL <https://aclanthology.org/P19-1340>.
- Jordan Kodner and Nitish Gupta. Overestimation of syntactic representation in neural language models. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pages 1757–1762, Online, July 2020. Association for Computational Linguistics. doi: 10.18653/v1/2020.acl-main.160. URL <https://www.aclweb.org/anthology/2020.acl-main.160>.
- Philippe Laban, Luke Dai, Lucas Bandarkar, and Marti A. Hearst. Can transformer models measure coherence in text: Re-thinking the shuffle test. In *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing (Volume 2: Short Papers)*, pages 1058–1064, Online, August 2021. Association for Computational Linguistics. doi: 10.18653/v1/2021.acl-short.134. URL <https://aclanthology.org/2021.acl-short.134>.
- Brenden Lake and Marco Baroni. Generalization without systematicity: On the compositional skills of sequence-to-sequence recurrent networks. In *International Conference on Machine Learning*, pages 2879–2888, 2018.
- Yair Lakretz, German Kruszewski, Theo Desbordes, Dieuwke Hupkes, Stanislas Dehaene, and Marco Baroni. The emergence of number and syntax units in LSTM language models. In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers)*, pages 11–20, Minneapolis, Minnesota, June 2019. Association for Computational Linguistics. doi: 10.18653/v1/N19-1002. URL <https://aclanthology.org/N19-1002>.
- Yair Lakretz, Dieuwke Hupkes, Alessandra Vergallito, Marco Marelli, Marco Baroni,

- and Stanislas Dehaene. Mechanisms for handling nested dependencies in neural-network language models and humans. *Cognition*, page 104699, 2021. ISSN 0010-0277. doi: <https://doi.org/10.1016/j.cognition.2021.104699>. URL <https://www.sciencedirect.com/science/article/pii/S0010027721001189>.
- Thomas K Landauer and Susan T Dumais. A solution to plato’s problem: The latent semantic analysis theory of acquisition, induction, and representation of knowledge. *Psychological review*, 104(2):211, 1997. URL <https://psycnet.apa.org/record/1997-03612-001>.
- Mike Lewis, Yinhan Liu, Naman Goyal, Marjan Ghazvininejad, Abdelrahman Mohamed, Omer Levy, Veselin Stoyanov, and Luke Zettlemoyer. BART: Denoising sequence-to-sequence pre-training for natural language generation, translation, and comprehension. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pages 7871–7880, Online, July 2020. Association for Computational Linguistics. doi: 10.18653/v1/2020.acl-main.703. URL <https://www.aclweb.org/anthology/2020.acl-main.703>.
- Peiguang Li, Hongfeng Yu, Wenkai Zhang, Guangluan Xu, and Xian Sun. SA-NLI: A supervised attention based framework for natural language inference. *Neurocomputing*, 2020. URL <https://doi.org/10.1016/j.neucom.2020.03.092>.
- Yongjie Lin, Yi Chern Tan, and Robert Frank. Open sesame: Getting inside BERT’s linguistic knowledge. In *Proceedings of the 2019 ACL Workshop BlackboxNLP: Analyzing and Interpreting Neural Networks for NLP*, pages 241–253, Florence, Italy, August 2019. Association for Computational Linguistics. doi: 10.18653/v1/W19-4825. URL <https://www.aclweb.org/anthology/W19-4825>.
- Tal Linzen and Marco Baroni. Syntactic structure from deep learning. *An-*

- nual Review of Linguistics, 2021. URL <https://doi.org/10.1146/annurev-linguistics-032020-051035>.
- Tal Linzen, Emmanuel Dupoux, and Yoav Goldberg. Assessing the ability of LSTMs to learn syntax-sensitive dependencies. *Transactions of the Association for Computational Linguistics*, 4:521–535, 2016. doi: 10.1162/tacl_a_00115. URL <https://www.aclweb.org/anthology/Q16-1037>.
- Yinhan Liu, Myle Ott, Naman Goyal, Jingfei Du, Mandar Joshi, Danqi Chen, Omer Levy, Mike Lewis, Luke Zettlemoyer, and Veselin Stoyanov. Roberta: A robustly optimized BERT pretraining approach. *arXiv*, 2019a. URL <https://arxiv.org/abs/1907.11692>.
- Yinhan Liu, Myle Ott, Naman Goyal, Jingfei Du, Mandar Joshi, Danqi Chen, Omer Levy, Mike Lewis, Luke Zettlemoyer, and Veselin Stoyanov. RoBERTa: A Robustly Optimized BERT Pretraining Approach. *arXiv:1907.11692 [cs]*, July 2019b. URL <http://arxiv.org/abs/1907.11692>.
- Bill MacCartney and Christopher D. Manning. Natural logic for textual inference. In *Proceedings of the ACL-PASCAL Workshop on Textual Entailment and Paraphrasing*, pages 193–200, Prague, June 2007. Association for Computational Linguistics. URL <https://aclanthology.org/W07-1431>.
- Christopher D Manning. Computational linguistics and deep learning. *Computational Linguistics*, 2015.
- Christopher D. Manning, Kevin Clark, John Hewitt, Urvashi Khandelwal, and Omer Levy. Emergent linguistic structure in artificial neural networks trained by self-supervision. *Proceedings of the National Academy of Sciences*, 117(48):30046–30054, 2020a. ISSN 0027-8424. doi: 10.1073/pnas.1907367117. URL <https://www.pnas.org/content/117/48/30046>.

Christopher D Manning, Kevin Clark, John Hewitt, Urvashi Khandelwal, and Omer Levy. Emergent linguistic structure in artificial neural networks trained by self-supervision. *Proceedings of the National Academy of Sciences*, 117(48):30046–30054, 2020b. URL <https://www.ncbi.nlm.nih.gov/pmc/articles/PMC7720155/>.

Rebecca Marvin and Tal Linzen. Targeted syntactic evaluation of language models. In *Proceedings of the 2018 Conference on Empirical Methods in Natural Language Processing*, pages 1192–1202, Brussels, Belgium, October-November 2018. Association for Computational Linguistics. doi: 10.18653/v1/D18-1151. URL <https://www.aclweb.org/anthology/D18-1151>.

R. Thomas McCoy, Junghyun Min, and Tal Linzen. BERTs of a feather do not generalize together: Large variability in generalization across models with similar test set performance. In *Proceedings of the Third BlackboxNLP Workshop on Analyzing and Interpreting Neural Networks for NLP*, pages 217–227, Online, November 2020. Association for Computational Linguistics. doi: 10.18653/v1/2020.blackboxnlp-1.21. URL <https://www.aclweb.org/anthology/2020.blackboxnlp-1.21>.

Tom McCoy, Ellie Pavlick, and Tal Linzen. Right for the wrong reasons: Diagnosing syntactic heuristics in natural language inference. In *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics*, pages 3428–3448, Florence, Italy, July 2019. Association for Computational Linguistics. doi: 10.18653/v1/P19-1334. URL <https://www.aclweb.org/anthology/P19-1334>.

Tomás Mikolov, Kai Chen, Greg Corrado, and Jeffrey Dean. Efficient estimation of word representations in vector space. In Yoshua Bengio and Yann LeCun, editors, *1st International Conference on Learning Representations, ICLR 2013, Scottsdale, Arizona, USA, May 2-4, 2013, Workshop Track Proceedings*, 2013a. URL <http://arxiv.org/abs/1301.3781>.

- Tomas Mikolov, Ilya Sutskever, Kai Chen, Greg S Corrado, and Jeff Dean. Distributed representations of words and phrases and their compositionality. In *Advances in Neural Information Processing Systems*, pages 3111–3119, 2013b. URL <https://arxiv.org/pdf/1310.4546.pdf>.
- Alexander Miller, Will Feng, Dhruv Batra, Antoine Bordes, Adam Fisch, Jiasen Lu, Devi Parikh, and Jason Weston. Parlai: A dialog research software platform. In *Proceedings of the 2017 Conference on Empirical Methods in Natural Language Processing: System Demonstrations*, pages 79–84, 2017.
- Francis Mollica, Matthew Siegelman, Evgeniia Diachek, Steven T Piantadosi, Zachary Mineroff, Richard Futrell, Hope Kean, Peng Qian, and Evelina Fedorenko. Composition is the core driver of the language-selective network. *Neurobiology of Language*, 1(1):104–134, 2020. URL <https://direct.mit.edu/nol/article/1/1/104/10024/Composition-is-the-Core-Driver-of-the-Language>.
- Richard Montague. Universal grammar. *Theoria*, 1970.
- Nasrin Mostafazadeh, Nathanael Chambers, Xiaodong He, Devi Parikh, Dhruv Batra, Lucy Vanderwende, Pushmeet Kohli, and James Allen. A corpus and cloze evaluation for deeper understanding of commonsense stories. In *Proceedings of NAACL-HLT*, pages 839–849, 2016.
- Aakanksha Naik, Abhilasha Ravichander, Norman Sadeh, Carolyn Rose, and Graham Neubig. Stress test evaluation for natural language inference. In *Proceedings of the 27th International Conference on Computational Linguistics*, pages 2340–2353, Santa Fe, New Mexico, USA, August 2018. Association for Computational Linguistics. URL <https://www.aclweb.org/anthology/C18-1198>.
- Aakanksha Naik, Abhilasha Ravichander, Carolyn Rose, and Eduard Hovy. Exploring numeracy in word embeddings. In *Proceedings of the 57th Annual Meeting of*

- the Association for Computational Linguistics*, pages 3374–3380, Florence, Italy, July 2019. Association for Computational Linguistics. doi: 10.18653/v1/P19-1329. URL <https://www.aclweb.org/anthology/P19-1329>.
- Nikita Nangia and Samuel R. Bowman. Human vs. muppet: A conservative estimate of human performance on the GLUE benchmark. In *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics*, pages 4566–4575, Florence, Italy, July 2019. Association for Computational Linguistics. doi: 10.18653/v1/P19-1449. URL <https://www.aclweb.org/anthology/P19-1449>.
- Tri Nguyen, Mir Rosenberg, Xia Song, Jianfeng Gao, Saurabh Tiwary, Rangan Majumder, and Li Deng. Ms marco: A human generated machine reading comprehension dataset. *arXiv preprint arXiv:1611.09268*, 2016.
- Yixin Nie, Adina Williams, Emily Dinan, Mohit Bansal, Jason Weston, and Douwe Kiela. Adversarial NLI: A new benchmark for natural language understanding. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pages 4885–4901, Online, July 2020. Association for Computational Linguistics. doi: 10.18653/v1/2020.acl-main.441. URL <https://www.aclweb.org/anthology/2020.acl-main.441>.
- Myle Ott, Sergey Edunov, Alexei Baevski, Angela Fan, Sam Gross, Nathan Ng, David Grangier, and Michael Auli. fairseq: A fast, extensible toolkit for sequence modeling. In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics (Demonstrations)*, pages 48–53, Minneapolis, Minnesota, June 2019. Association for Computational Linguistics. doi: 10.18653/v1/N19-4009. URL <https://www.aclweb.org/anthology/N19-4009>.
- Isabel Papadimitriou and Dan Jurafsky. Learning Music Helps You Read: Using transfer to study linguistic structure in language models. In *Proceedings of the 2020 Con-*

- ference on Empirical Methods in Natural Language Processing (EMNLP), pages 6829–6839, Online, November 2020. Association for Computational Linguistics. doi: 10.18653/v1/2020.emnlp-main.554. URL <https://aclanthology.org/2020.emnlp-main.554>.
- Isabel Papadimitriou, Ethan A. Chi, Richard Futrell, and Kyle Mahowald. Deep subjecthood: Higher-order grammatical features in multilingual BERT. In *Proceedings of the 16th Conference of the European Chapter of the Association for Computational Linguistics: Main Volume*, pages 2522–2532, Online, April 2021. Association for Computational Linguistics. URL <https://www.aclweb.org/anthology/2021.eacl-main.215>.
- Prasanna Parthasarathi, Koustuv Sinha, Joelle Pineau, and Adina Williams. Sometimes we want translationese. *arXiv preprint arXiv:2104.07623*, 2021. URL <https://arxiv.org/abs/2104.07623>.
- Adam Paszke, Sam Gross, Soumith Chintala, Gregory Chanan, Edward Yang, Zachary DeVito, Zeming Lin, Alban Desmaison, Luca Antiga, and Adam Lerer. Automatic differentiation in pytorch. 2017.
- Ethan Perez, Douwe Kiela, and Kyunghyun Cho. Rissanen Data Analysis: Examining Dataset Characteristics via Description Length. In *Proceedings of the Thirty-eighth International Conference on Machine Learning (ICML)*, March 2021.
- Matthew E. Peters, Mark Neumann, Mohit Iyyer, Matt Gardner, Christopher Clark, Kenton Lee, and Luke Zettlemoyer. Deep contextualized word representations. In *Proceedings of the 2018 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long Papers)*, pages 2227–2237, New Orleans, Louisiana, June 2018a. Association for Computation Linguistics.

- tional Linguistics. doi: 10.18653/v1/N18-1202. URL <https://aclanthology.org/N18-1202>.
- Matthew E Peters, Mark Neumann, Mohit Iyyer, Matt Gardner, Christopher Clark, Kenton Lee, and Luke Zettlemoyer. Deep contextualized word representations. In *NAACL*, 2018b.
- Thang M Pham, Trung Bui, Long Mai, and Anh Nguyen. Out of order: How important is the sequential order of words in a sentence in natural language understanding tasks? *arXiv preprint arXiv:2012.15180*, 2020a. URL <https://arxiv.org/abs/2012.15180>.
- Thang M. Pham, Trung Bui, Long Mai, and Anh Nguyen. Out of Order: How important is the sequential order of words in a sentence in Natural Language Understanding tasks? *arXiv:2012.15180 [cs]*, December 2020b.
- Tiago Pimentel, Naomi Saphra, Adina Williams, and Ryan Cotterell. Pareto probing: Trading off accuracy for complexity. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 3138–3153, Online, November 2020a. Association for Computational Linguistics. doi: 10.18653/v1/2020.emnlp-main.254. URL <https://aclanthology.org/2020.emnlp-main.254>.
- Tiago Pimentel, Josef Valvoda, Rowan Hall Maudslay, Ran Zmigrod, Adina Williams, and Ryan Cotterell. Information-theoretic probing for linguistic structure. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pages 4609–4622, Online, July 2020b. Association for Computational Linguistics. doi: 10.18653/v1/2020.acl-main.420. URL <https://aclanthology.org/2020.acl-main.420>.
- Adam Poliak, Jason Naradowsky, Aparajita Haldar, Rachel Rudinger, and Benjamin Van Durme. Hypothesis only baselines in natural language inference. In *Proceed-*

- ings of the Seventh Joint Conference on Lexical and Computational Semantics*, pages 180–191, New Orleans, Louisiana, June 2018. Association for Computational Linguistics. doi: 10.18653/v1/S18-2023. URL <https://www.aclweb.org/anthology/S18-2023>.
- Carl Pollard and Ivan A Sag. *Head-driven phrase structure grammar*. University of Chicago Press, 1994. URL <https://web.stanford.edu/group/cslipublications/cslipublications/site/0226674479.shtml>.
- Grusha Prasad, Marten van Schijndel, and Tal Linzen. Using priming to uncover the organization of syntactic representations in neural language models. In *Proceedings of the 23rd Conference on Computational Natural Language Learning (CoNLL)*, pages 66–76, Hong Kong, China, November 2019. Association for Computational Linguistics. doi: 10.18653/v1/K19-1007. URL <https://www.aclweb.org/anthology/K19-1007>.
- Liina Pylkkänen, Douglas K Bemis, and Estibaliz Blanco Elorrieta. Building phrases in language production: An meg study of simple composition. *Cognition*, 2014.
- J R Quinlan. Learning logical definitions from relations. *Mach. Learn.*, 5(3):239–266, August 1990. ISSN 0885-6125. doi: 10.1007/BF00117105.
- Alec Radford, Jeffrey Wu, Rewon Child, David Luan, Dario Amodei, and Ilya Sutskever. Language models are unsupervised multitask learners. *OpenAI blog*, 2019. URL https://d4mucfpksywv.cloudfront.net/better-language-models/language_models_are_unsupervised_multitask_learners.pdf.
- Pranav Rajpurkar, Jian Zhang, Konstantin Lopyrev, and Percy Liang. Squad: 100,000+ questions for machine comprehension of text. In *Proceedings of the 2016 Conference on Empirical Methods in Natural Language Processing*, pages 2383–2392, 2016a.

Pranav Rajpurkar, Jian Zhang, Konstantin Lopyrev, and Percy Liang. SQuAD: 100,000+ questions for machine comprehension of text. In *Proceedings of the 2016 Conference on Empirical Methods in Natural Language Processing*, pages 2383–2392, Austin, Texas, November 2016b. Association for Computational Linguistics. doi: 10.18653/v1/D16-1264. URL <https://aclanthology.org/D16-1264>.

Shauli Ravfogel, Yoav Goldberg, and Francis Tyers. Can LSTM learn to capture agreement? the case of Basque. In *Proceedings of the 2018 EMNLP Workshop BlackboxNLP: Analyzing and Interpreting Neural Networks for NLP*, pages 98–107, Brussels, Belgium, November 2018. Association for Computational Linguistics. doi: 10.18653/v1/W18-5412. URL <https://www.aclweb.org/anthology/W18-5412>.

Shauli Ravfogel, Yoav Goldberg, and Tal Linzen. Studying the inductive biases of RNNs with synthetic variations of natural languages. In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers)*, pages 3532–3542, Minneapolis, Minnesota, June 2019. Association for Computational Linguistics. doi: 10.18653/v1/N19-1356. URL <https://www.aclweb.org/anthology/N19-1356>.

Abhilasha Ravichander, Aakanksha Naik, Carolyn Rose, and Eduard Hovy. EQUATE: A benchmark evaluation framework for quantitative reasoning in natural language inference. In *Proceedings of the 23rd Conference on Computational Natural Language Learning (CoNLL)*, pages 349–361, Hong Kong, China, November 2019. Association for Computational Linguistics. doi: 10.18653/v1/K19-1033. URL <https://www.aclweb.org/anthology/K19-1033>.

Matthew Richardson, Christopher JC Burges, and Erin Renshaw. McTest: A challenge dataset for the open-domain machine comprehension of text. In *Proceedings of the 2013 Conference on Empirical Methods in Natural Language Processing*, pages 193–203, 2013.

- Jorma Rissanen. Universal coding, information, prediction, and estimation. *IEEE Transactions on Information theory*, 30(4):629–636, 1984. URL <https://ieeexplore.ieee.org/document/1056936>.
- Anna Rogers, Olga Kovaleva, and Anna Rumshisky. A primer in BERTology: What we know about how BERT works. *Transactions of the Association for Computational Linguistics*, 8:842–866, 2020a. doi: 10.1162/tacl_a_00349. URL <https://aclanthology.org/2020.tacl-1.54>.
- Anna Rogers, Olga Kovaleva, and Anna Rumshisky. A primer in BERTology: What we know about how BERT works. *Transactions of the Association for Computational Linguistics*, 8:842–866, 2020b. doi: 10.1162/tacl_a_00349. URL <https://www.aclweb.org/anthology/2020.tacl-1.54>.
- Julian Salazar, Davis Liang, Toan Q. Nguyen, and Katrin Kirchhoff. Masked Language Model Scoring. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pages 2699–2712, Online, 2020. Association for Computational Linguistics. doi: 10.18653/v1/2020.acl-main.240.
- Victor Sanh, Lysandre Debut, Julien Chaumond, and Thomas Wolf. Distilbert, a distilled version of bert: smaller, faster, cheaper and lighter, 2020. URL <https://arxiv.org/abs/1910.01108>.
- Adam Santoro, David Raposo, David G Barrett, Mateusz Malinowski, Razvan Pascanu, Peter Battaglia, and Timothy Lillicrap. A simple neural network module for relational reasoning. In *Advances in neural information processing systems*, pages 4967–4976, 2017.
- Adam Santoro, Ryan Faulkner, David Raposo, Jack Rae, Mike Chrzanowski, Theophane Weber, Daan Wierstra, Oriol Vinyals, Razvan Pascanu, and Timothy Lillicrap.

- Relational recurrent neural networks. In *Advances in Neural Information Processing Systems*, pages 7310–7321, 2018.
- Yves Schabes, Anne Abeille, and Aravind K. Joshi. Parsing strategies with ‘lexicalized’ grammars: Application to Tree Adjoining Grammars. In *Coling Budapest 1988 Volume 2: International Conference on Computational Linguistics*, 1988. URL <https://www.aclweb.org/anthology/C88-2121>.
- Eckart Scheerer. Early german approaches to experimental reading research: The contributions of wilhelm wundt and ernst meumann. *Psychological Research*, 1981.
- Christian Scheible and Hinrich Schütze. Cutting recursive autoencoder trees. In Yoshua Bengio and Yann LeCun, editors, *1st International Conference on Learning Representations (ICLR) Scottsdale, Arizona, USA, May 2-4, 2013, Conference Track Proceedings*, 2013. URL <https://arxiv.org/pdf/1301.2811.pdf>.
- Sheng Shen, Alexei Baevski, Ari Morcos, Kurt Keutzer, Michael Auli, and Douwe Kiela. Reservoir transformers. In *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing (Volume 1: Long Papers)*, pages 4294–4309, Online, August 2021. Association for Computational Linguistics. doi: 10.18653/v1/2021.acl-long.331. URL <https://aclanthology.org/2021.acl-long.331>.
- Natalia Silveira, Timothy Dozat, Marie-Catherine De Marneffe, Samuel R Bowman, Miriam Connor, John Bauer, and Christopher D Manning. A gold standard dependency corpus for english. In *LREC*, pages 2897–2904. Citeseer, 2014. URL http://www.lrec-conf.org/proceedings/lrec2014/pdf/1089_Paper.pdf.
- Koustuv Sinha, Shagun Sodhani, Jin Dong, Joelle Pineau, and William L. Hamilton. CLUTRR: A diagnostic benchmark for inductive reasoning from text. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and*

the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP), pages 4506–4515, Hong Kong, China, November 2019. Association for Computational Linguistics. doi: 10.18653/v1/D19-1458. URL <https://aclanthology.org/D19-1458>.

Koustuv Sinha, Robin Jia, Dieuwke Hupkes, Joelle Pineau, Adina Williams, and Douwe Kiela. Masked language modeling and the distributional hypothesis: Order word matters pre-training for little. In *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*, pages 2888–2913, Online and Punta Cana, Dominican Republic, November 2021a. Association for Computational Linguistics. doi: 10.18653/v1/2021.emnlp-main.230. URL <https://aclanthology.org/2021.emnlp-main.230>.

Koustuv Sinha, Robin Jia, Dieuwke Hupkes, Joelle Pineau, Adina Williams, and Douwe Kiela. Masked language modeling and the distributional hypothesis: Order word matters pre-training for little. *arXiv preprint arXiv:2104.06644*, 2021b. URL <https://arxiv.org/abs/2104.06644>.

Koustuv Sinha, Prasanna Parthasarathi, Joelle Pineau, and Adina Williams. UnNatural Language Inference. In *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing (Volume 1: Long Papers)*, pages 7329–7346, Online, August 2021c. Association for Computational Linguistics. doi: 10.18653/v1/2021.acl-long.569. URL <https://aclanthology.org/2021.acl-long.569>.

Koustuv Sinha, Prasanna Parthasarathi, Joelle Pineau, and Adina Williams. UnNatural Language Inference. In *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing (Volume 1: Long Papers)*, pages 7329–7346, Online, August 2021d. As-

- sociation for Computational Linguistics. doi: 10.18653/v1/2021.acl-long.569. URL <https://aclanthology.org/2021.acl-long.569>.
- Joshua Snell and Jonathan Grainger. The sentence superiority effect revisited. *Cognition*, 2017.
- Joshua Snell and Jonathan Grainger. Word position coding in reading is noisy. *Psychonomic bulletin & review*, 26(2):609–615, 2019.
- Richard Socher, Alex Perelygin, Jean Wu, Jason Chuang, Christopher D Manning, Andrew Y Ng, and Christopher Potts. Recursive deep models for semantic compositionality over a sentiment treebank. In *Proceedings of the 2013 conference on empirical methods in natural language processing*, pages 1631–1642, 2013. URL https://nlp.stanford.edu/~socherr/EMNLP2013_RNTN.pdf.
- Shagun Sodhani, Sarath Chandar, and Yoshua Bengio. On Training Recurrent Neural Networks for Lifelong Learning. *arXiv e-prints*, November 2018.
- Yu Su, Huan Sun, Brian Sadler, Mudhakar Srivatsa, Izzeddin Gur, Zenghui Yan, and Xifeng Yan. On generating characteristic-rich question sets for qa evaluation. In *Proceedings of the 2016 Conference on Empirical Methods in Natural Language Processing*, pages 562–572, 2016.
- Whitney Tabor. *Syntactic innovation: A connectionist model*. PhD thesis, 1994. URL <https://www.proquest.com/openview/0a8f7e8a71e058b12053b545ca857fb2/1>.
- Ian Tenney, Dipanjan Das, and Ellie Pavlick. BERT rediscovers the classical NLP pipeline. In *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics*, pages 4593–4601, Florence, Italy, July 2019. Association for Computational Linguistics. doi: 10.18653/v1/P19-1452. URL <https://www.aclweb.org/anthology/P19-1452>.

Hiroshi Toyota. Changes in the constraints of semantic and syntactic congruity on memory across three age groups. *Perceptual and Motor Skills*, 2001. URL <https://pubmed.ncbi.nlm.nih.gov/11453195/>.

Adam Trischler, Tong Wang, Xingdi Yuan, Justin Harris, Alessandro Sordoni, Philip Bachman, and Kaheer Suleman. NewsQA: A machine comprehension dataset. *arXiv preprint*, pages 1–12, 2016. doi: 10.1101/B978-0-12-800077-9/00020-7.

Masatoshi Tsuchiya. Performance impact caused by hidden bias of training data for recognizing textual entailment. In *Proceedings of the Eleventh International Conference on Language Resources and Evaluation (LREC 2018)*, Miyazaki, Japan, May 2018. European Language Resources Association (ELRA). URL <https://www.aclweb.org/anthology/L18-1239>.

Dmitry Ulyanov, Andrea Vedaldi, and Victor Lempitsky. Deep image prior. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pages 9446–9454, 2018. URL <https://arxiv.org/pdf/1711.10925.pdf>.

Marten van Schijndel and Tal Linzen. A neural model of adaptation in reading. In *Proceedings of the 2018 Conference on Empirical Methods in Natural Language Processing*, pages 4704–4710, Brussels, Belgium, October–November 2018. Association for Computational Linguistics. doi: 10.18653/v1/D18-1499. URL <https://www.aclweb.org/anthology/D18-1499>.

Marten van Schijndel, Aaron Mueller, and Tal Linzen. Quantity doesn't buy quality syntax with neural language models. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, pages 5831–5837, Hong Kong, China, November 2019. Association for Computational Linguistics. doi: 10.18653/v1/D19-1592. URL <https://www.aclweb.org/anthology/D19-1592>.

Lucy Vanderwende and William B Dolan. What syntax can contribute in the entailment task. In *Machine Learning Challenges Workshop*. Springer, 2005. URL https://dl.acm.org/doi/10.1007/11736790_11.

Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N Gomez, Łukasz Kaiser, and Illia Polosukhin. Attention is all you need. In I. Guyon, U. V. Luxburg, S. Bengio, H. Wallach, R. Fergus, S. Vishwanathan, and R. Garnett, editors, *Advances in Neural Information Processing Systems*, volume 30. Curran Associates, Inc., 2017. URL <https://proceedings.neurips.cc/paper/2017/file/3f5ee243547dee91fb053c1c4a845aa-Paper.pdf>.

Petar Veličković, Guillem Cucurull, Arantxa Casanova, Adriana Romero, Pietro Liò, and Yoshua Bengio. Graph attention networks. In *International Conference on Learning Representations*, 2018. URL <https://openreview.net/forum?id=rJXMpikCZ>.

Elena Voita and Ivan Titov. Information-theoretic probing with minimum description length. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 183–196, Online, November 2020. Association for Computational Linguistics. doi: 10.18653/v1/2020.emnlp-main.14. URL <https://aclanthology.org/2020.emnlp-main.14>.

Eric Wallace, Shi Feng, Nikhil Kandpal, Matt Gardner, and Sameer Singh. Universal adversarial triggers for attacking and analyzing NLP. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, pages 2153–2162, Hong Kong, China, November 2019. Association for Computational Linguistics. doi: 10.18653/v1/D19-1221. URL <https://www.aclweb.org/anthology/D19-1221>.

Alex Wang, Amanpreet Singh, Julian Michael, Felix Hill, Omer Levy, and Samuel Bow-

- man. GLUE: A multi-task benchmark and analysis platform for natural language understanding. In *Proceedings of the 2018 EMNLP Workshop BlackboxNLP: Analyzing and Interpreting Neural Networks for NLP*, pages 353–355, Brussels, Belgium, November 2018. Association for Computational Linguistics. doi: 10.18653/v1/W18-5446. URL <https://www.aclweb.org/anthology/W18-5446>.
- Alex Wang, Yada Pruksachatkun, Nikita Nangia, Amanpreet Singh, Julian Michael, Felix Hill, Omer Levy, and Samuel R. Bowman. SuperGLUE: A stickier benchmark for general-purpose language understanding systems. In Hanna M. Wallach, Hugo Larochelle, Alina Beygelzimer, Florence d’Alché-Buc, Emily B. Fox, and Roman Garnett, editors, *NeurIPS*, 2019. URL <https://papers.nips.cc/paper/2019/hash/4496bf24afe7fab6f046bf4923da8de6-Abstract.html>.
- Z. Wang, L. Li, D. D. Zeng, and Y. Chen. Attention-based multi-hop reasoning for knowledge graph. In *2018 IEEE International Conference on Intelligence and Security Informatics (ISI)*, pages 211–213, Nov 2018. doi: 10.1109/ISI.2018.8587330.
- Alex Warstadt and Samuel R Bowman. Can neural networks acquire a structural bias from raw linguistic data? In *Proceedings of the 42nd Annual Virtual Meeting of the Cognitive Science Society*, 2020. URL <https://arxiv.org/abs/2007.06761>.
- Alex Warstadt, Yu Cao, Ioana Grosu, Wei Peng, Hagen Blix, Yining Nie, Anna Alsop, Shikha Bordia, Haokun Liu, Alicia Parrish, Sheng-Fu Wang, Jason Phang, Anhad Mohananey, Phu Mon Htut, Paloma Jeretic, and Samuel R. Bowman. Investigating BERT’s knowledge of language: Five analysis methods with NPIs. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, pages 2877–2887, Hong Kong, China, November 2019a. Association for Computational Linguistics. doi: 10.18653/v1/D19-1286. URL <https://www.aclweb.org/anthology/D19-1286>.

Alex Warstadt, Amanpreet Singh, and Samuel R Bowman. Neural network acceptability judgments. *Transactions of the Association for Computational Linguistics*, 7:625–641, 2019b. URL <https://arxiv.org/pdf/1805.12471.pdf>.

Alex Warstadt, Amanpreet Singh, and Samuel R. Bowman. Neural network acceptability judgments. *Transactions of the Association for Computational Linguistics*, 7:625–641, March 2019c. doi: 10.1162/tacl_a_00290. URL <https://www.aclweb.org/anthology/Q19-1040>.

Alex Warstadt, Alicia Parrish, Haokun Liu, Anhad Mohananey, Wei Peng, Sheng-Fu Wang, and Samuel R. Bowman. BLiMP: The benchmark of linguistic minimal pairs for English. *Transactions of the Association for Computational Linguistics*, 8:377–392, 2020a. doi: 10.1162/tacl_a_00321. URL <https://www.aclweb.org/anthology/2020.tacl-1.25>.

Alex Warstadt, Alicia Parrish, Haokun Liu, Anhad Mohananey, Wei Peng, Sheng-Fu Wang, and Samuel R. Bowman. BLiMP: The benchmark of linguistic minimal pairs for English. *Transactions of the Association for Computational Linguistics*, 8:377–392, 2020b. doi: 10.1162/tacl_a_00321. URL <https://aclanthology.org/2020.tacl-1.25>.

Alex Warstadt, Yian Zhang, Xiaocheng Li, Haokun Liu, and Samuel R. Bowman. Learning which features matter: RoBERTa acquires a preference for linguistic generalizations (eventually). In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 217–235, Online, November 2020c. Association for Computational Linguistics. doi: 10.18653/v1/2020.emnlp-main.16. URL <https://aclanthology.org/2020.emnlp-main.16>.

Johannes Welbl, Pontus Stenetorp, and Sebastian Riedel. Constructing datasets for

- multi-hop reading comprehension across documents. *Transactions of the Association of Computational Linguistics*, 6:287–302, 2018.
- Yun Wen, Joshua Snell, and Jonathan Grainger. Parallel, cascaded, interactive processing of words during sentence reading. *Cognition*, 2019.
- Jason Weston, Antoine Bordes, Sumit Chopra, Alexander M Rush, Bart van Merriënboer, Armand Joulin, and Tomas Mikolov. Towards AI-Complete question answering: A set of prerequisite toy tasks. 2015. ISSN 0378-7753. doi: 10.1016/j.jpowsour.2014.09.131.
- Aaron Steven White, Rachel Rudinger, Kyle Rawlins, and Benjamin Van Durme. Lexicosyntactic inference in neural models. In *Proceedings of the 2018 Conference on Empirical Methods in Natural Language Processing*, pages 4717–4724, Brussels, Belgium, October–November 2018. Association for Computational Linguistics. doi: 10.18653/v1/D18-1501. URL <https://www.aclweb.org/anthology/D18-1501>.
- John Wieting and Douwe Kiela. No training required: Exploring random encoders for sentence classification. In *7th International Conference on Learning Representations, ICLR 2019, New Orleans, LA, USA, May 6-9, 2019*. OpenReview.net, 2019. URL <https://openreview.net/forum?id=BkgPajAcY7>.
- Ethan Wilcox, Roger Levy, Takashi Morita, and Richard Futrell. What do RNN language models learn about filler–gap dependencies? In *Proceedings of the 2018 EMNLP Workshop BlackboxNLP: Analyzing and Interpreting Neural Networks for NLP*, pages 211–221, Brussels, Belgium, November 2018. Association for Computational Linguistics. doi: 10.18653/v1/W18-5423. URL <https://www.aclweb.org/anthology/W18-5423>.
- Adina Williams, Andrew Drozdov, and Samuel R. Bowman. Do latent tree learning models identify meaningful structure in sentences? *Transactions of the Association for*

- Computational Linguistics*, 6:253–267, 2018a. doi: 10.1162/tacl_a_00019. URL <https://www.aclweb.org/anthology/Q18-1019>.
- Adina Williams, Andrew Drozdov, and Samuel R Bowman. Do latent tree learning models identify meaningful structure in sentences? *Transactions of the Association for Computational Linguistics*, 6:253–267, 2018b. URL <https://www.aclweb.org/anthology/Q18-1019/>.
- Adina Williams, Nikita Nangia, and Samuel Bowman. A broad-coverage challenge corpus for sentence understanding through inference. In *Proceedings of the 2018 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long Papers)*, pages 1112–1122, New Orleans, Louisiana, June 2018c. Association for Computational Linguistics. doi: 10.18653/v1/N18-1101. URL <https://www.aclweb.org/anthology/N18-1101>.
- Adina Williams, Nikita Nangia, and Samuel Bowman. A broad-coverage challenge corpus for sentence understanding through inference. In *Proceedings of the 2018 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long Papers)*, volume 1, pages 1112–1122, 2018d.
- Ludwig Wittgenstein. *Tractatus Logico-Philosophicus*. Harcourt, Brace & Company, Inc., 1922.
- Thomas Wolf. Some additional experiments extending the tech report "Assessing BERT's Syntactic Abilities" by Yoav Goldberg. page 7, 2019a. URL <https://huggingface.co/bert-syntax/extending-bert-syntax.pdf>.
- Thomas Wolf. Some additional experiments extending the tech report "assessing berts syntactic abilities" by yoav goldberg. Technical report, HuggingFace, 2019b. URL <https://huggingface.co/bert-syntax/extending-bert-syntax.pdf>.

Thomas Wolf, Julien Chaumond, Lysandre Debut, Victor Sanh, Clement Delangue, Anthony Moi, Pierrick Cistac, Morgan Funtowicz, Joe Davison, Sam Shleifer, et al. Transformers: State-of-the-art natural language processing. In *EMNLP: System Demonstrations*, 2020.

Yonghui Wu, Mike Schuster, Zhifeng Chen, Quoc V. Le, Mohammad Norouzi, Wolfgang Macherey, Maxim Krikun, Yuan Cao, Qin Gao, Klaus Macherey, Jeff Klingner, Apurva Shah, Melvin Johnson, Xiaobing Liu, Łukasz Kaiser, Stephan Gouws, Yoshikiyo Kato, Taku Kudo, Hideto Kazawa, Keith Stevens, George Kurian, Nishant Patil, Wei Wang, Cliff Young, Jason Smith, Jason Riesa, Alex Rudnick, Oriol Vinyals, Greg Corrado, Macduff Hughes, and Jeffrey Dean. Google’s neural machine translation system: Bridging the gap between human and machine translation, 2016.

Zhiyong Wu, Yun Chen, Ben Kao, and Qun Liu. Perturbed masking: Parameter-free probing for analyzing and interpreting BERT. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pages 4166–4176, Online, July 2020. Association for Computational Linguistics. doi: 10.18653/v1/2020.acl-main.383. URL <https://www.aclweb.org/anthology/2020.acl-main.383>.

Saining Xie, Alexander Kirillov, Ross Girshick, and Kaiming He. Exploring randomly wired neural networks for image recognition. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, pages 1284–1293, 2019. URL <https://arxiv.org/pdf/1904.01569.pdf>.

Wenhan Xiong, Thien Hoang, and William Yang Wang. Deeppath: A reinforcement learning method for knowledge graph reasoning. In *Proceedings of the 2017 Conference on Empirical Methods in Natural Language Processing*, pages 564–573, 2017.

Wenhan Xiong, Mo Yu, Shiyu Chang, Xiaoxiao Guo, and William Yang Wang. One-

- shot relational learning for knowledge graphs. In *Proceedings of the 2018 Conference on Empirical Methods in Natural Language Processing*, pages 1980–1990, 2018.
- Lang Yu and Allyson Ettinger. On the interplay between fine-tuning and composition in transformers. In *Findings of the Association for Computational Linguistics: ACL-IJCNLP 2021*, pages 2279–2293, Online, August 2021. Association for Computational Linguistics. doi: 10.18653/v1/2021.findings-acl.201. URL <https://aclanthology.org/2021.findings-acl.201>.
- Kelly Zhang and Samuel Bowman. Language modeling teaches you more than translation does: Lessons learned through auxiliary syntactic task analysis. In *Proceedings of the 2018 EMNLP Workshop BlackboxNLP: Analyzing and Interpreting Neural Networks for NLP*, pages 359–361, Brussels, Belgium, November 2018. Association for Computational Linguistics. doi: 10.18653/v1/W18-5448. URL <https://www.aclweb.org/anthology/W18-5448>.
- Tianyi Zhang, Felix Wu, Arzoo Katiyar, Kilian Q. Weinberger, and Yoav Artzi. Revisiting few-sample BERT fine-tuning. In *9th International Conference on Learning Representations, ICLR 2021, Virtual Event, Austria, May 3-7, 2021*. OpenReview.net, March 2021. URL <https://openreview.net/forum?id=c01IH43yUF>.
- Yu Zhang, Zhenghua Li, and Min Zhang. Efficient second-order TreeCRF for neural dependency parsing. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pages 3295–3305, Online, July 2020. Association for Computational Linguistics. doi: 10.18653/v1/2020.acl-main.302. URL <https://aclanthology.org/2020.acl-main.302>.
- Yuan Zhang, Jason Baldridge, and Luheng He. PAWS: Paraphrase adversaries from word scrambling. In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Vol-*

ume 1 (Long and Short Papers), pages 1298–1308, Minneapolis, Minnesota, June 2019. Association for Computational Linguistics. doi: 10.18653/v1/N19-1131. URL <https://aclanthology.org/N19-1131>.

Han Zhao, Zhengdong Lu, and Pascal Poupart. Self-adaptive hierarchical sentence model. In *Proceedings of the 24th International Conference on Artificial Intelligence*, pages 4069–4076, 2015. URL <https://arxiv.org/abs/1504.05070>.

Wangchunshu Zhou, Canwen Xu, Tao Ge, Julian McAuley, Ke Xu, and Furu Wei. Bert loses patience: Fast and robust inference with early exit. In H. Larochelle, M. Ranzato, R. Hadsell, M. F. Balcan, and H. Lin, editors, *Advances in Neural Information Processing Systems*, volume 33, pages 18330–18341. Curran Associates, Inc., 2020. URL <https://proceedings.neurips.cc/paper/2020/file/d4dd111a4fd973394238aca5c05bebe3-Paper.pdf>.

Yukun Zhu, Ryan Kiros, Rich Zemel, Ruslan Salakhutdinov, Raquel Urtasun, Antonio Torralba, and Sanja Fidler. Aligning books and movies: Towards story-like visual explanations by watching movies and reading books. In *Proceedings of the IEEE international conference on computer vision*, pages 19–27, 2015. URL <https://arxiv.org/pdf/1506.06724.pdf>.

Glossary

Transformers A class of models first derived by Vaswani et al. 2017. 2

Acronyms

LLMs Large Language Models. 2

NLU Natural Language Understanding. 4, 5, 24, 32–34, 59, 60

Chapter 8

Appendix

8.1 Org mode auto save

Run the following snippet to auto save and compile in org mode.

```
(defun kdm/org-save-and-export ()
  (interactive)
  (if (and (eq major-mode 'org-mode)
            (ido-local-file-exists-p (concat (file-name-sans-extension (buffer-name)
                                                               (org-latex-export-to-latex)))))

    (add-hook 'after-save-hook 'kdm/org-save-and-export))
```

8.2 Remove “parts” from report

```
(add-to-list 'org-latex-classes
             ' ("report-noparts"
                 "\documentclass[11pt]{report}"
                 ("\\chapter{\\%s}" . "\\chapter*{\\%s}"))
```

```
("\\"section{\%s}" . "\\"section*{\%s}")
("\\"subsection{\%s}" . "\\"subsection*{\%s}")
("\\"subsubsection{\%s}" . "\\"subsubsection*{\%s}"))
```

8.3 Add newpage before a heading

```
(defun org/get-headline-string-element (headline backend info)
  (let ((prop-point (next-property-change 0 headline)))
    (if prop-point (plist-get (text-properties-at prop-point headline) :pa

(defun org/ensure-latex-clearpage (headline backend info)
  (when (org-export-derived-backend-p backend 'latex)
    (let ((elmnt (org/get-headline-string-element headline backend info)))
      (when (member "newpage" (org-element-property :tags elmnt))
        (concat "\\\clearpage\n" headline)))))

(add-to-list 'org-export-filter-headline-functions
             'org/ensure-latex-clearpage)
```

8.4 Glossary and Acronym build using Latexmk

Add the following snippet in the file “`~/.latexmkrc`”: (Source: <https://tex.stackexchange.com/a/44316>)

```
add_cus_dep('glo', 'gls', 0, 'run_makeglossaries');
add_cus_dep('acn', 'acr', 0, 'run_makeglossaries');
```

```

sub run_makeglossaries {
    my ($base_name, $path) = fileparse( $_[0] ); #handle -outdir param by :
    pushd $path; # ... cd-ing into folder first, then running makeglossaries

    if ( $silent ) {
        system "makeglossaries -q '$base_name'"; #unix
        # system "makeglossaries", "-q", "$base_name"; #windows
    }
    else {
        system "makeglossaries '$base_name'"; #unix
        # system "makeglossaries", "$base_name"; #windows
    };
}

popd; # ... and cd-ing back again
}

push @generated_exts, 'glo', 'gls', 'glg';
push @generated_exts, 'acn', 'acr', 'alg';
$clean_ext .= ' %R.ist %R.xdy';

```

8.5 Citation style buffer local

```
(set (make-local-variable 'bibtex-completion-format-citation-functions)
     '((org-mode . my/bibtex-completion-format-citation-org-default-cite)
       ...))
```

8.6 Org latex compiler options

```
(setq org-latex-pdf-process (list "latexmk -f -pdf -%latex -interaction=nonstopmode"))
```

Original value

```
(setq org-latex-pdf-process (list "latexmk -f -pdf %f"))
```

Let us try Fast compile <https://gist.github.com/yig/ba124dfbc8f63762f222>.

```
(setq org-latex-pdf-process (list "latexmk-fast %f"))
```

- Doesn't seem to work from Emacs.
- I need to change the save function to only export in tex. Then, have a separate process run latexmk.
- Using the python package `when-changed` to watch the `thesis.tex` file for change.
- Usage:

```
when-changed thesis.tex latexmk -f -pdf -interaction=nonstopmode -output-d
```

- The pdf does not update. It seems to but not always? No it does. For some reason, compilation takes ages.
- Works with `when-changed`!